




iSMAC



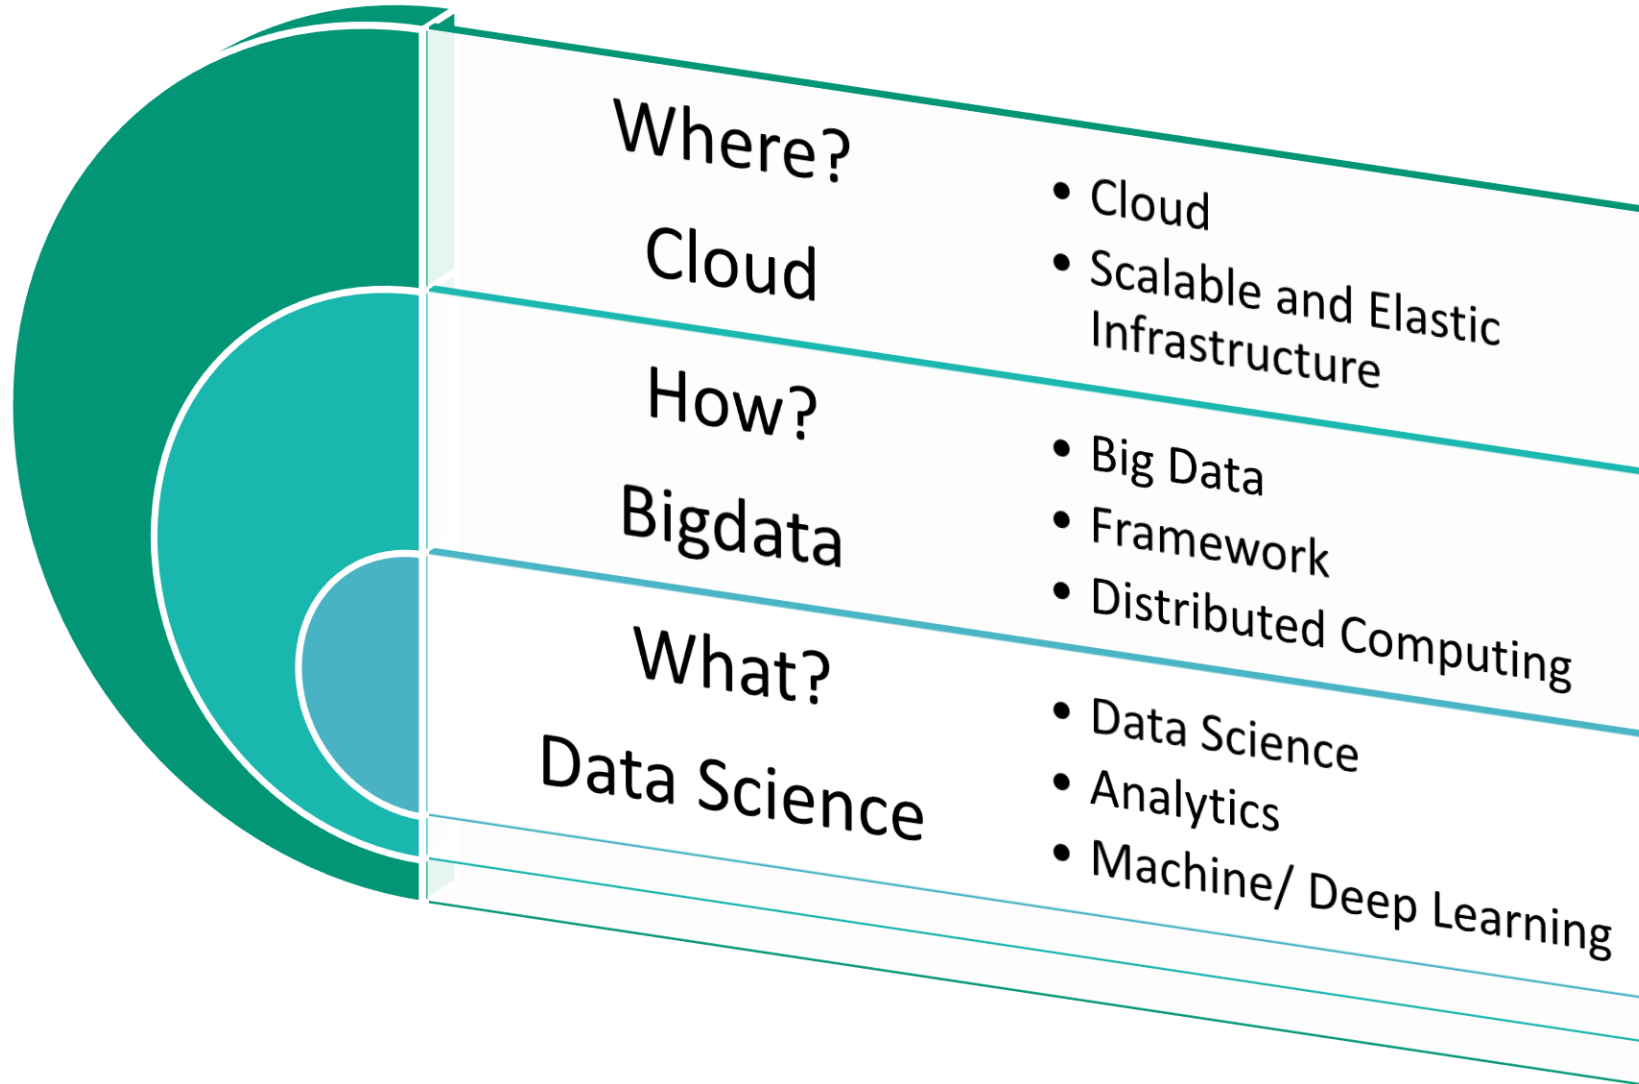
Data Growth



Magic Triangle



Where? How? What?



Spectrum of Analytics

Descriptive

- What happened? - PAST
- Descriptive Statistics
- Data Clustering

Diagnostic

- Why did it happen? / Why is it happening?
- Sensitivity Analysis

Predictive

- What will happen?
- Linear and Logistic Regression

Prescriptive

- What should I do? What should happen?
- Simulation
- Non Linear Programming

Data Science Venn

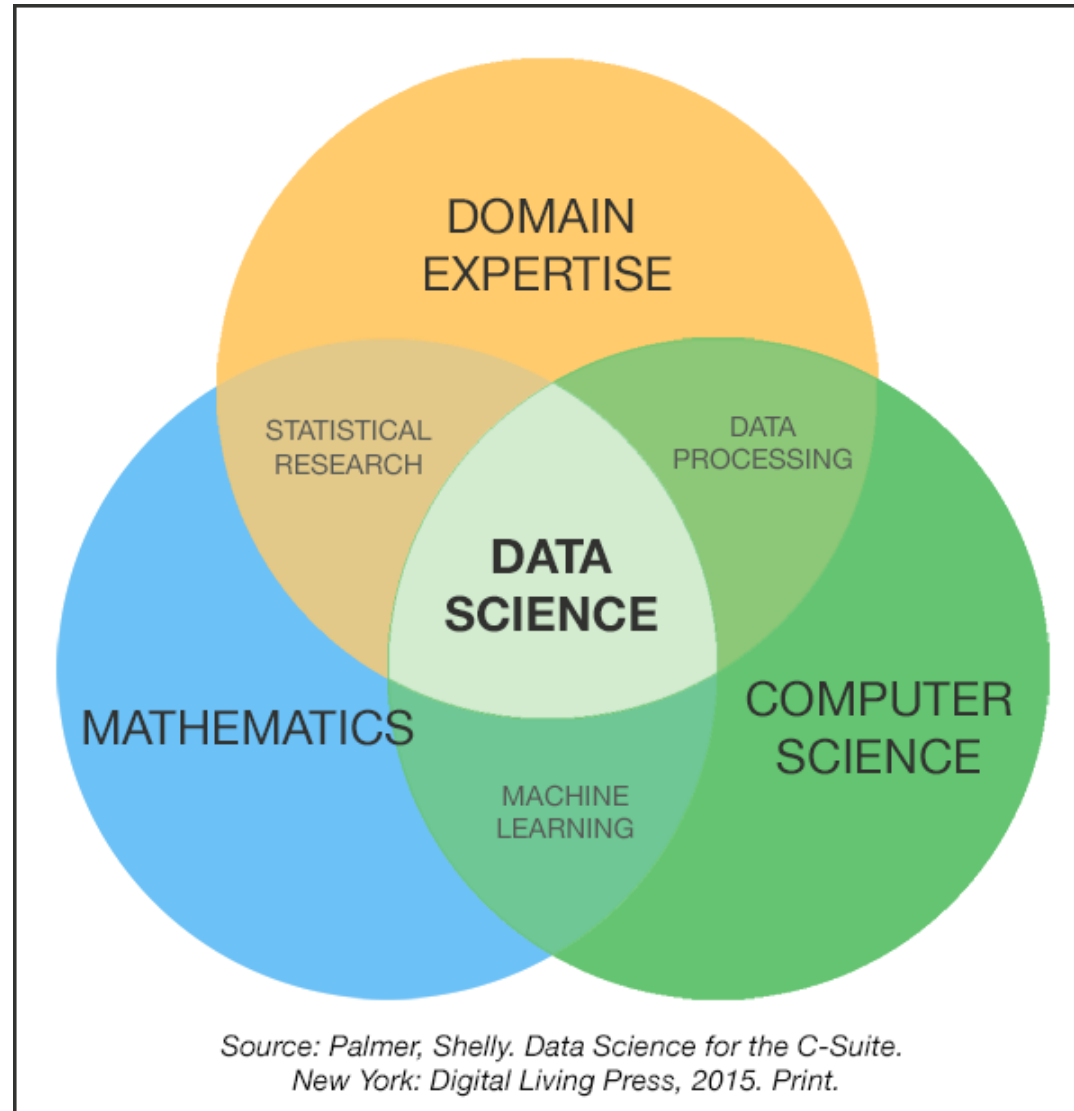






Image Source : <https://pixabay.com>



Algorithm



Model

Birds





Animals




Human






Model

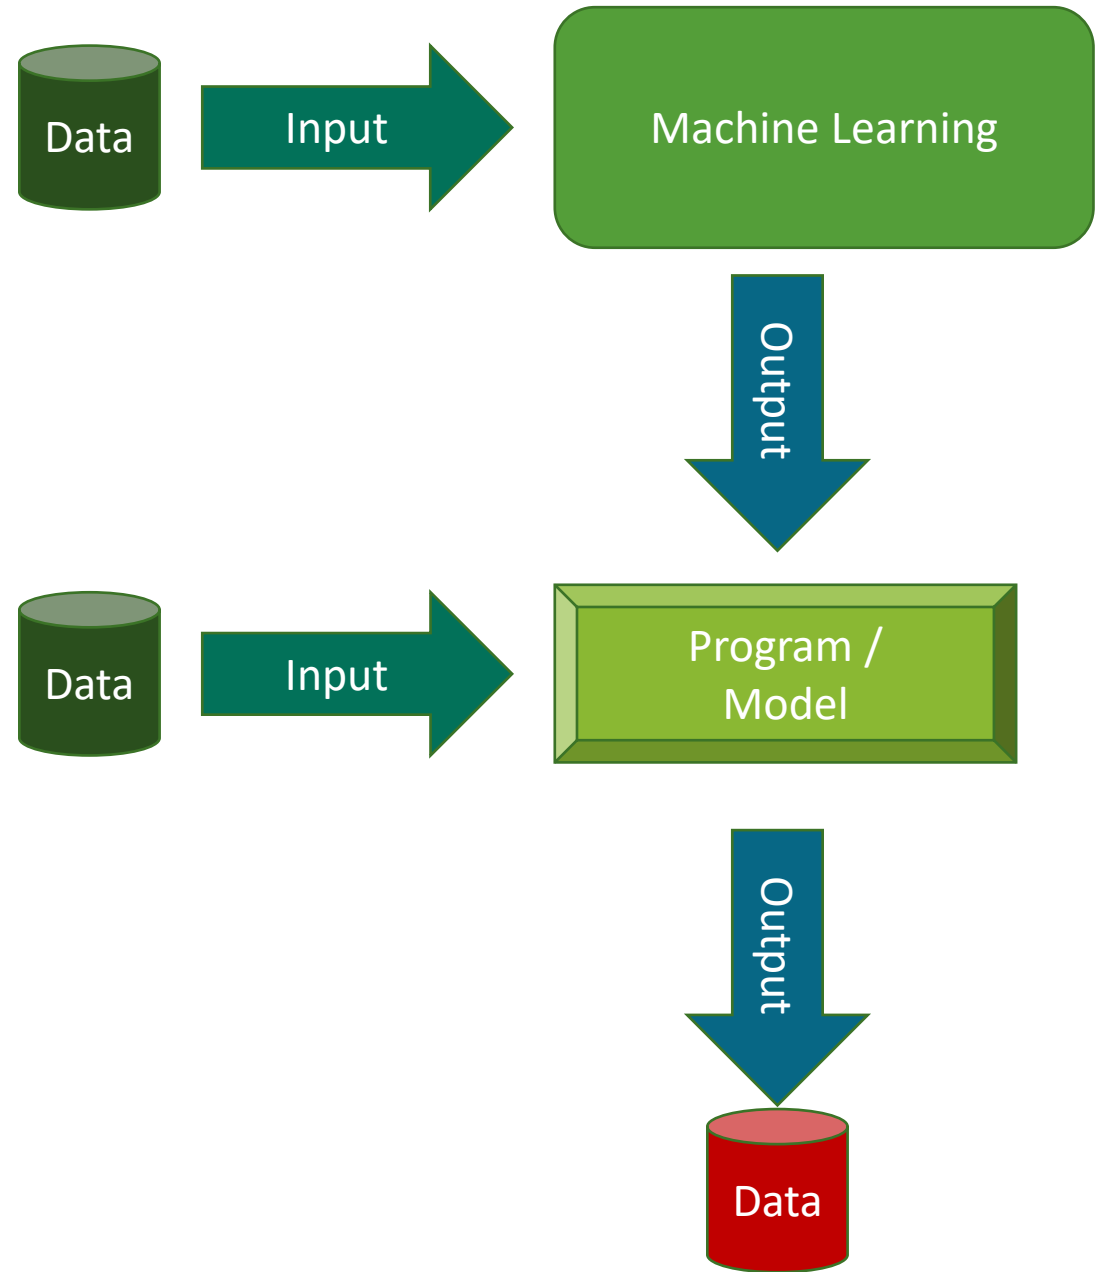



Bird – 7%
Animal – 86%
Human – 6%
Don't know – 1%

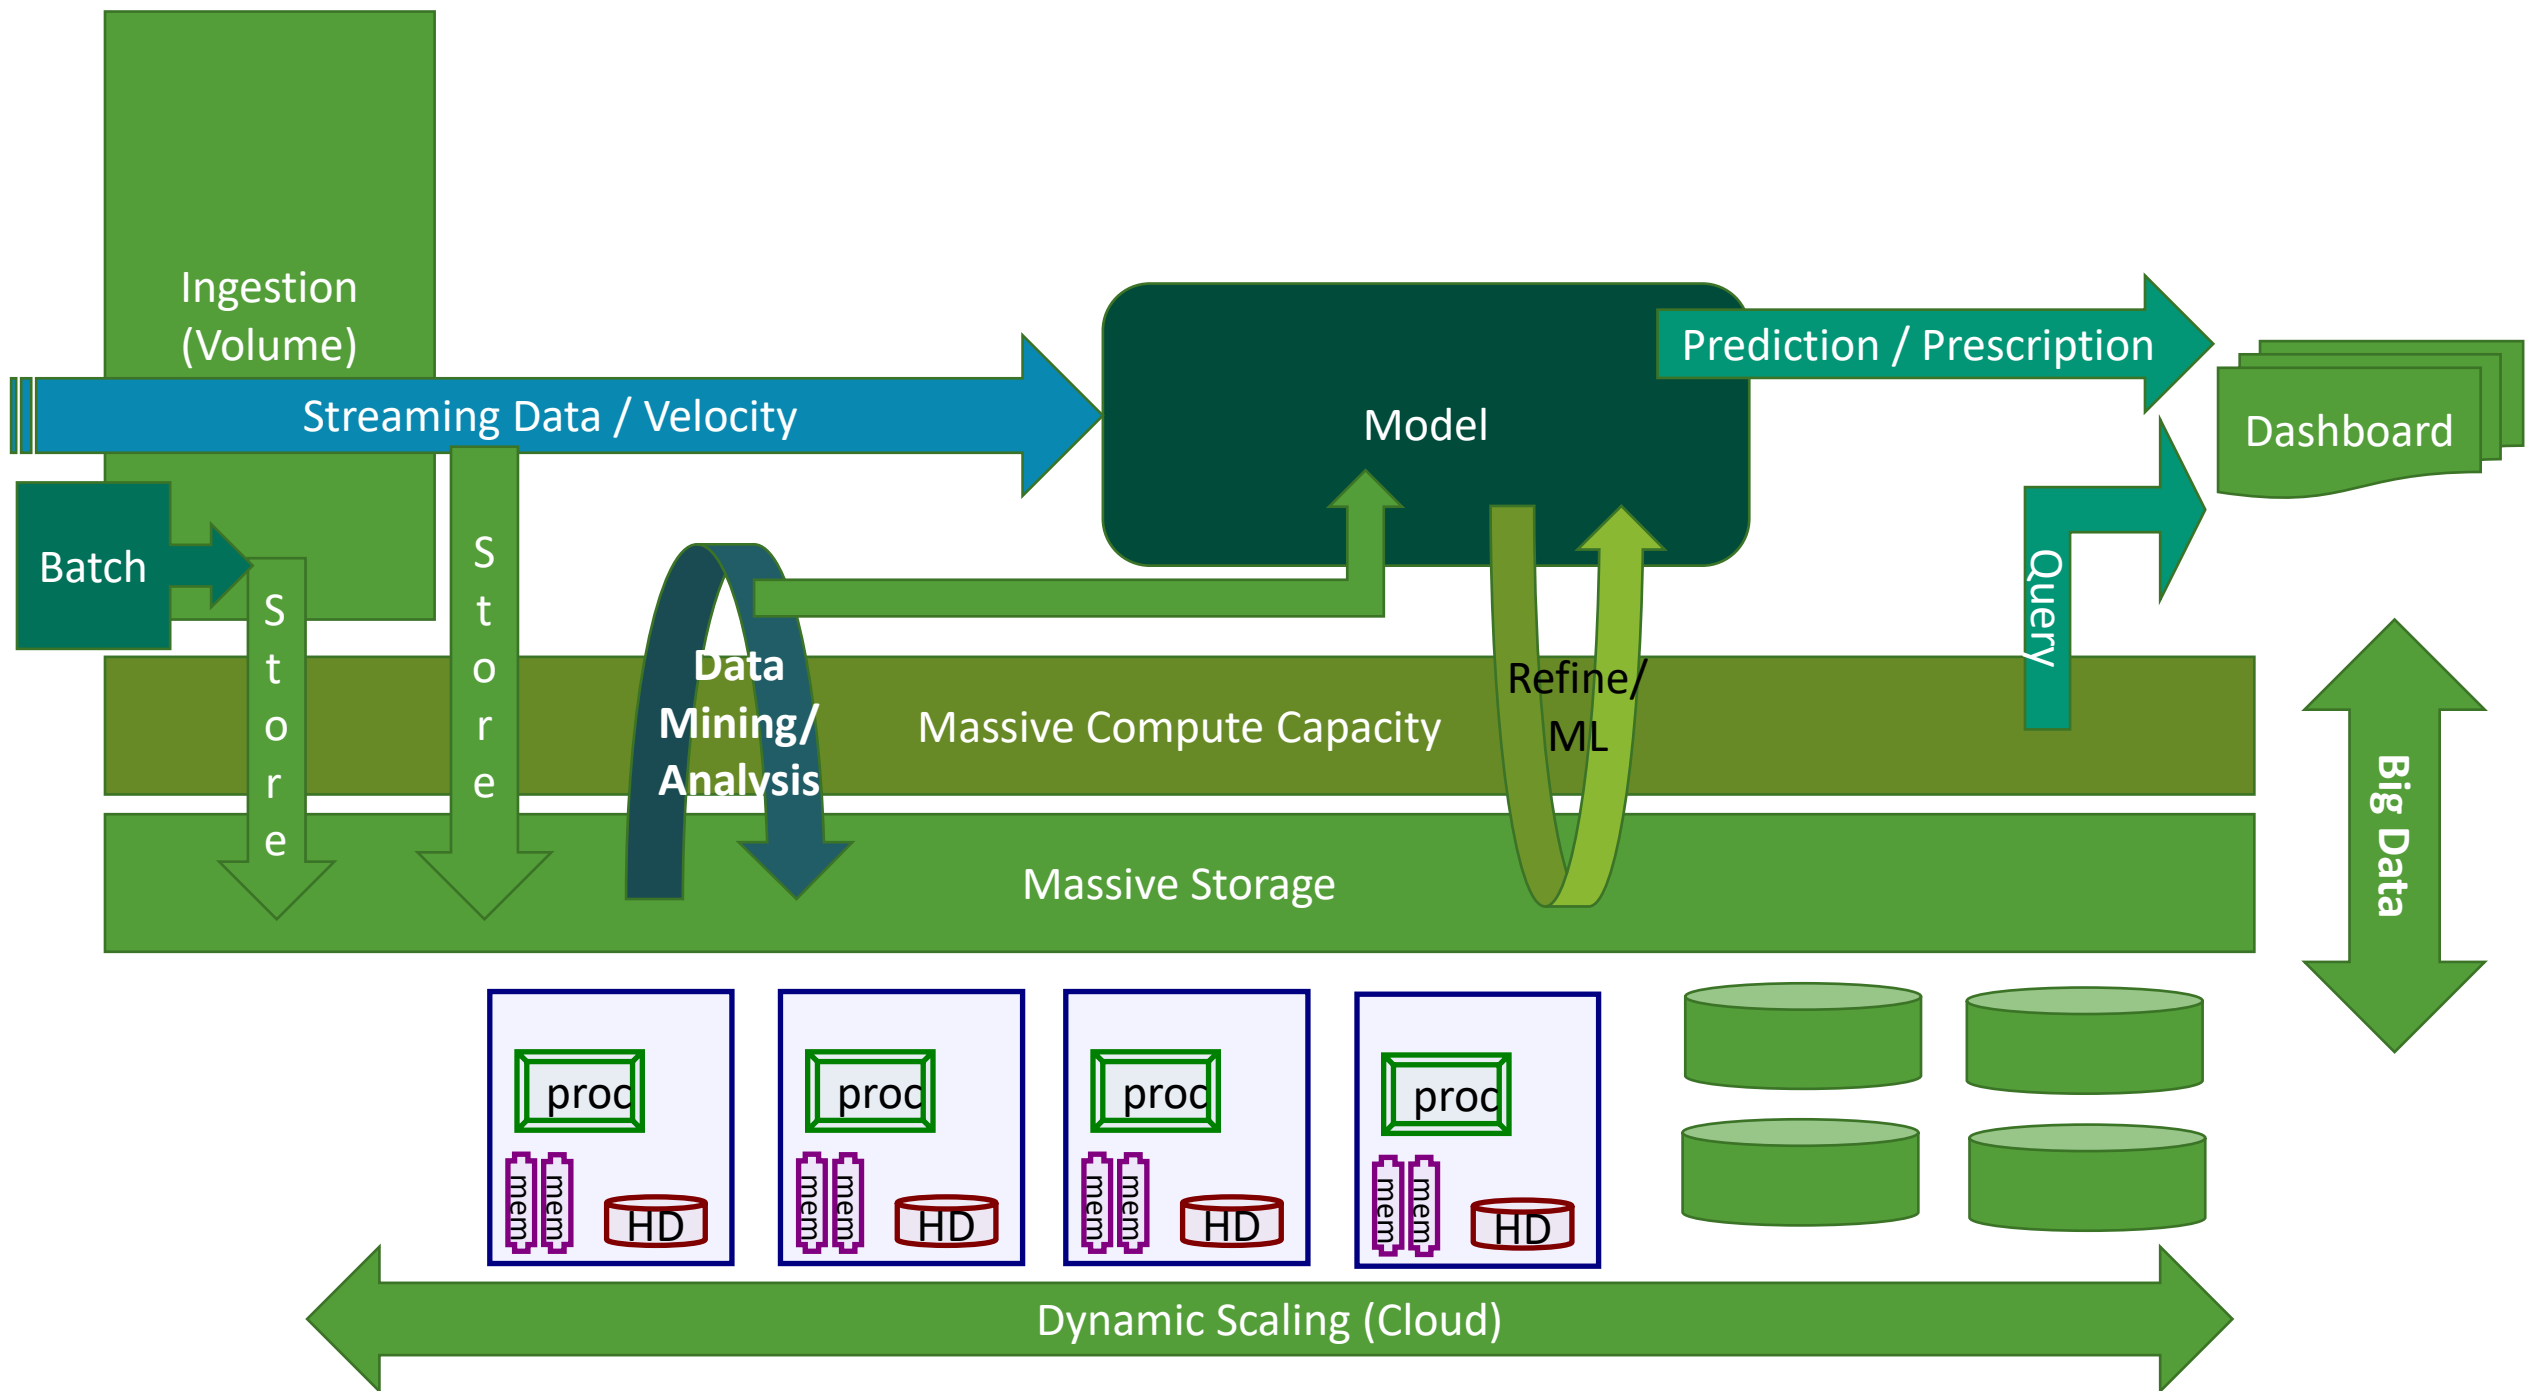
Traditional way



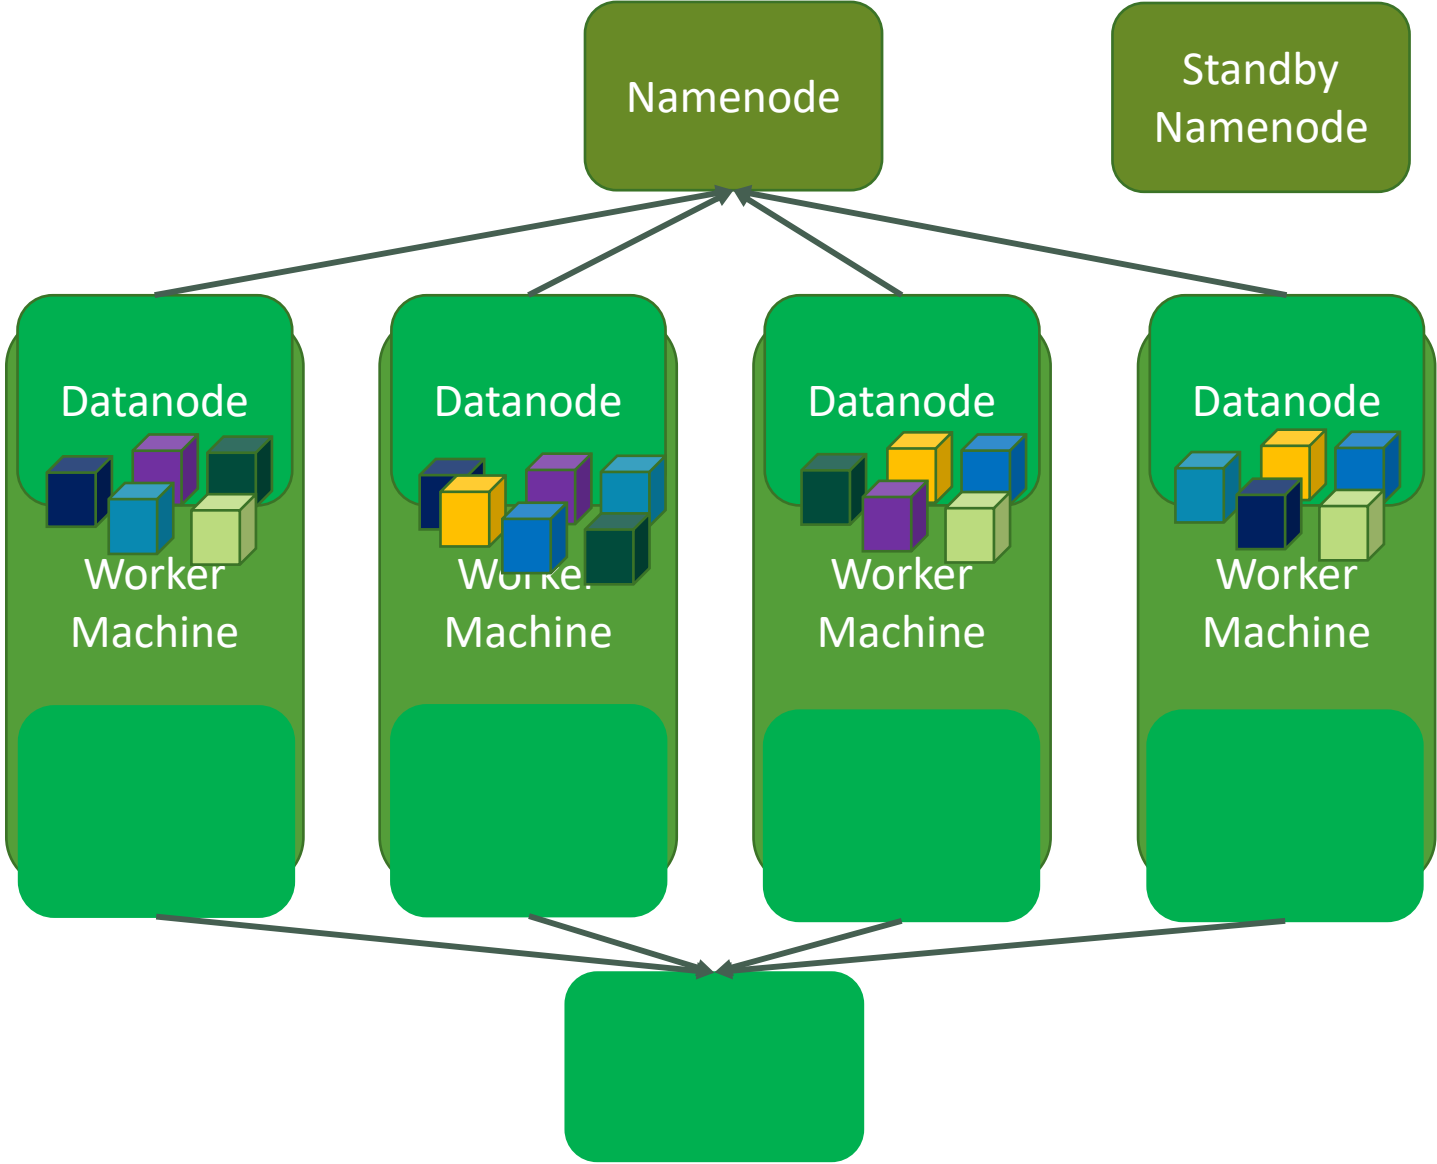
Machine Learning way









HDFS



YARN



HDFS and YARN


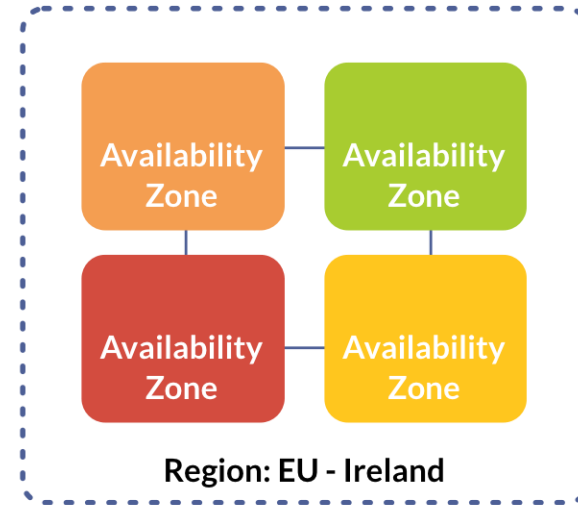
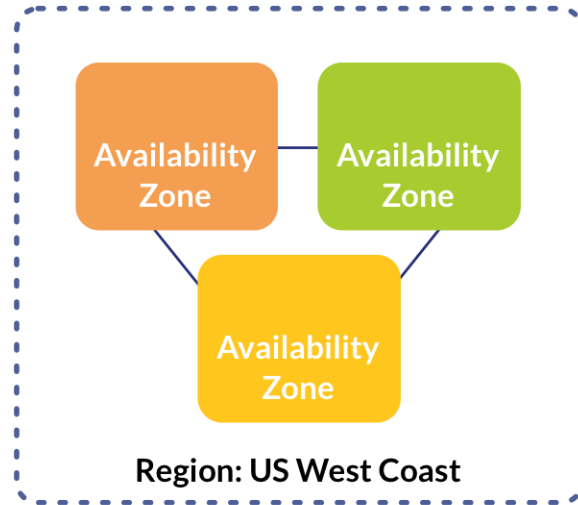


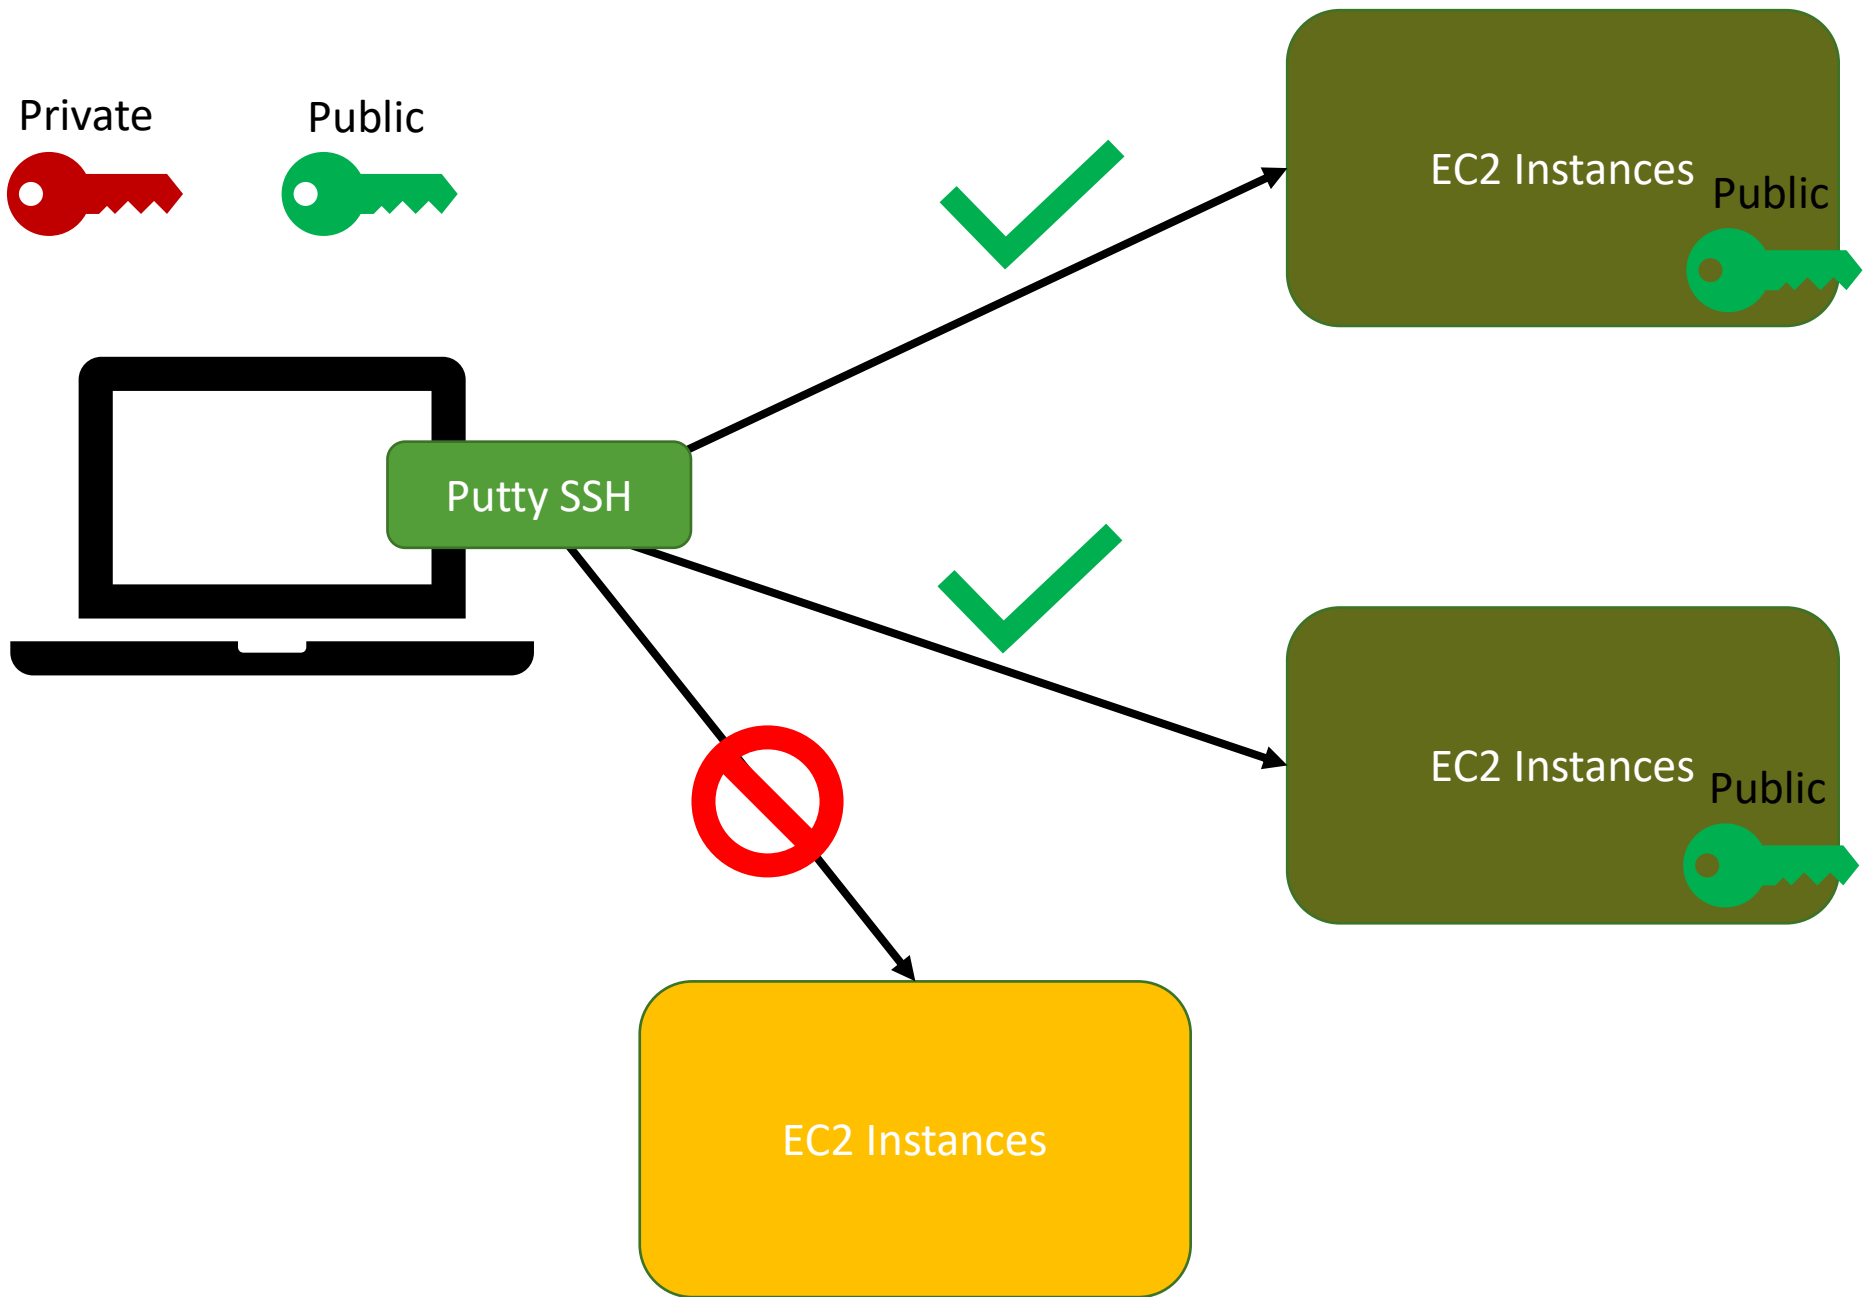


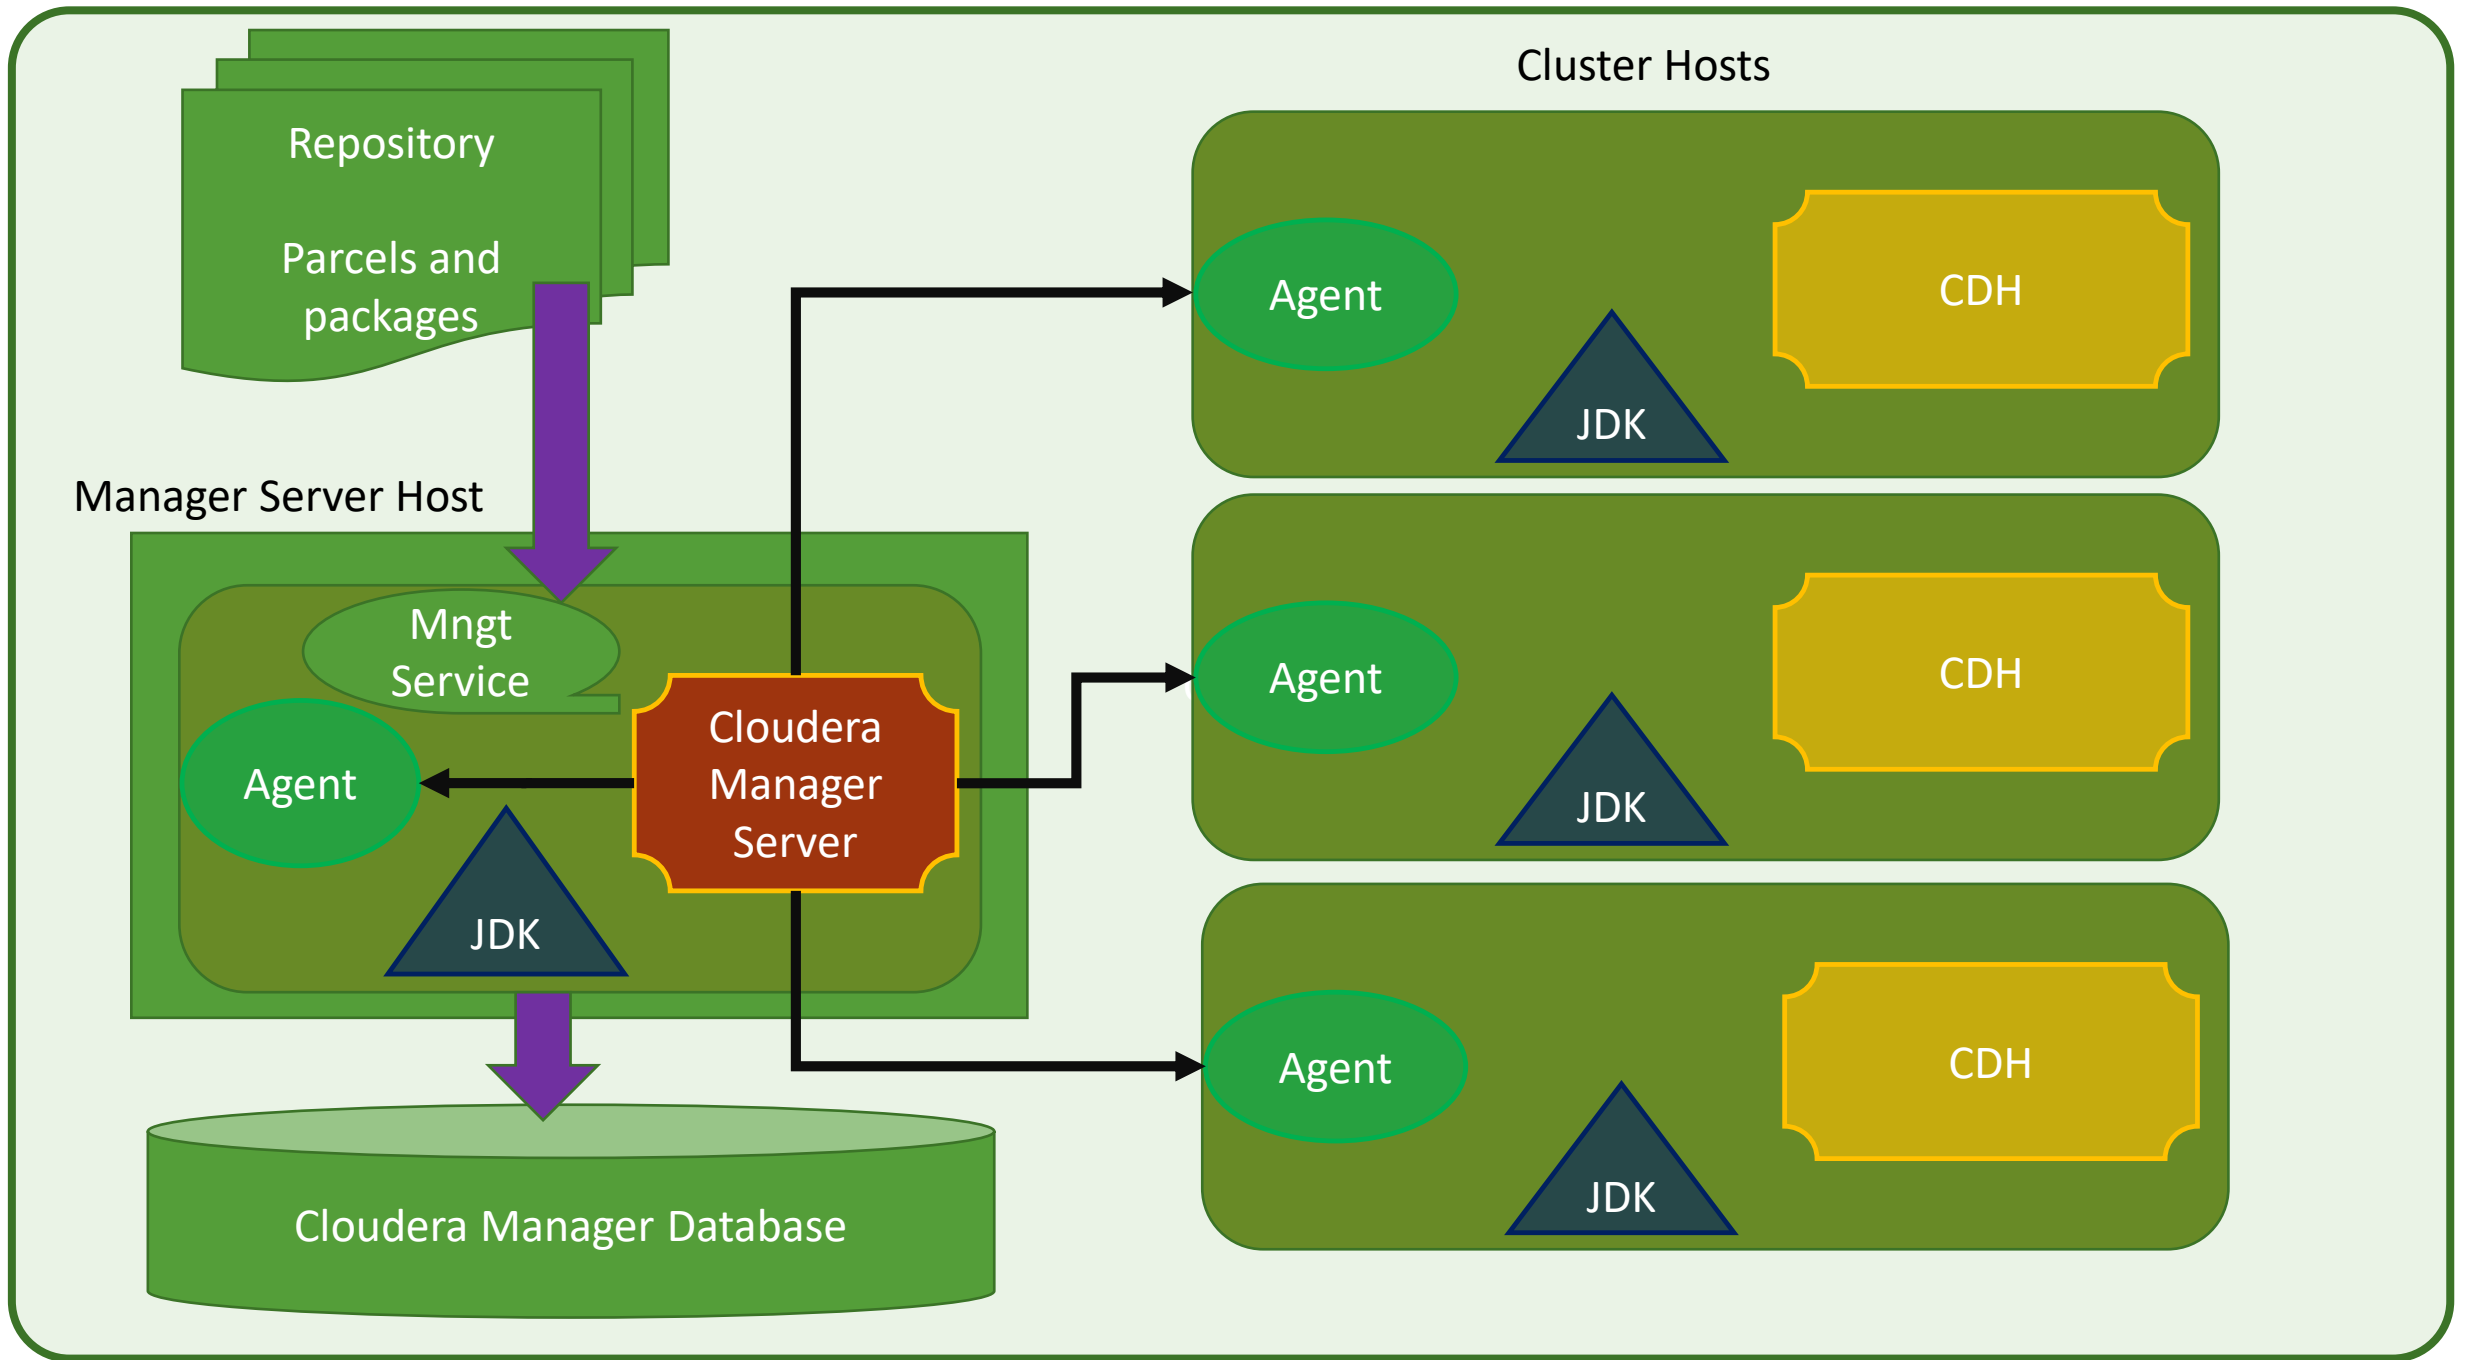
Hadoop EcoSystem							
Data Visualization							
	SAS Visual Analytics	Tableau	Qlick	SAP Lumira	R		
	D3.js	iCharts	Timeline JS	Apache Zeppelin	Pentaho		
System Deployment							
	Apache Ambari	Apache Mesos	Marathon	Apache Bigtop	HortonWorks HOYA	Deploop	
	Apache Eagle	Cloudera Manager	Brooklyn	Apache Helix	Buildoop	SequenceIQ Cloudbreak	
	Myriad						
Data Ingestion	Service Programming		Scheduling & DR	Security	Frameworks	Metadata	Machine Learning
Apache Flume	Apache Thrift	Apache Karaf	Apache Oozie	Apache Sentry	Jumbune	Metascope	Apache Mahout
Apache Sqoop	Apache Zookeeper	Twitter Elephant Bin	Linkedin Azkaban	Apache Knox Gatew	Spring XD	Apache Tika	WEKA
Facebook Scribe	Apache Avro	Linkedin Norbert	Apache Falcon	Apache Ranger	Cask Data App Platform		Cloudera Oryx
Apache Chukwa	Apache Curator		Shedoscope				Deeplearning4j
Apache Kafka							MADlib
Netflix Suro	Distributed Programming				SQL on Hadoop		H2O
Apache Samza	Apache Ignite	Apache Flink	Apache Twill	Kangaroo	Apache Hive	Facebook Presto	Sparkling Water
Cloudera Morphline	Apache Mapreduce	Apache Apex	Damballa PARKOUR	TinkerPop	Apache Hcatalog	Datasalt Splout SQL	Apache SystemML
HIHO	Apache Pig	Netflix Pigpen	Apache Hama	Pachyderm Mapred	Apache Trafodion	Apache Tajo	
Apache Nifi	JAQL	AMPLAB SIMR	Datasalt Pangool	Apache Beam	Apache HAWQ	Apache MRQL	
Apache ManifoldCF	Apache Spark	Facebook Corona	Apache Tez		Apache Drill	Kylin	
	Apache Storm	Apache REEF	Apache DataFu		Cloudera Impala		
	NoSQL Databases					NewSQL Databases	Deep Learning
	Wide Column	Document	Key-Value	Graph	Stream Data Model	TokuDB	Caffe
	Apache Hbase	MongoDB	Redis	Giraph	EventStore	HandlerSocket	Microsoft CNTK
	Apache Cassandra	RethinkDB	Linkedin Voldemort	Neo4J		Akiban Server	Tensor Flow
	Hypertable	ArangoDB	RocksDB	TitanDB		Drizzle	Theano
	Apache Accumulo	CouchDB	OpenTSDB	OrientDB		Haeinsa	Torch
	Apache Kudu	DynamoDB				SenseiDB	Mxnet
	Apache Parquet	Gemfire				Sky	Chainer
						BayesDB	Keras
						InfluxDB	
						VoldDB	
						SAP HANA	
Distributed File System							
	Apache HDFS	Quantcast File Syste	Lustre File System	GridGain			
	Red Hat GlusterFS	Ceph File System	Alluxio	XtreemFS			

AWS Regions and AZs

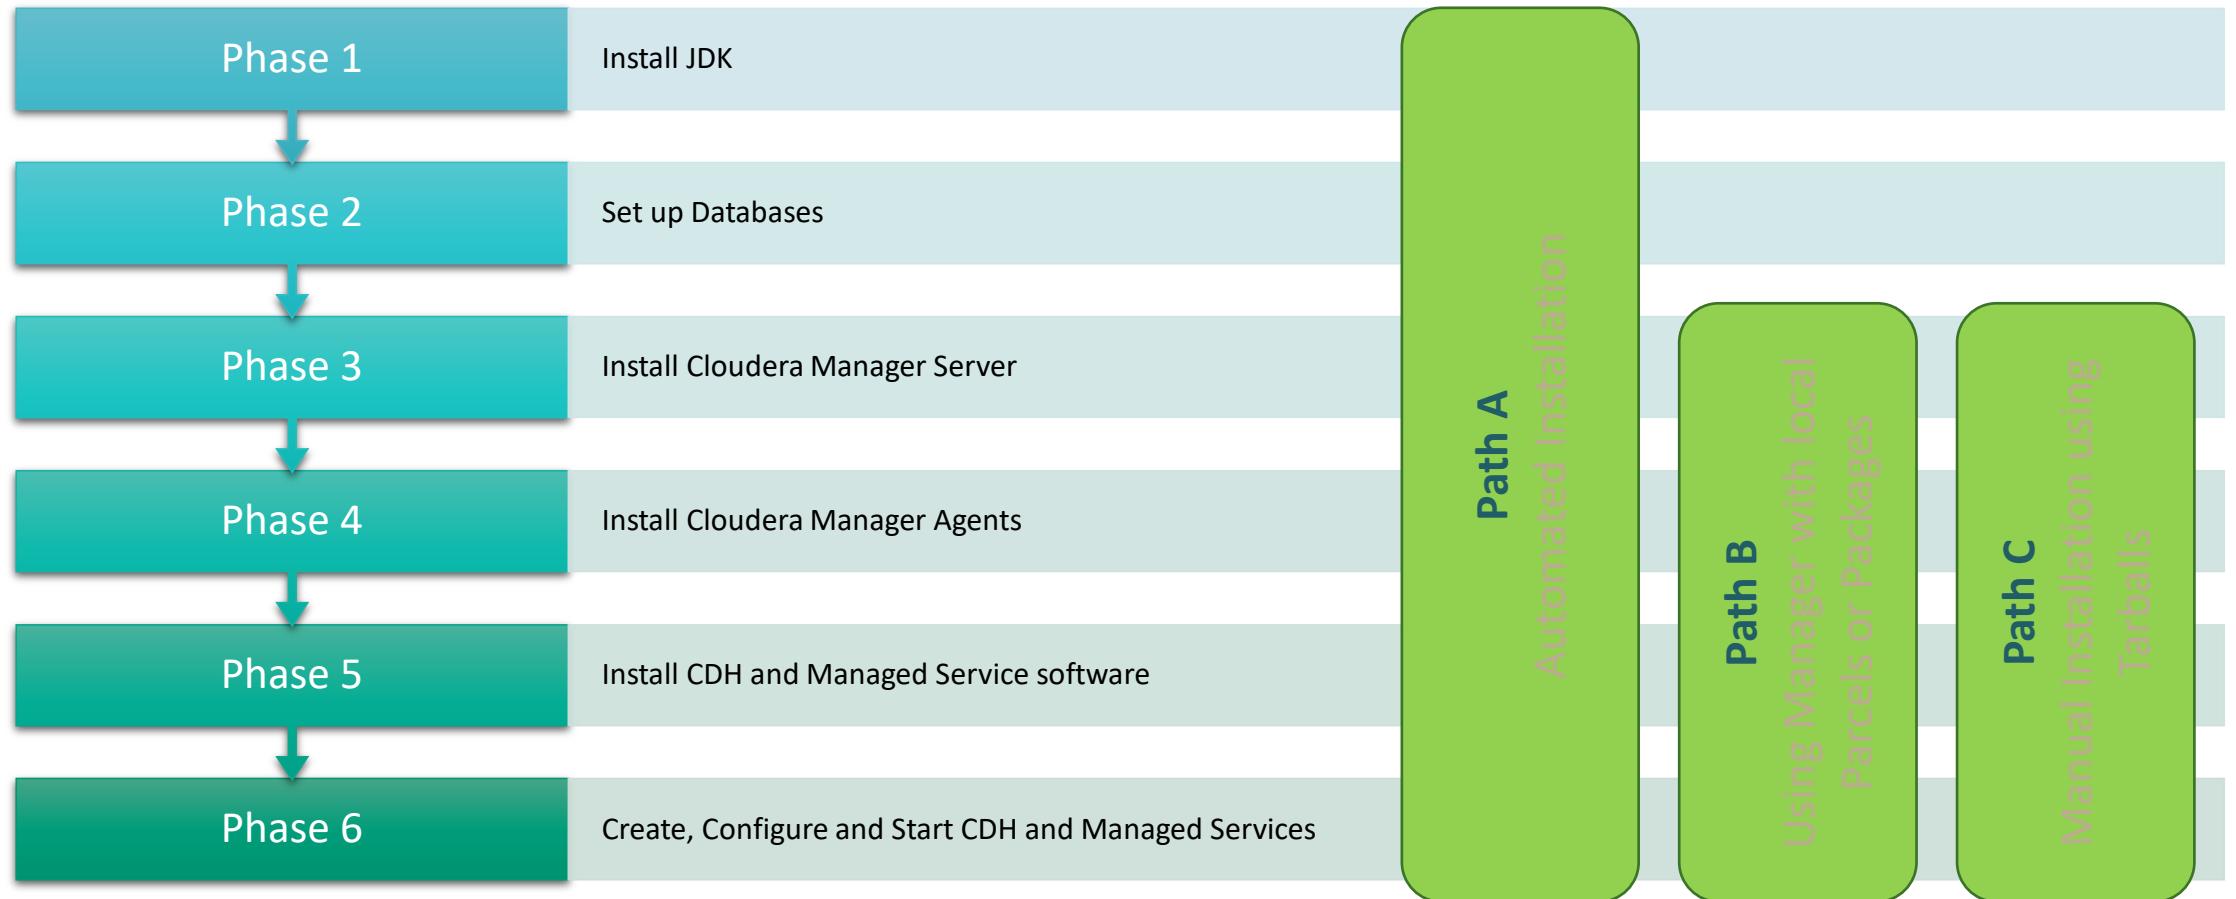








Cloudera Installation Phases and Paths

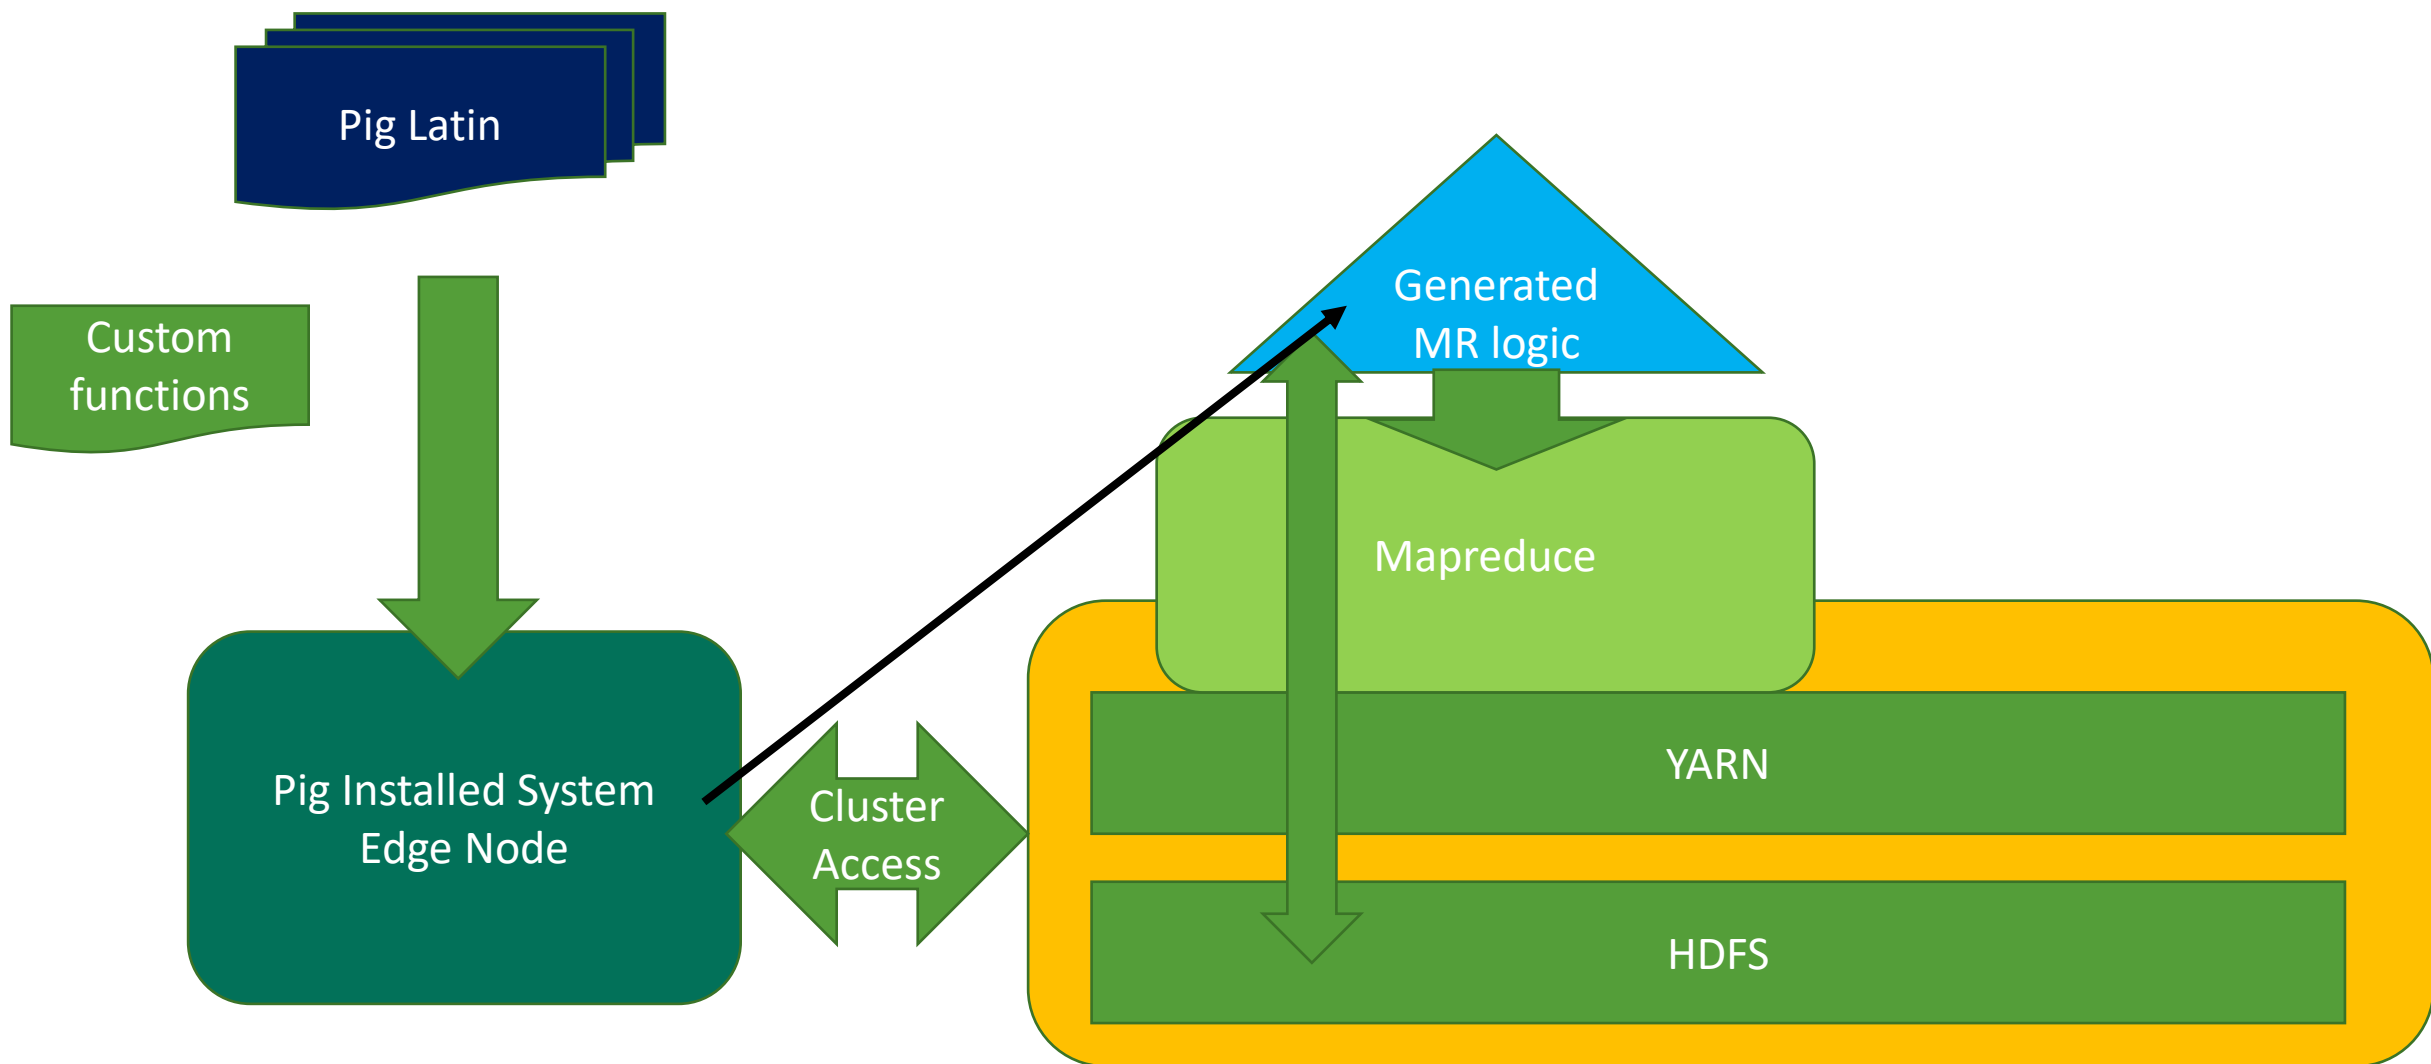



Advantages of using Parcels

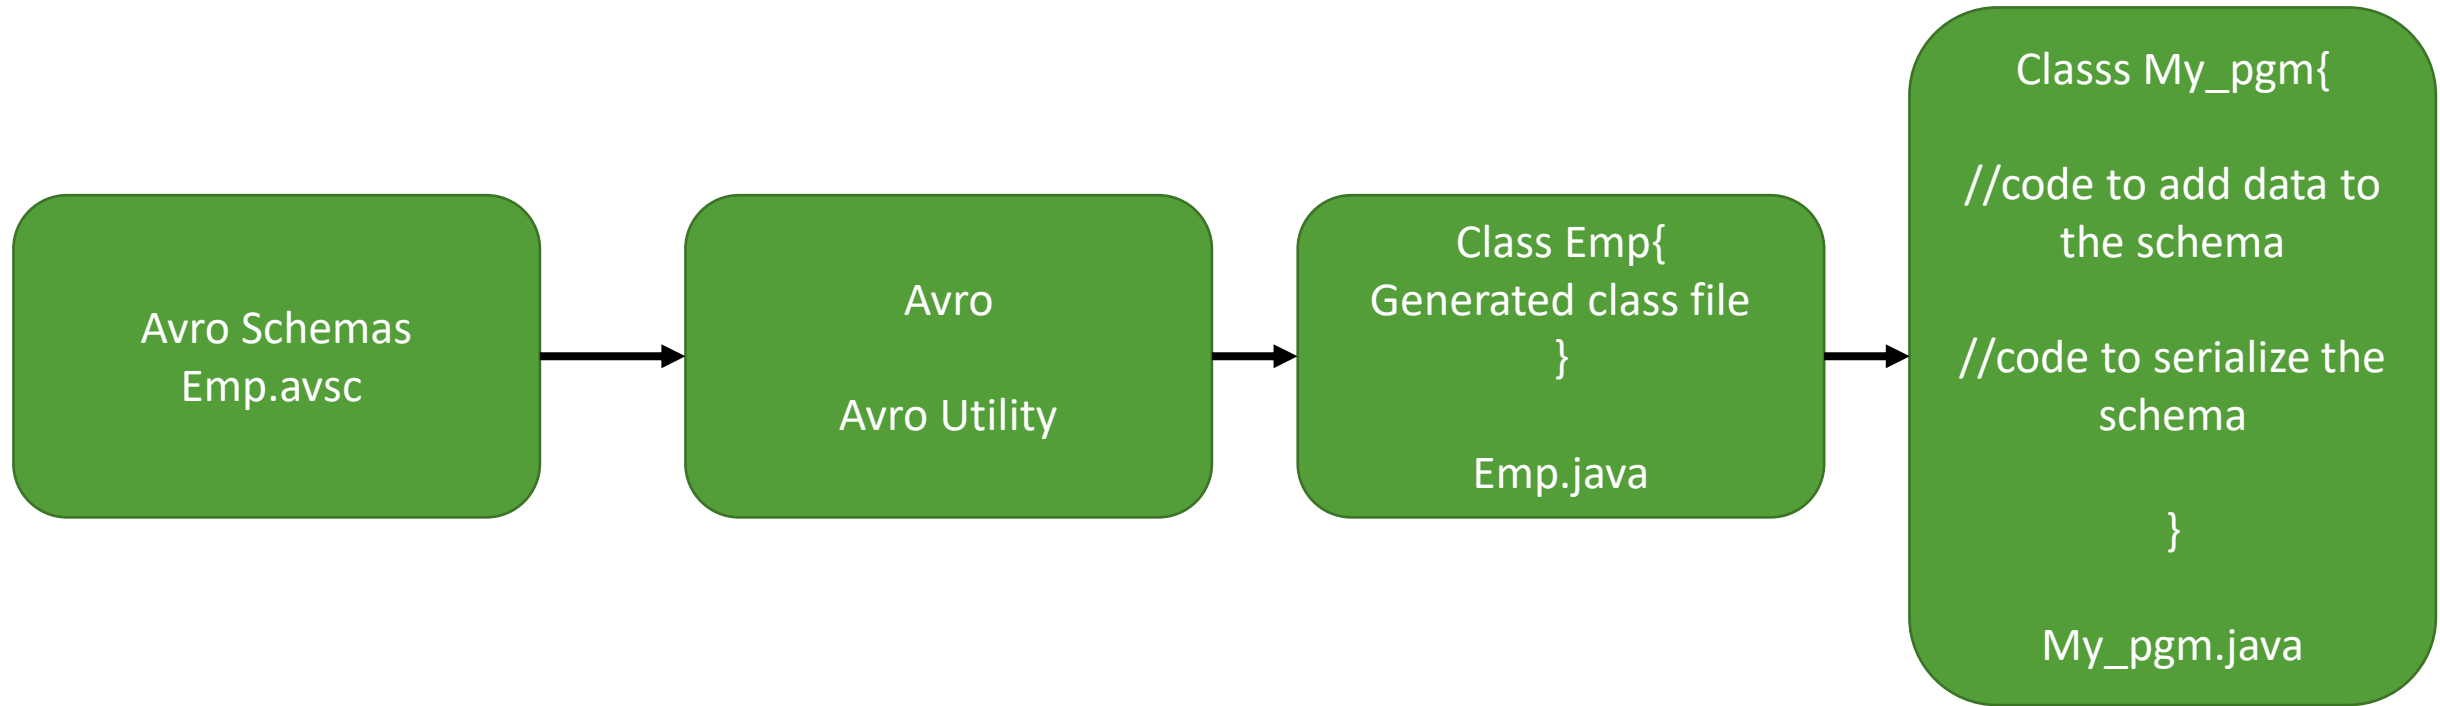
- Distribution of CDH as a single object - Instead of having a separate package for each part of CDH, **parcels have just a single object to install**
- Internal consistency - All CDH components are matched, **eliminating the possibility of installing parts from different versions of CDH**
- Installation **outside of /usr**
- Installation of CDH without sudo - Parcel installation is handled by the Cloudera Manager Agent running as root or another user, so **you can install CDH without sudo.**

Types of Hosts and their roles


- **Master hosts** run Hadoop master processes such as the HDFS NameNode and YARN Resource Manager.
- **Worker hosts** primarily run DataNodes and other distributed processes such as Impalad, Datanode, Node Manager.
- **Utility hosts** run other cluster processes that are not master processes such as Cloudera Manager and the Hive Metastore.
- **Edge hosts** are client access points for launching jobs in the cluster. The number of Edge hosts required varies depending on the type and size of the workloads.








MySQL Database Installation and Overview

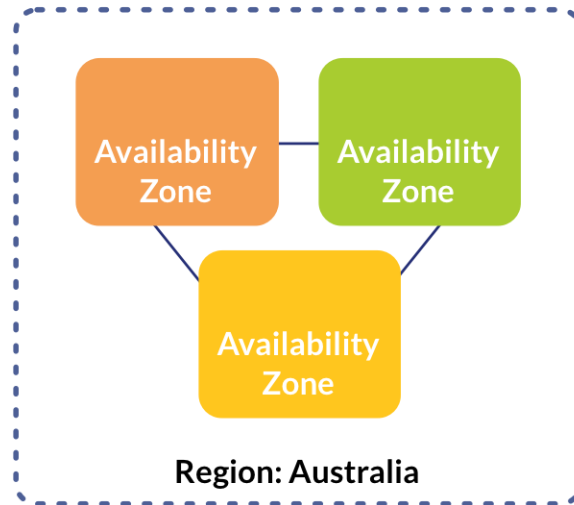





Repository Setup

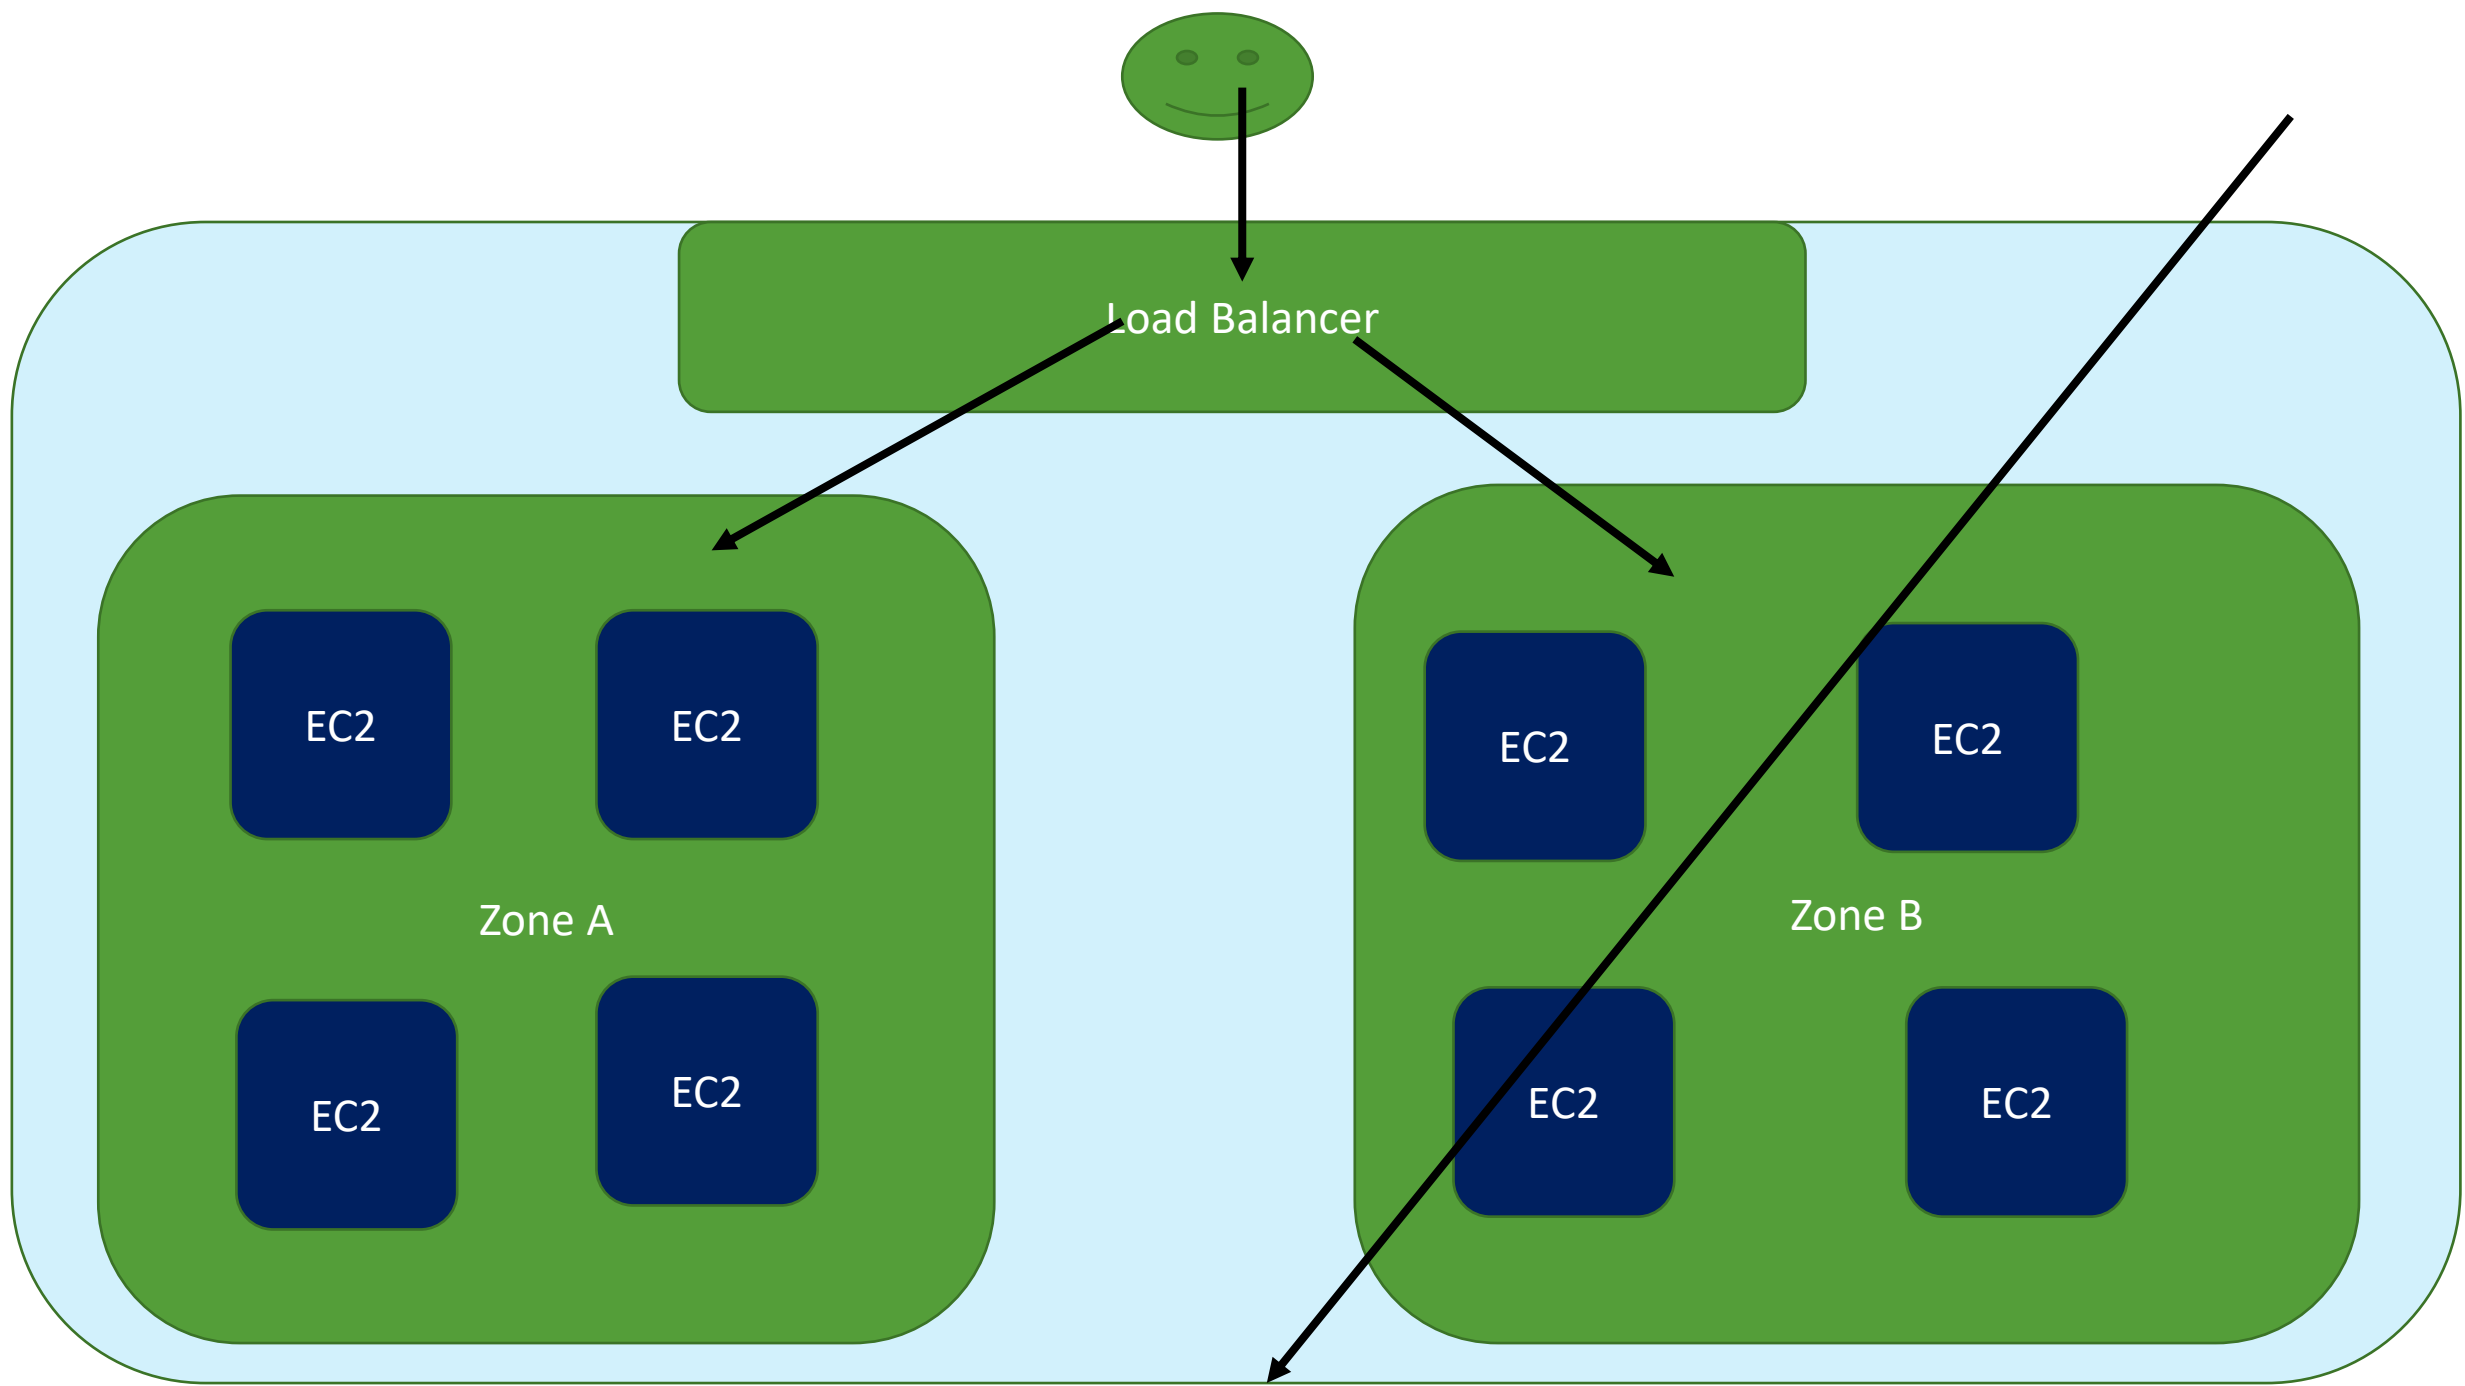


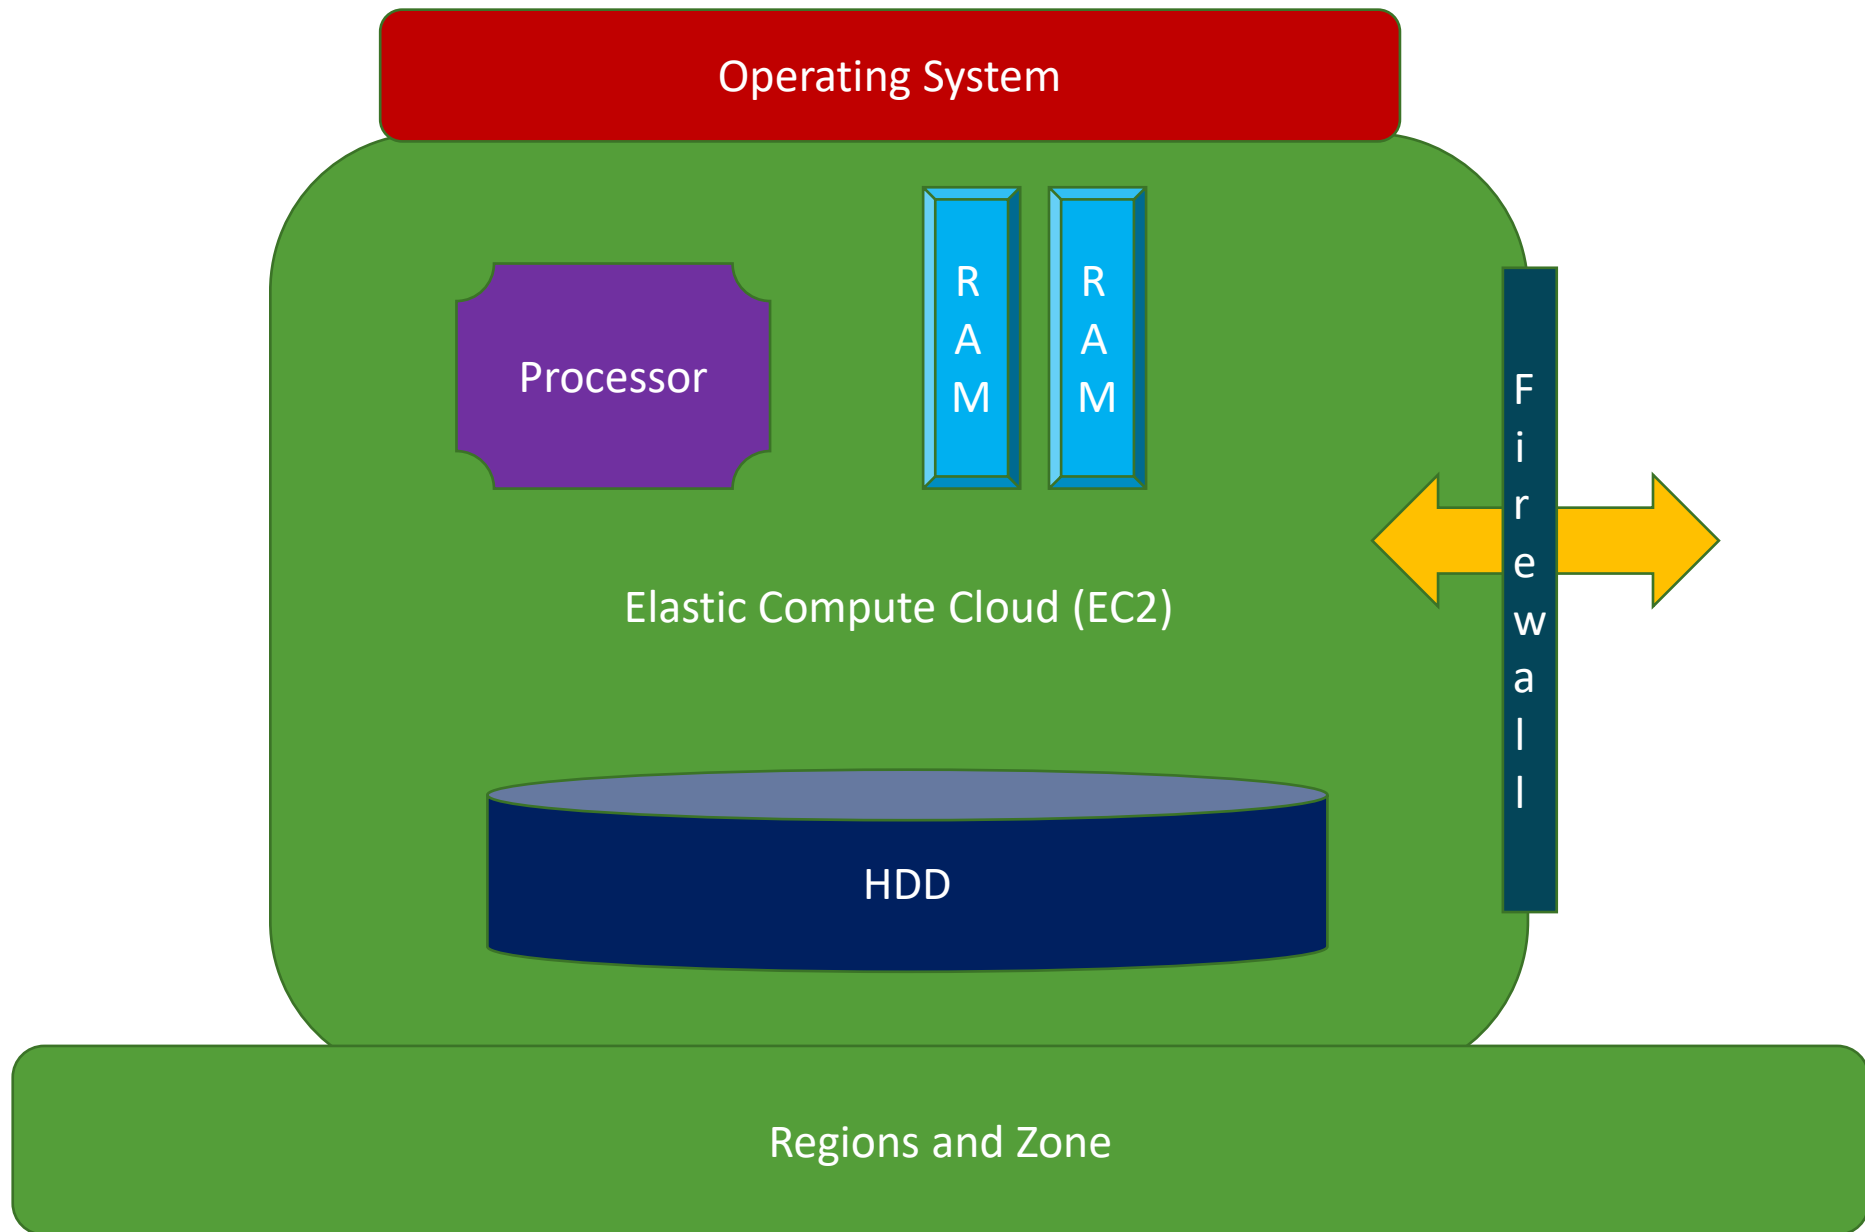
AWS Regions and AZs

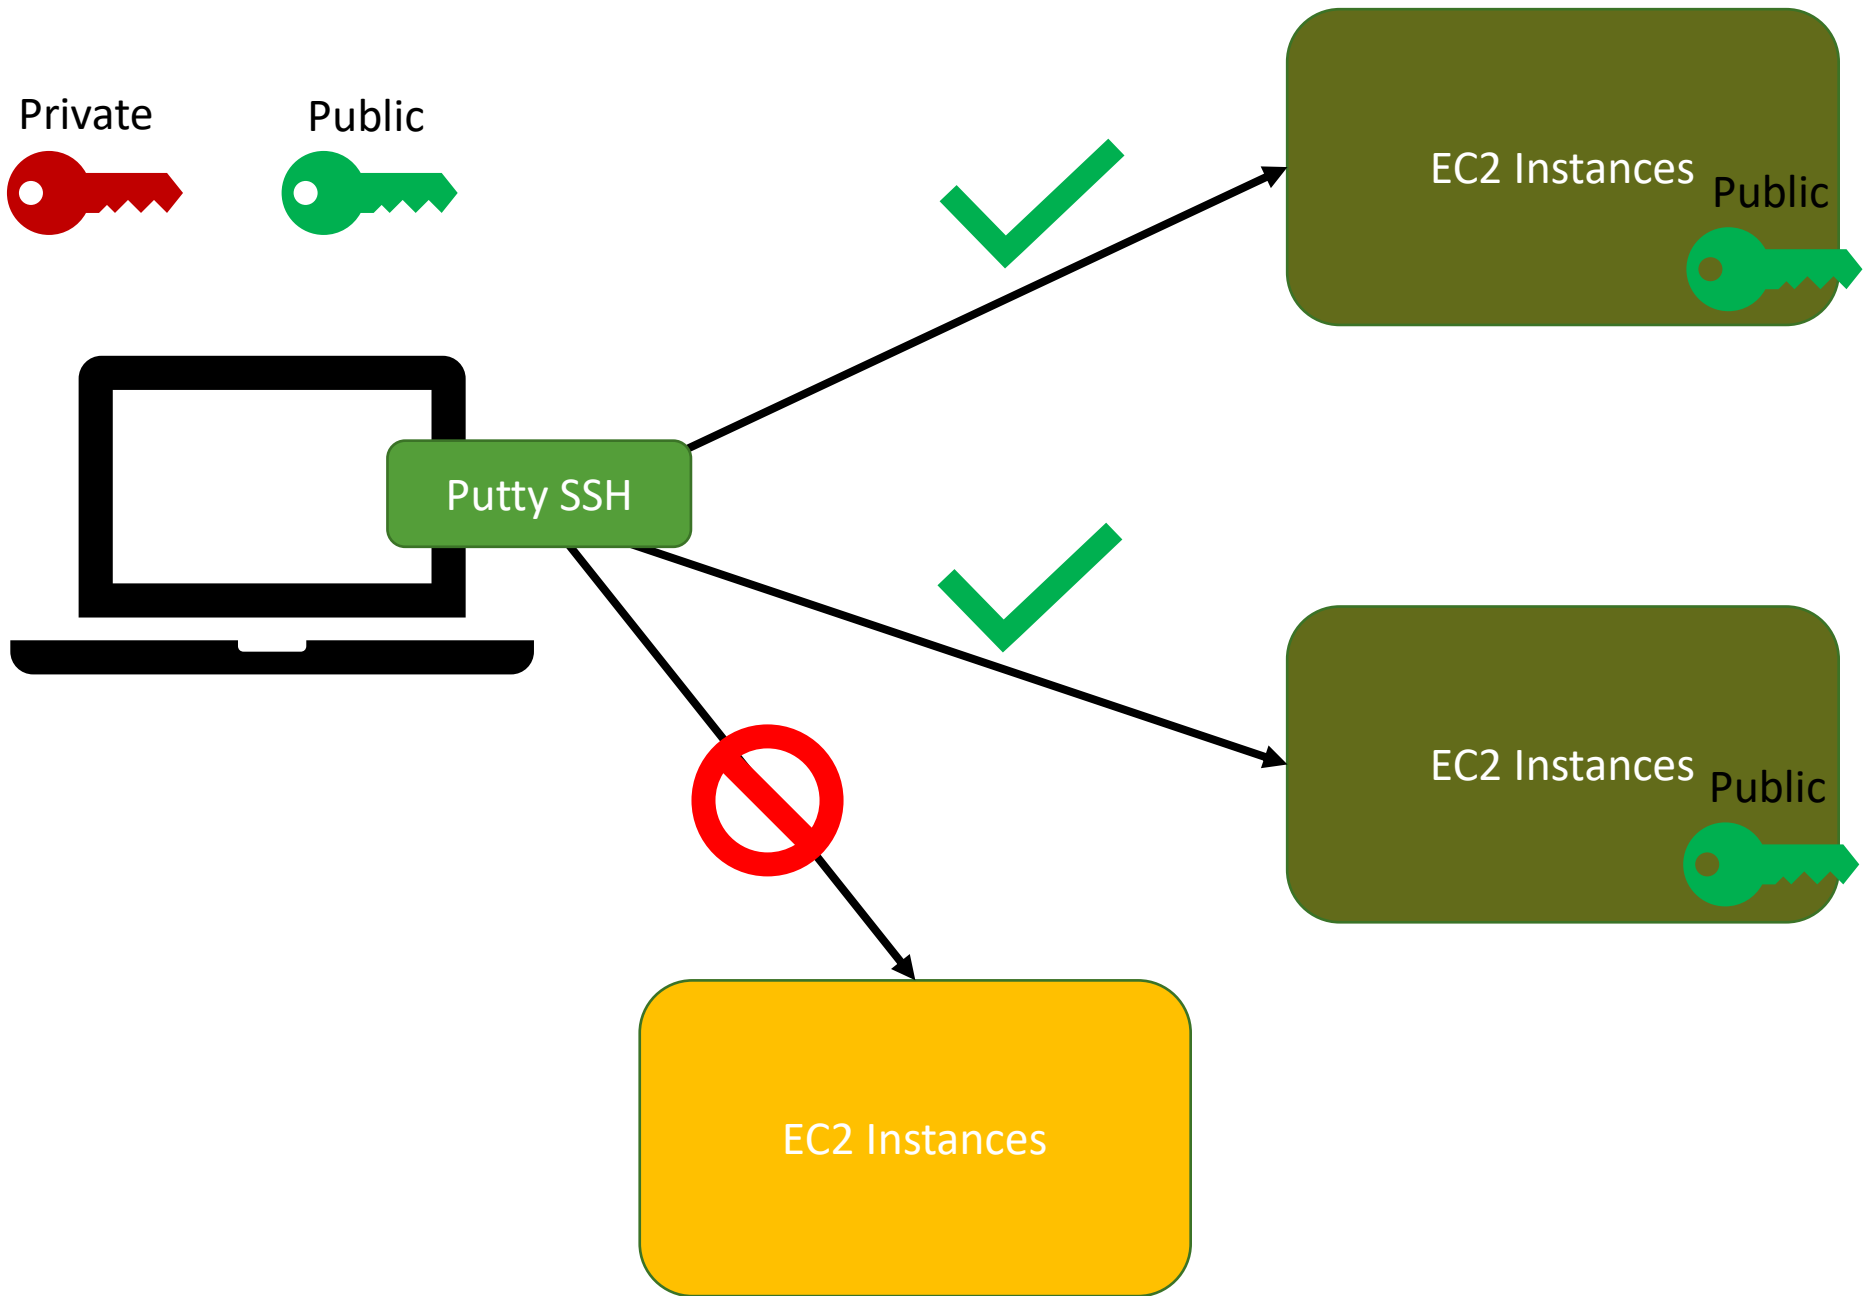


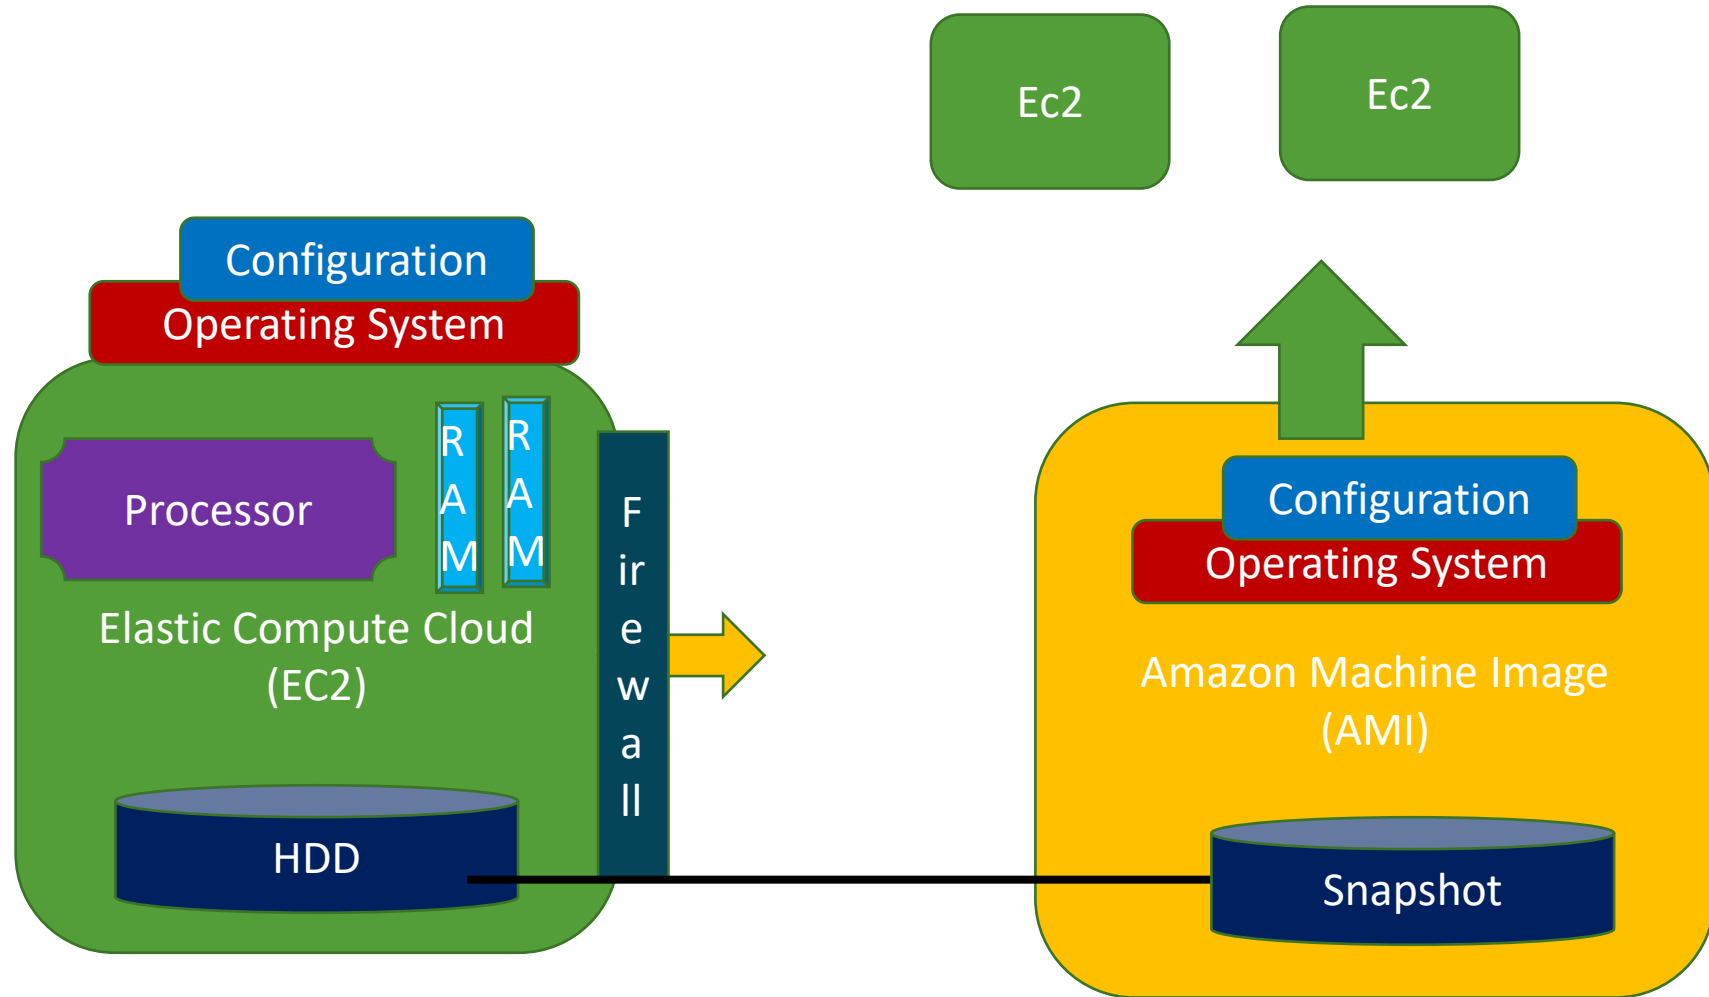













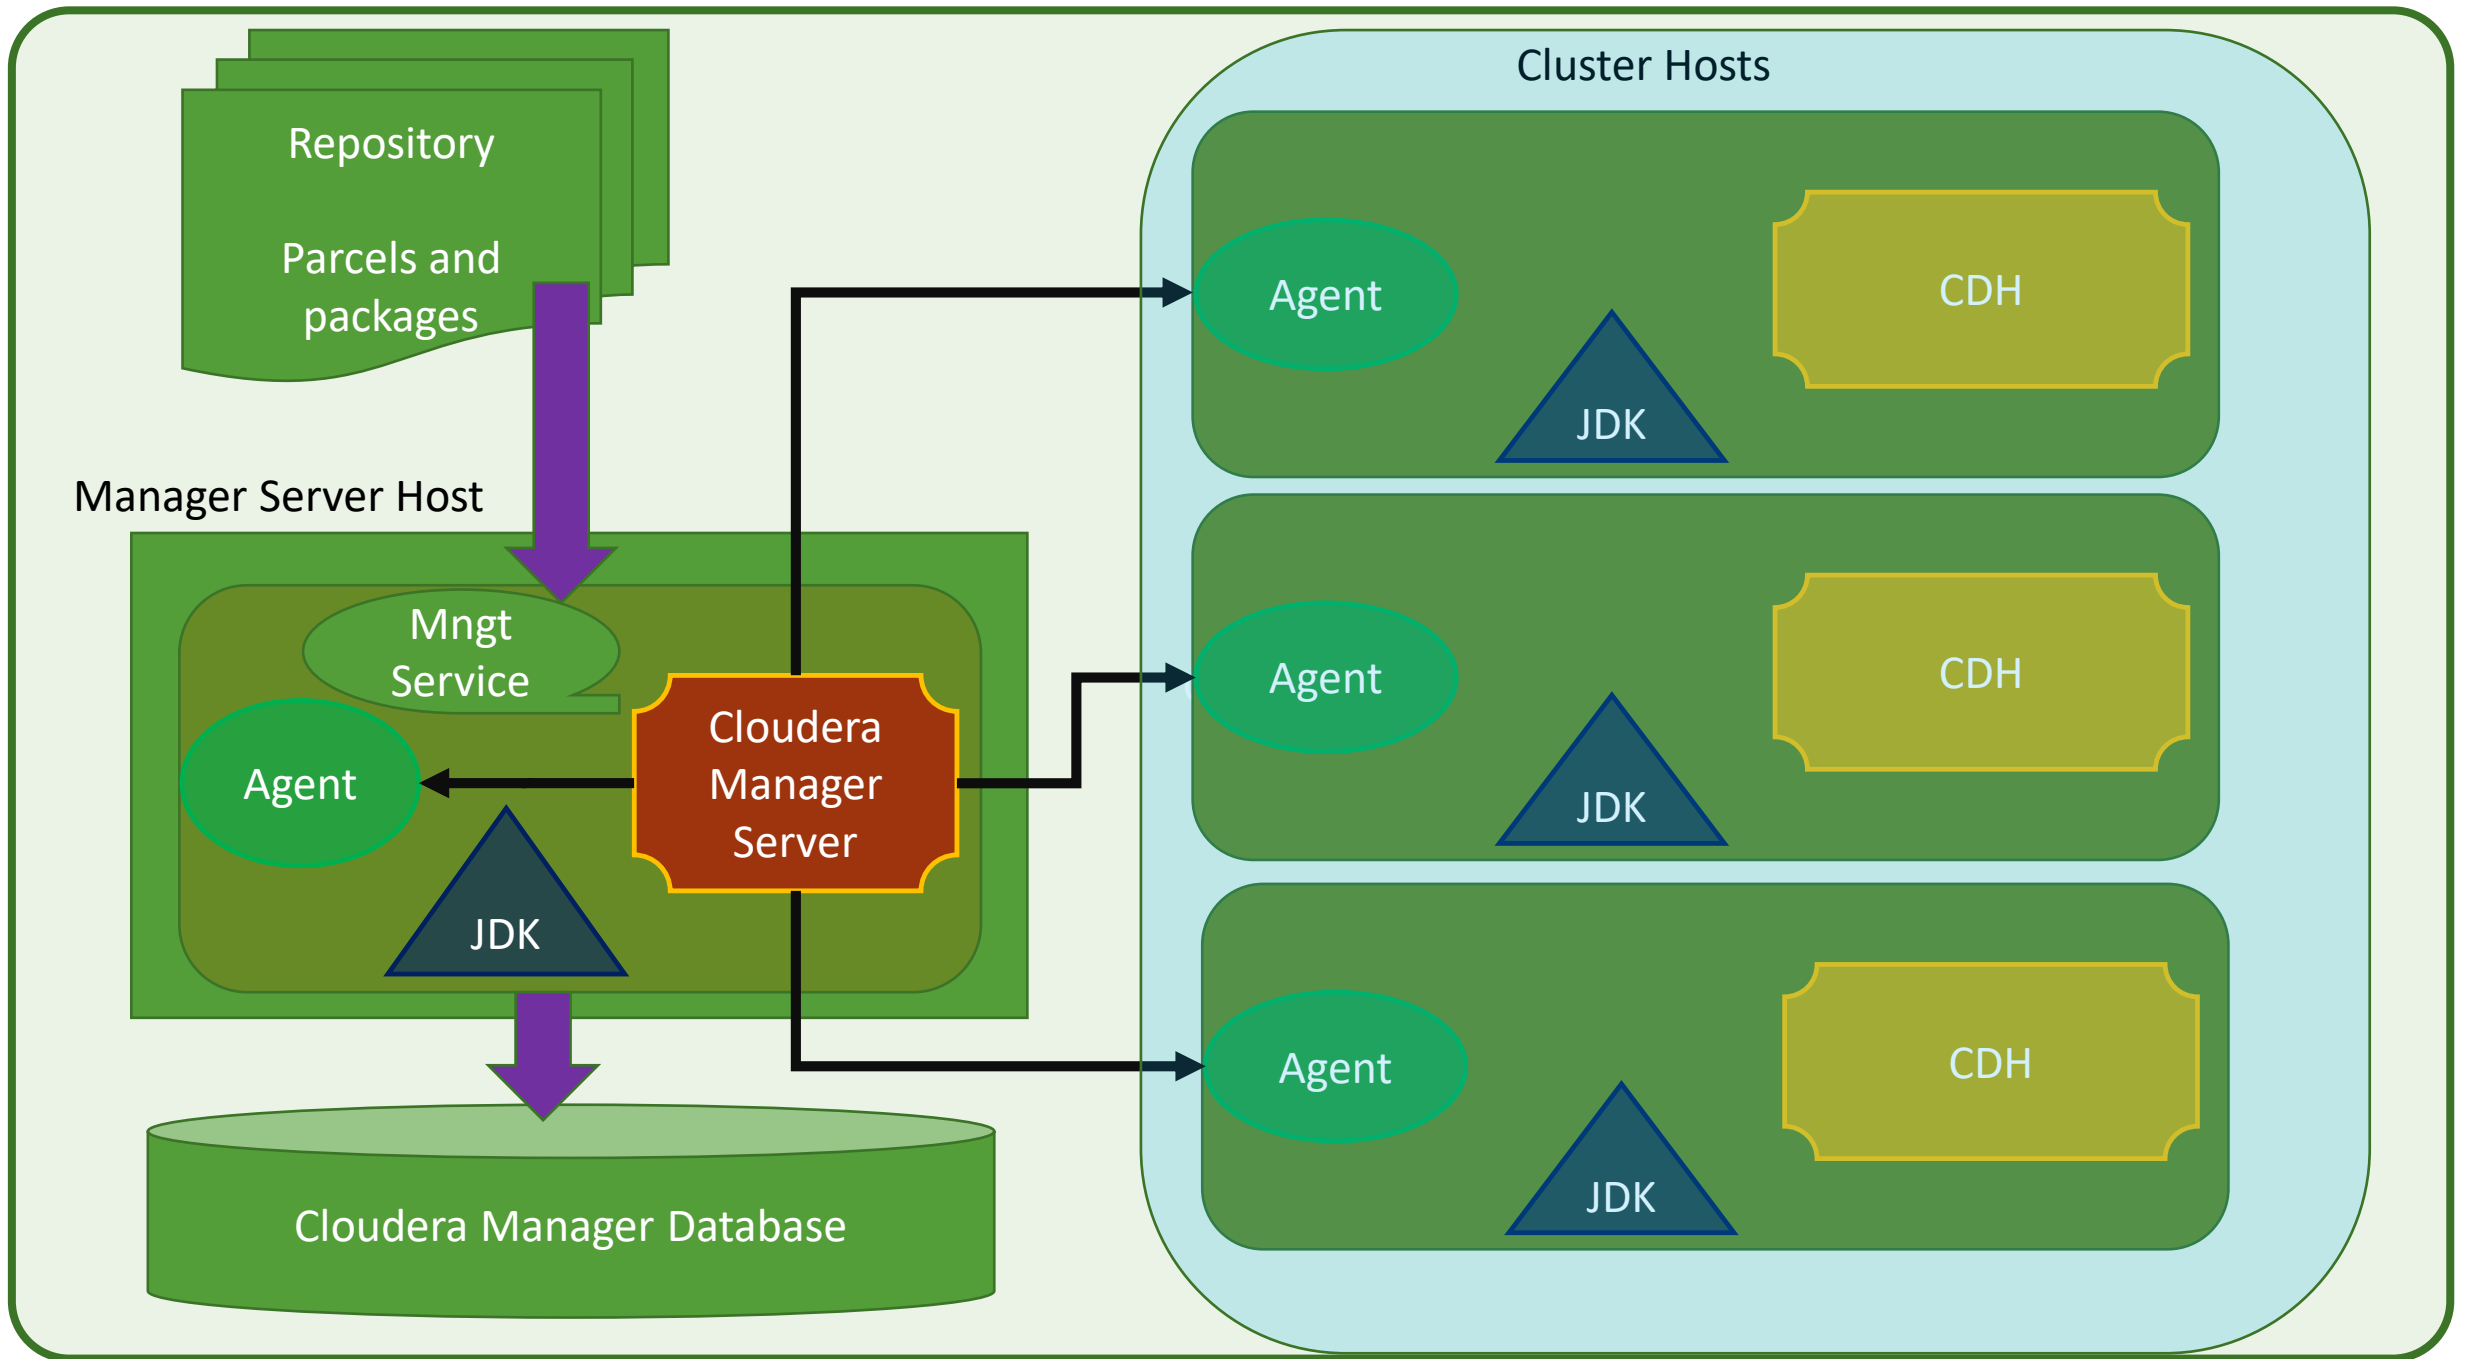




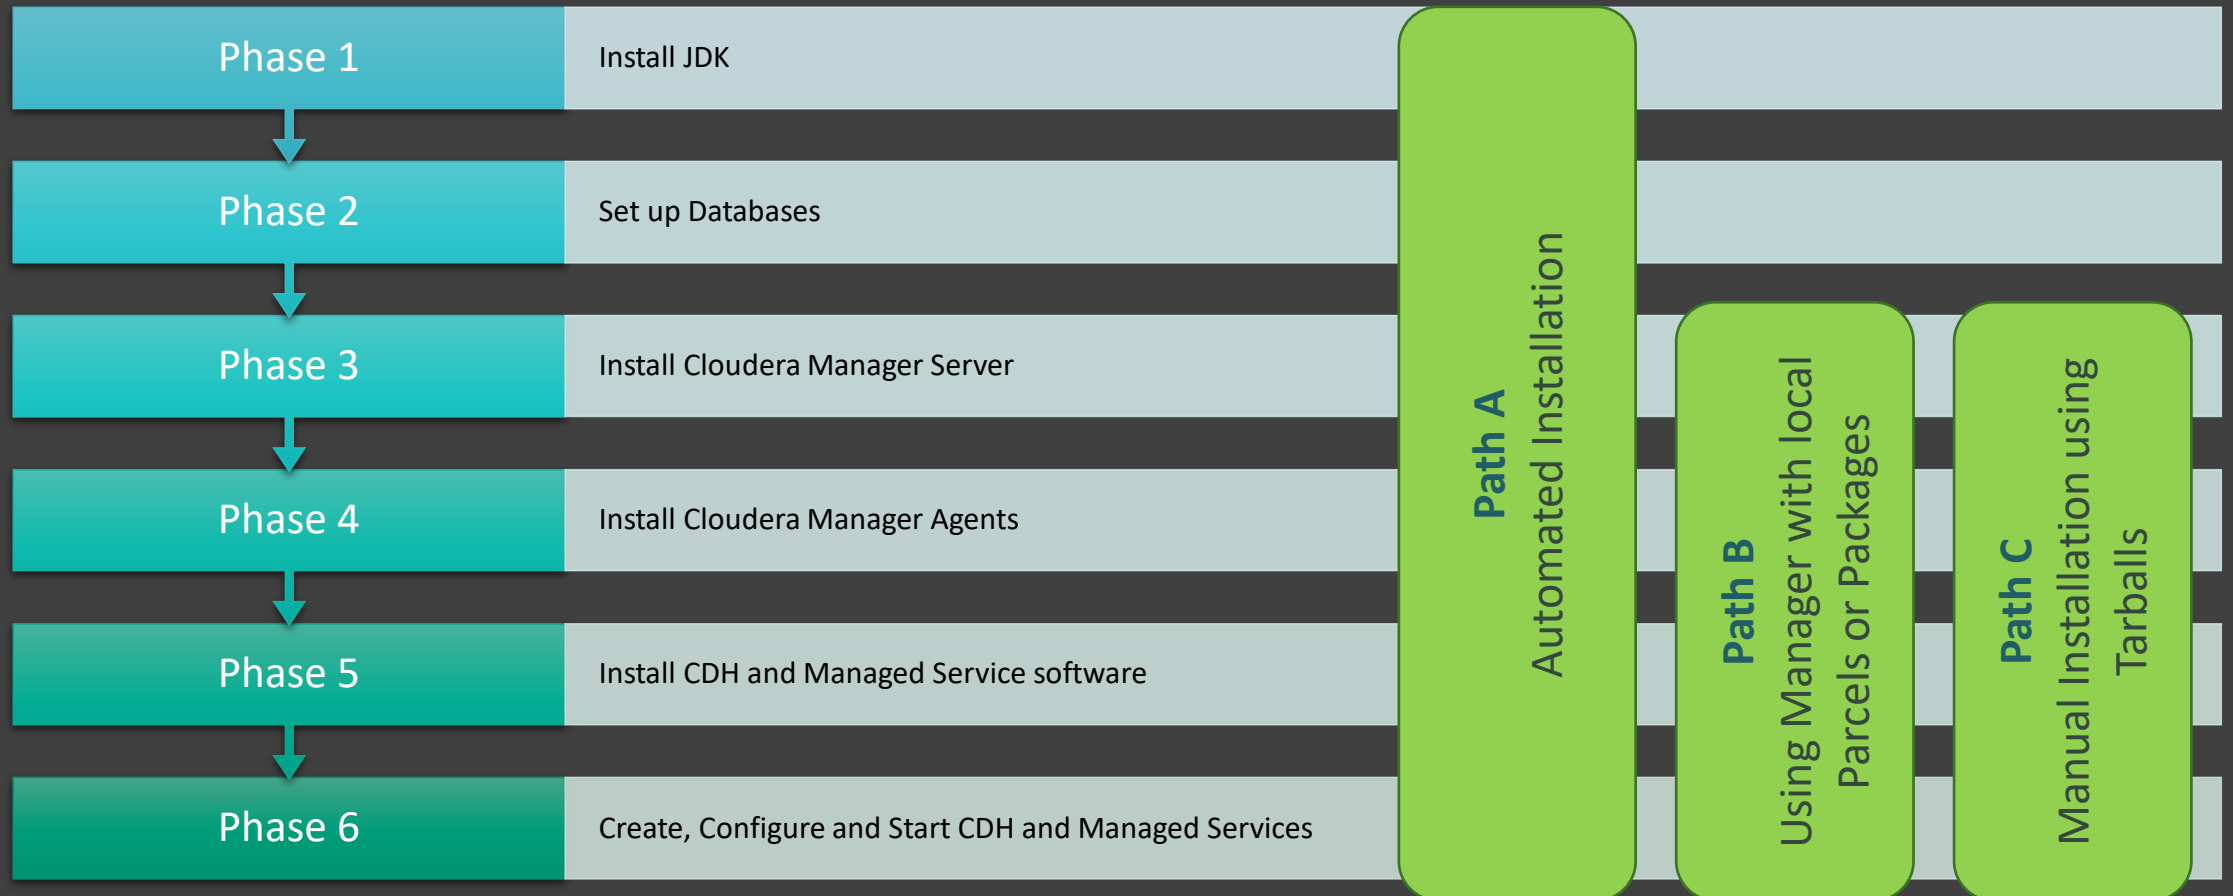




Disable SELinux
TurnOff Iptables
Resize Volume
Change swappiness



Cloudera Installation Phases and Paths




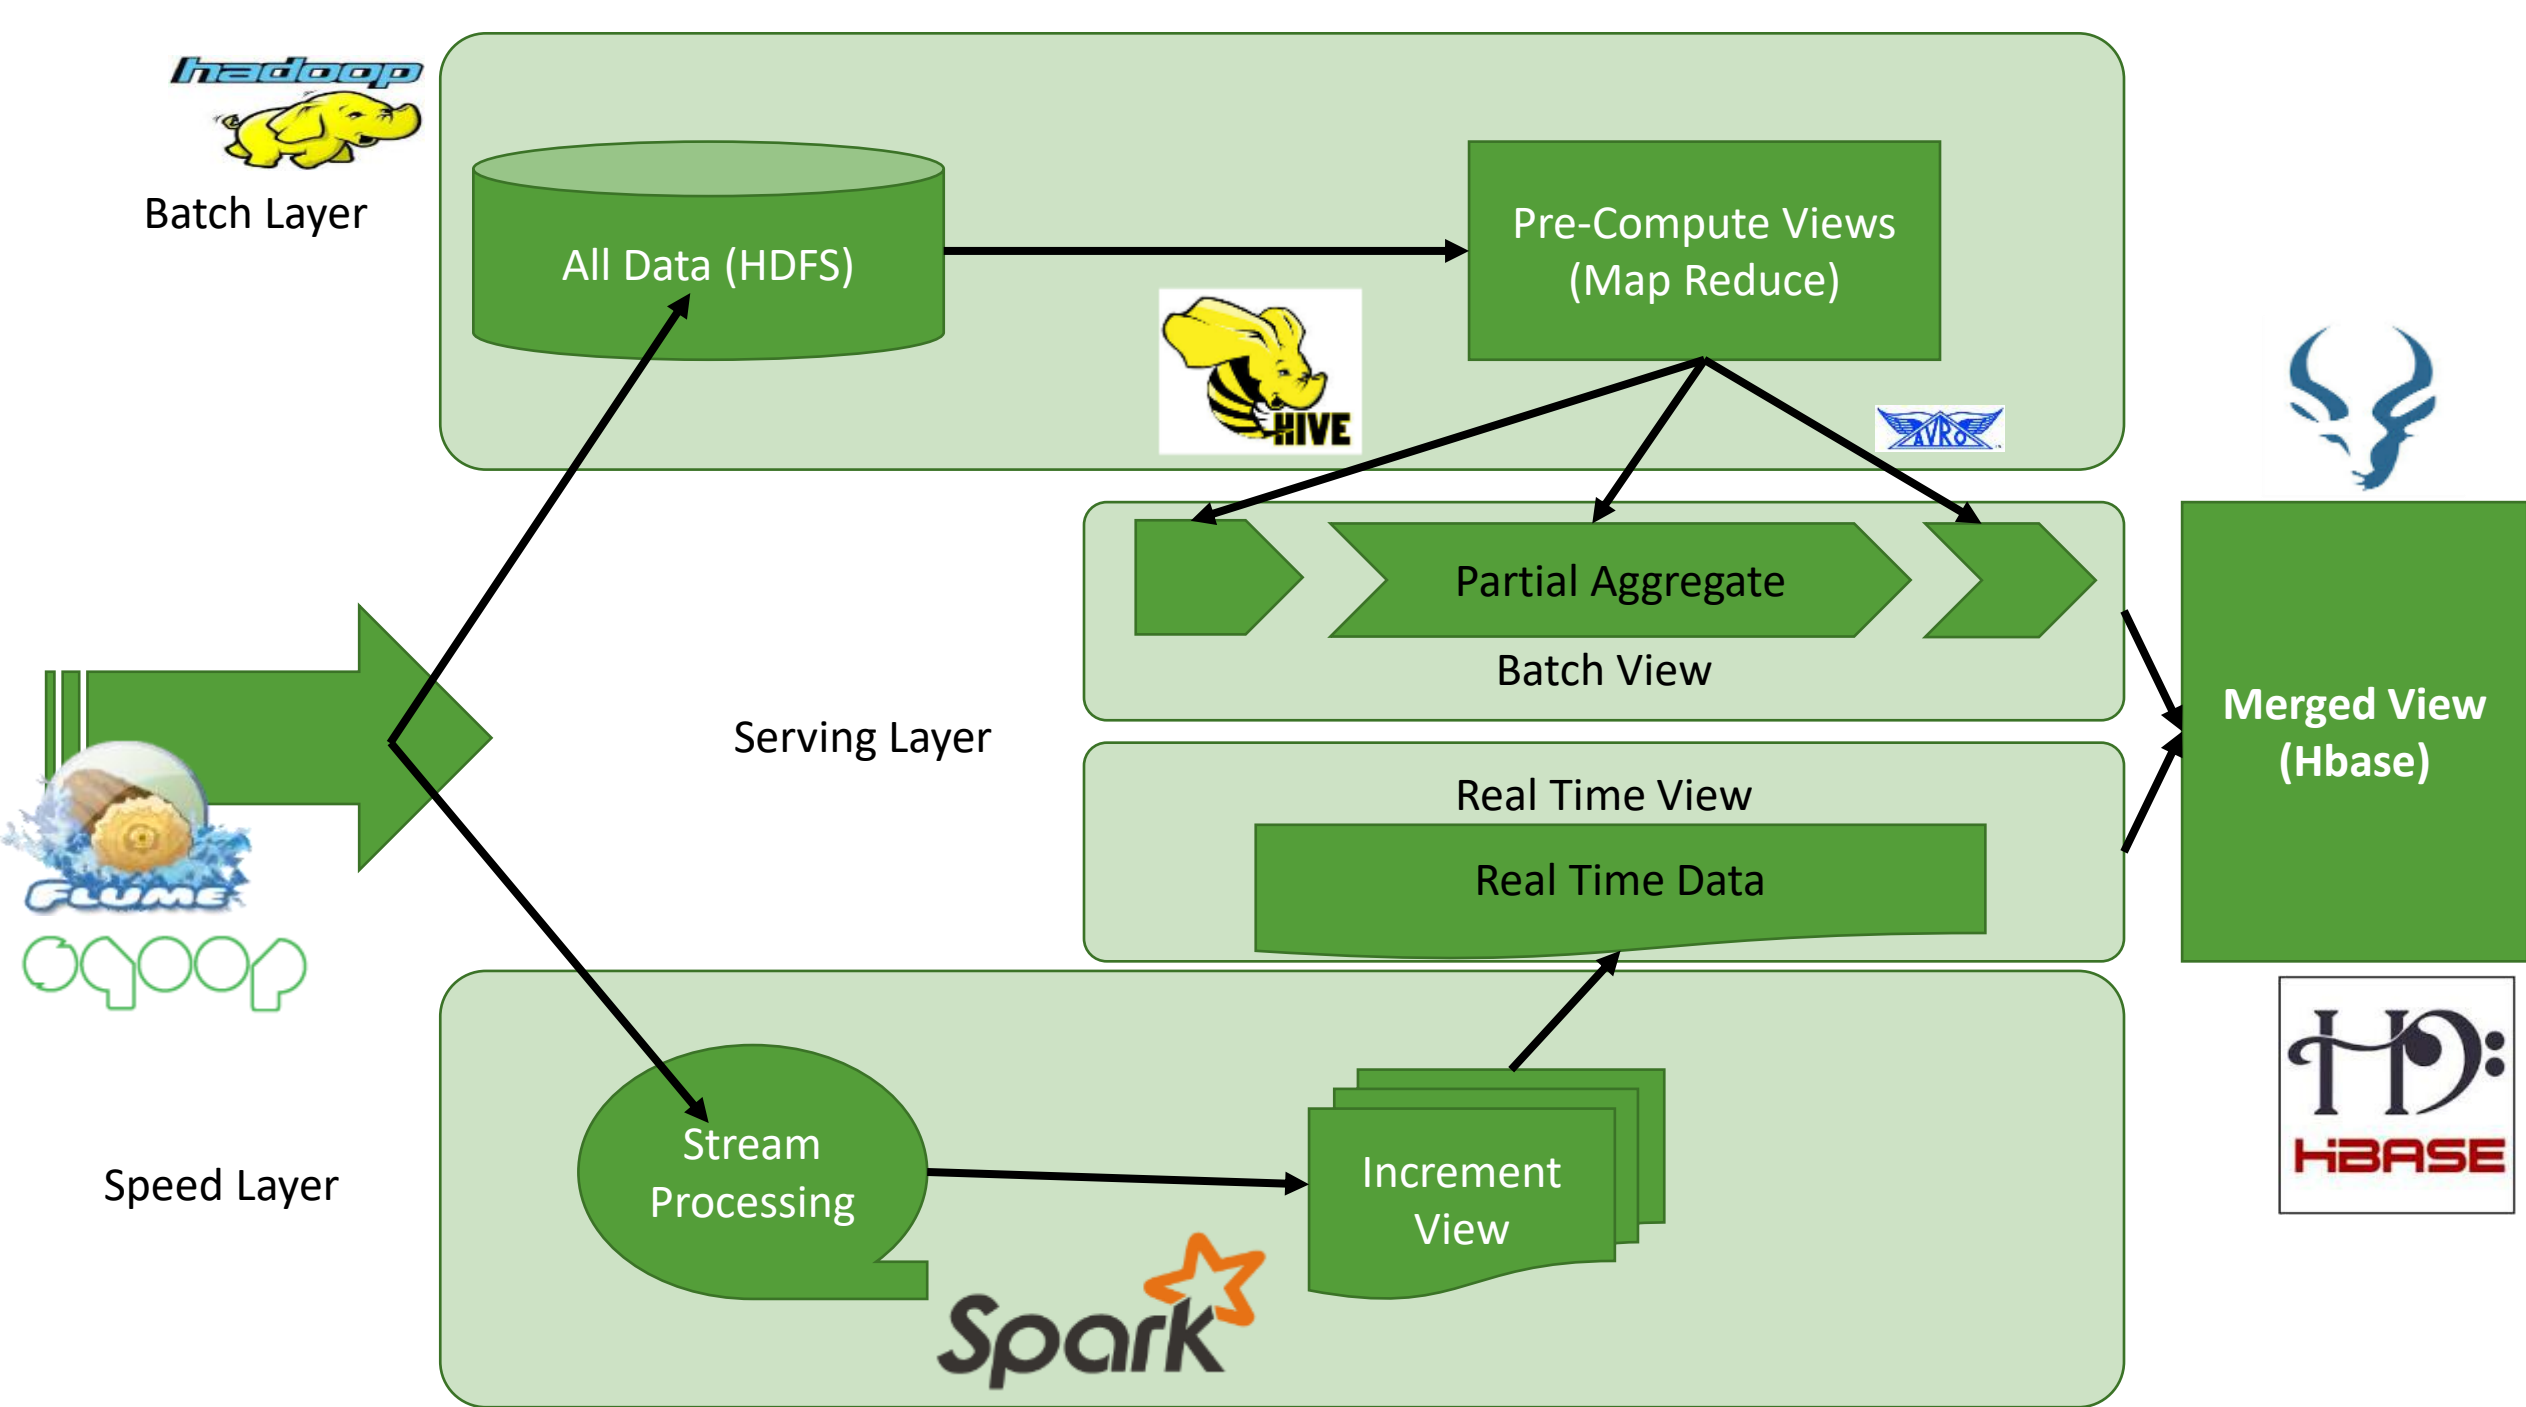
Advantages of using Parcels

- Distribution of CDH as a single object - Instead of having a separate package for each part of CDH, **parcels have just a single object to install**
- Internal consistency - All CDH components are matched, **eliminating the possibility of installing parts from different versions of CDH**
- Installation **outside of /usr**
- Installation of CDH without sudo - Parcel installation is handled by the Cloudera Manager Agent running as root or another user, so **you can install CDH without sudo.**

Types of Hosts and their roles

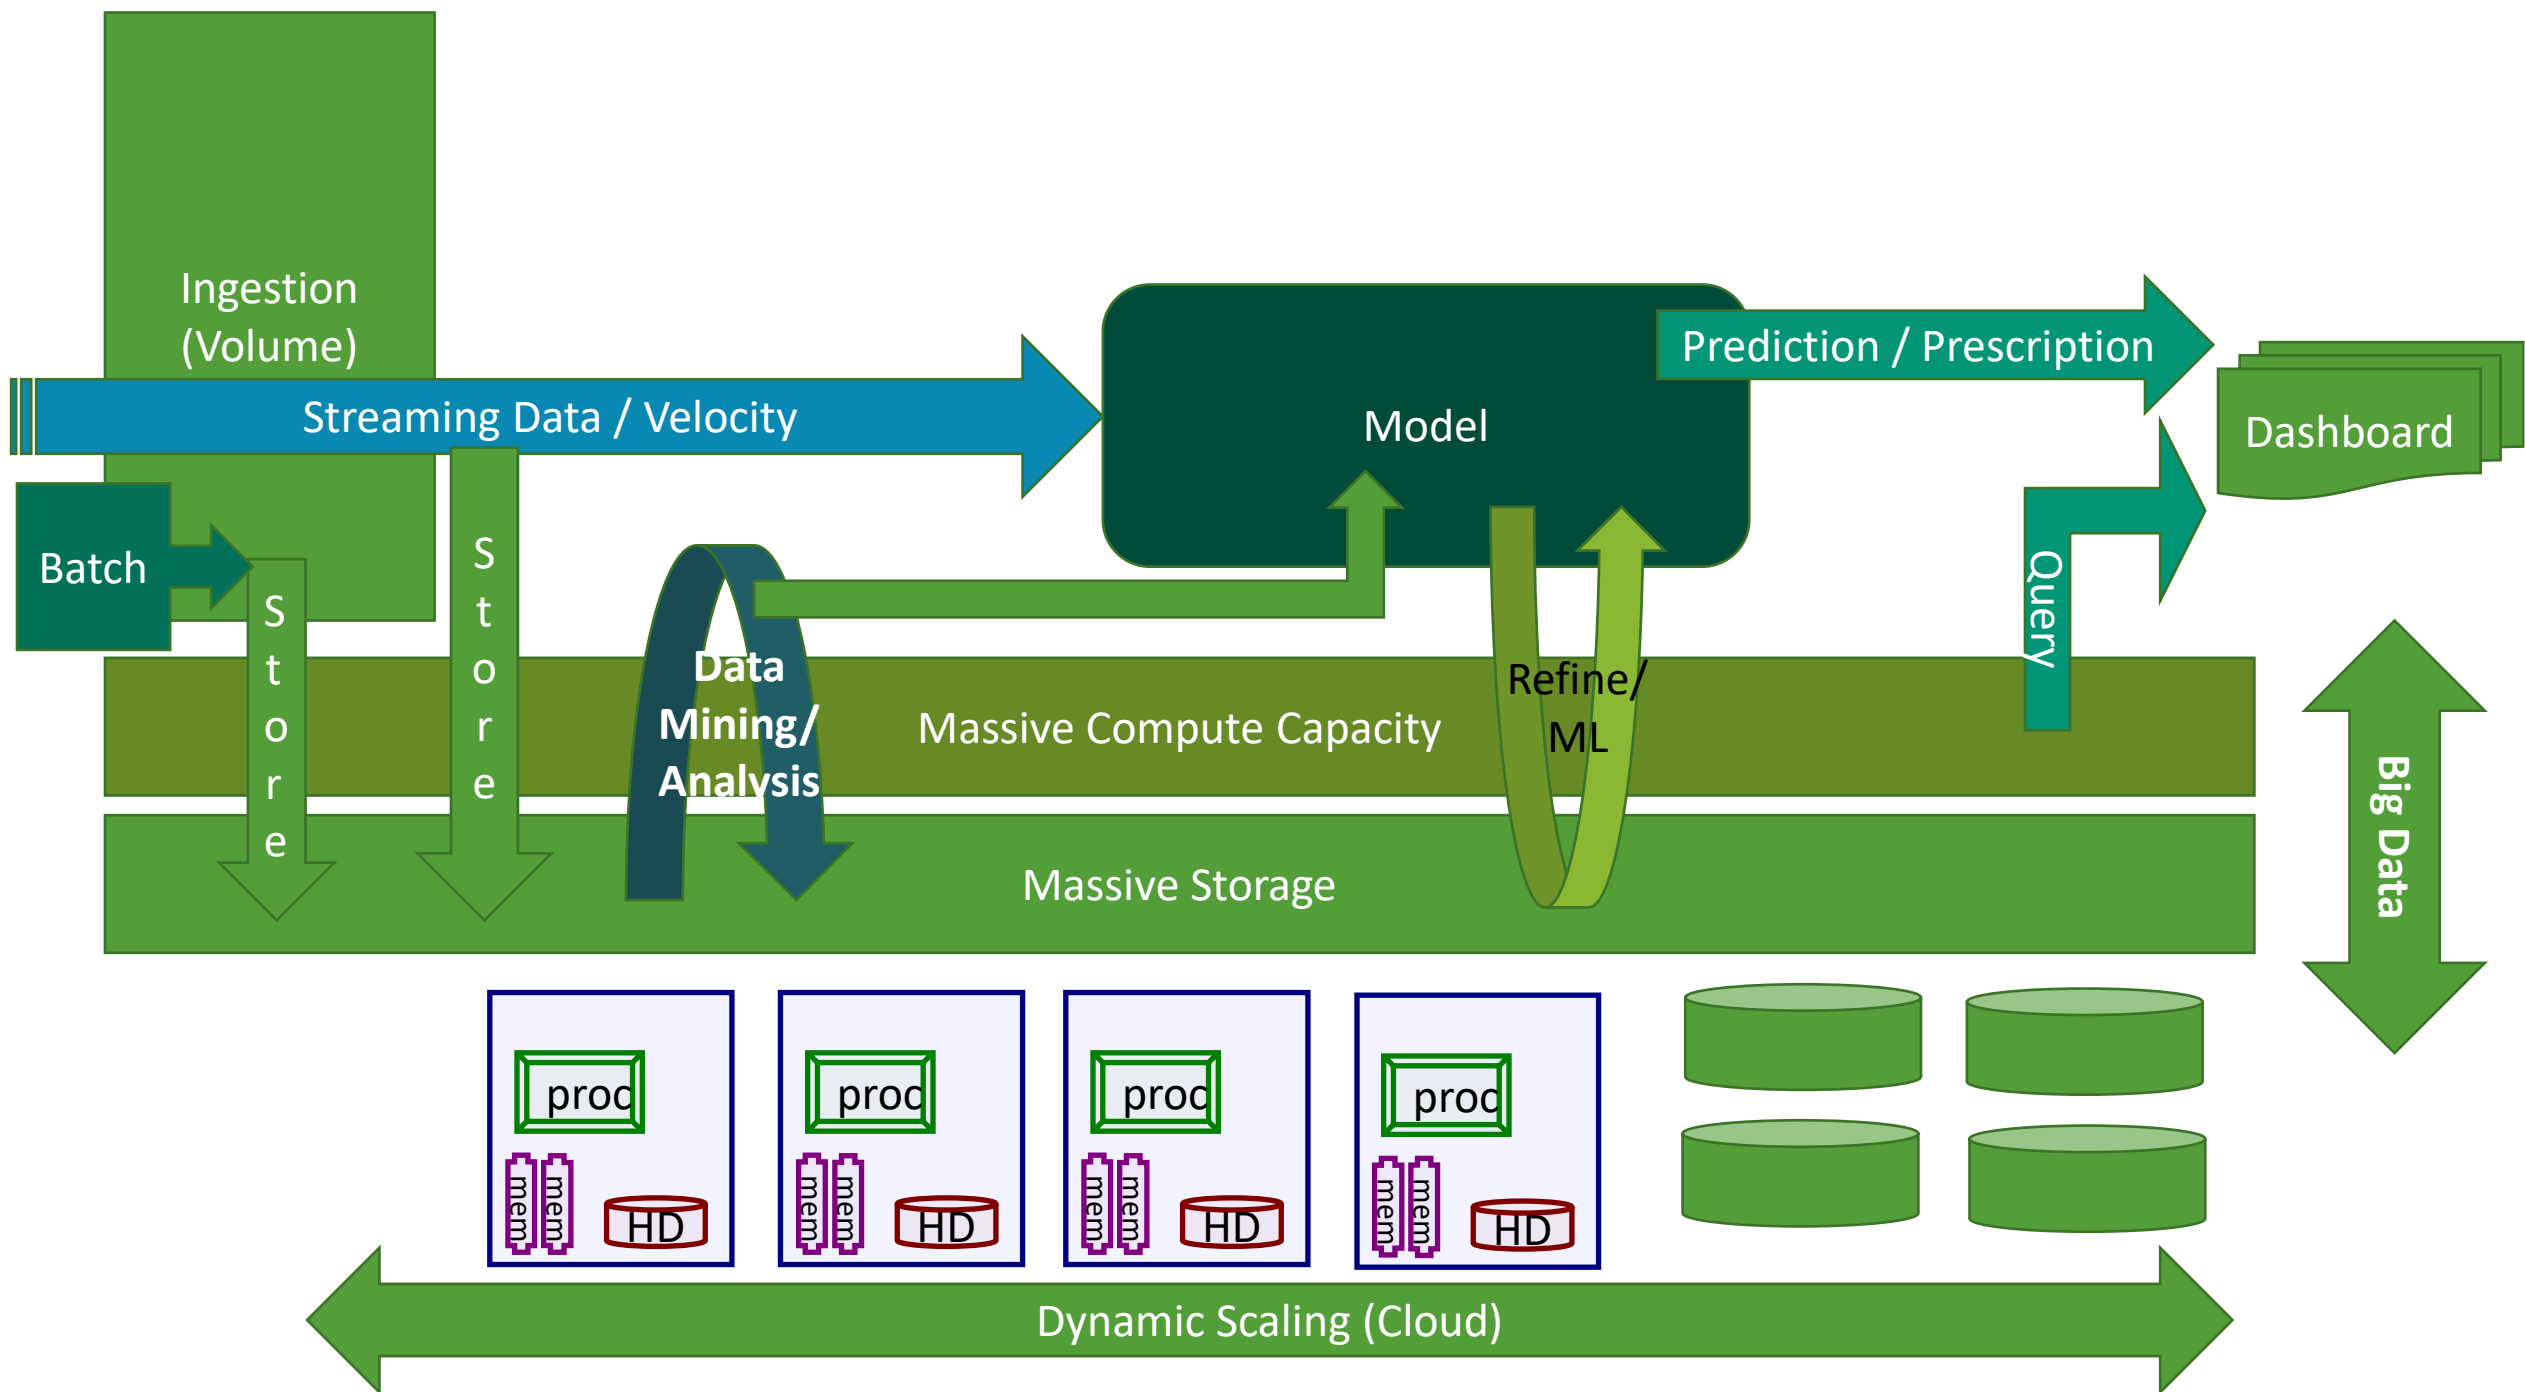
- **Master hosts** run Hadoop master processes such as the HDFS NameNode and YARN Resource Manager.
- **Worker hosts** primarily run DataNodes and other distributed processes such as Impalad, Datanode, Node Manager.
- **Utility hosts** run other cluster processes that are not master processes such as Cloudera Manager and the Hive Metastore.
- **Edge hosts** are client access points for launching jobs in the cluster. The number of Edge hosts required varies depending on the type and size of the workloads.

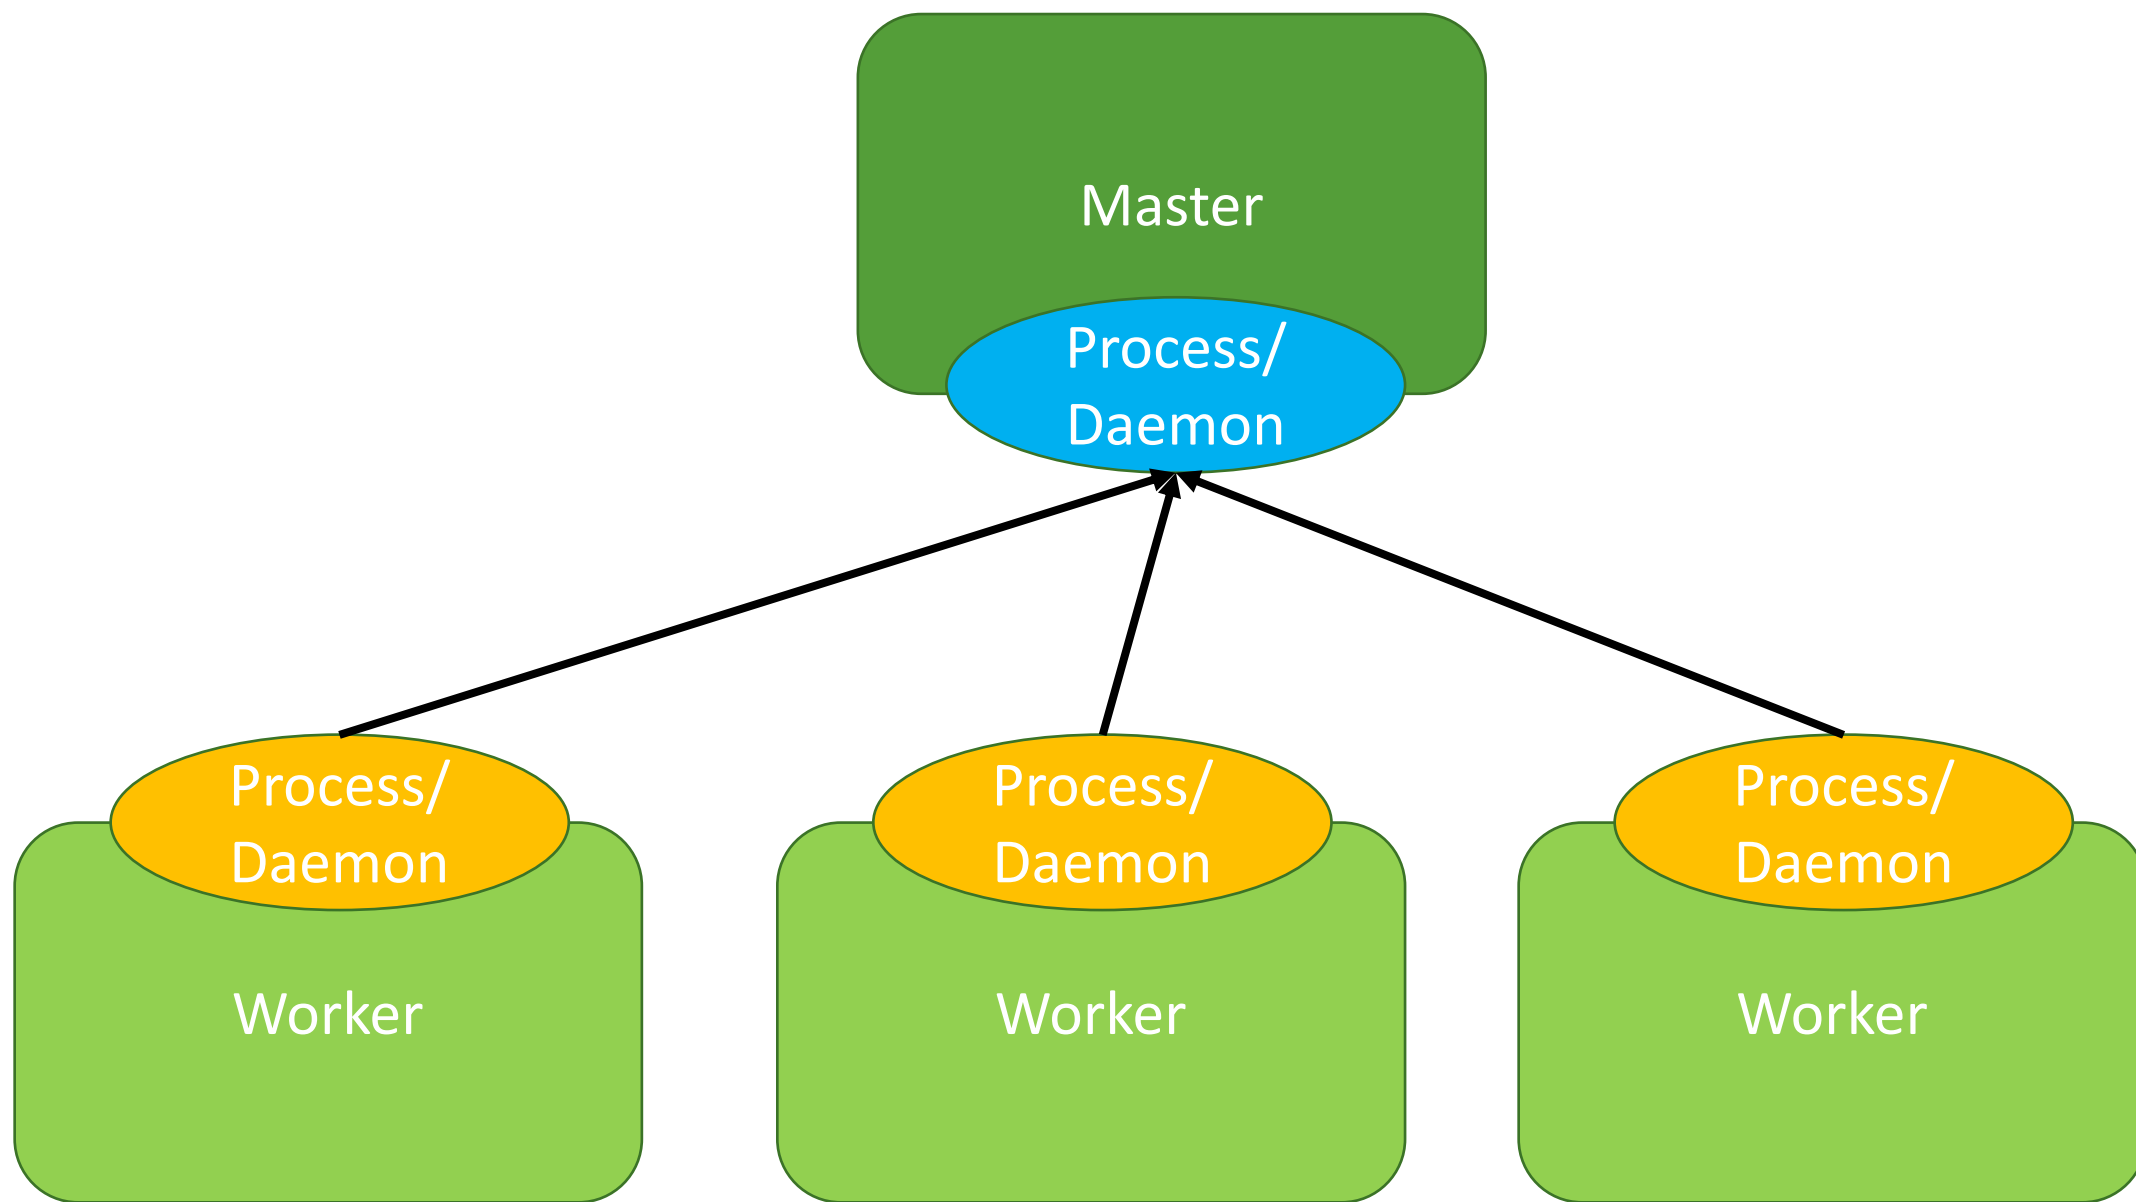






Hadoop EcoSystem							
Data Visualization							
	SAS Visual Analytics	Tableau	Qlick	SAP Lumira	R		
	D3.js	iCharts	Timeline JS	Apache Zeppelin	Pentaho		
System Deployment							
	Apache Ambari	Apache Mesos	Marathon	Apache Bigtop	HortonWorks HOYA	Deploop	
	Apache Eagle	Cloudera Manager	Brooklyn	Apache Helix	Buildoop	SequenceIQ Cloudbreak	
	Myriad						
Data Ingestion	Service Programming		Scheduling & DR	Security	Frameworks	Metadata	Machine Learning
Apache Flume	Apache Thrift	Apache Karaf	Apache Oozie	Apache Sentry	Jumbune	Metascope	Apache Mahout
Apache Sqoop	Apache Zookeeper	Twitter Elephant Bin	Linkedin Azkaban	Apache Knox Gatew	Spring XD	Apache Tika	WEKA
Facebook Scribe	Apache Avro	Linkedin Norbert	Apache Falcon	Apache Ranger	Cask Data App Platform		Cloudera Oryx
Apache Chukwa	Apache Curator		Shedoscope				Deeplearning4j
Apache Kafka							MADlib
Netflix Suro	Distributed Programming				SQL on Hadoop		H2O
Apache Samza	Apache Ignite	Apache Flink	Apache Twill	Kangaroo	Apache Hive	Facebook Presto	Sparkling Water
Cloudera Morphline	Apache Mapreduce	Apache Apex	Damballa PARKOUR	TinkerPop	Apache Hcatalog	Datasalt Splout SQL	Apache SystemML
HIHO	Apache Pig	Netflix Pigpen	Apache Hama	Pachyderm Mapred	Apache Trafodion	Apache Tajo	
Apache Nifi	JAQL	AMPLAB SIMR	Datasalt Pangool	Apache Beam	Apache HAWQ	Apache MRQL	
Apache ManifoldCF	Apache Spark	Facebook Corona	Apache Tez		Apache Drill	Kylin	
	Apache Storm	Apache REEF	Apache DataFu		Cloudera Impala		
	NoSQL Databases					NewSQL Databases	Deep Learning
	Wide Column	Document	Key-Value	Graph	Stream Data Model	TokuDB	Caffe
	Apache Hbase	MongoDB	Redis	Giraph	EventStore	HandlerSocket	Microsoft CNTK
	Apache Cassandra	RethinkDB	Linkedin Voldemort	Neo4J		Akiban Server	Tensor Flow
	Hypertable	ArangoDB	RocksDB	TitanDB		Drizzle	Theano
	Apache Accumulo	CouchDB	OpenTSDB	OrientDB		Haeinsa	Torch
	Apache Kudu	DynamoDB				SenseiDB	Mxnet
	Apache Parquet	Gemfire				Sky	Chainer
						BayesDB	Keras
						InfluxDB	
						VoldDB	
						SAP HANA	
Distributed File System							
	Apache HDFS	Quantcast File Syste	Lustre File System	GridGain			
	Red Hat GlusterFS	Ceph File System	Alluxio	XtreemFS			





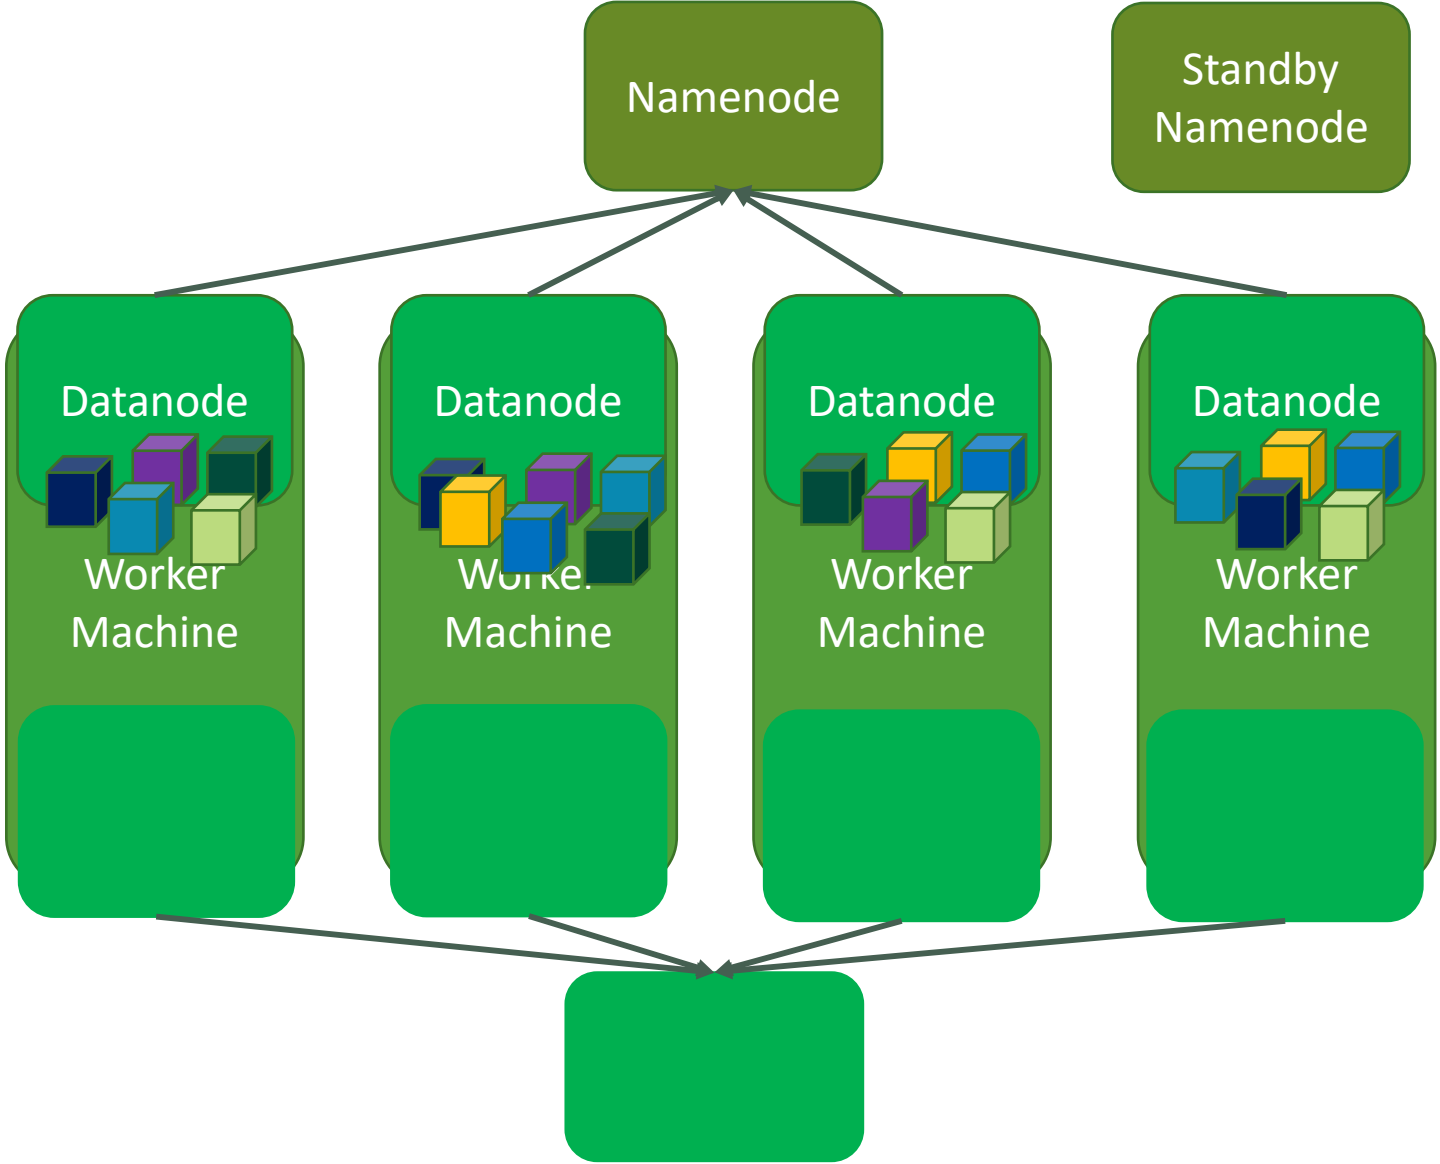


YARN – Yet Another Resource Negotiator

The diagram consists of a large green rounded rectangle containing two blue rounded rectangles stacked vertically. The top blue rectangle contains the text 'YARN – Yet Another Resource Negotiator' and the bottom blue rectangle contains the text 'HDFS – Hadoop Distributed File System'.

HDFS – Hadoop Distributed File System

HDFS




200 MB


128 MB

72 MB

YARN



HDFS and YARN





Hadoop EcoSystem							
Data Visualization							
	SAS Visual Analytics	Tableau	Qlick	SAP Lumira	R		
	D3.js	iCharts	Timeline JS	Apache Zeppelin	Pentaho		
System Deployment							
	Apache Ambari	Apache Mesos	Marathon	Apache Bigtop	HortonWorks HOYA	Deploop	
	Apache Eagle	Cloudera Manager	Brooklyn	Apache Helix	Buildoop	SequenceIQ Cloudbreak	
	Myriad						
Data Ingestion	Service Programming		Scheduling & DR	Security	Frameworks	Metadata	Machine Learning
Apache Flume	Apache Thrift	Apache Karaf	Apache Oozie	Apache Sentry	Jumbune	Metascope	Apache Mahout
Apache Sqoop	Apache Zookeeper	Twitter Elephant Bin	Linkedin Azkaban	Apache Knox Gatew	Spring XD	Apache Tika	WEKA
Facebook Scribe	Apache Avro	Linkedin Norbert	Apache Falcon	Apache Ranger	Cask Data App Platform		Cloudera Oryx
Apache Chukwa	Apache Curator		Shedoscope				Deeplearning4j
Apache Kafka							MADlib
Netflix Suro	Distributed Programming				SQL on Hadoop		H2O
Apache Samza	Apache Ignite	Apache Flink	Apache Twill	Kangaroo	Apache Hive	Facebook Presto	Sparkling Water
Cloudera Morphline	Apache Mapreduce	Apache Apex	Damballa PARKOUR	TinkerPop	Apache Hcatalog	Datasalt Splout SQL	Apache SystemML
HIHO	Apache Pig	Netflix Pigpen	Apache Hama	Pachyderm Mapred	Apache Trafodion	Apache Tajo	
Apache Nifi	JAQL	AMPLAB SIMR	Datasalt Pangool	Apache Beam	Apache HAWQ	Apache MRQL	
Apache ManifoldCF	Apache Spark	Facebook Corona	Apache Tez		Apache Drill	Kylin	
	Apache Storm	Apache REEF	Apache DataFu		Cloudera Impala		
	NoSQL Databases					NewSQL Databases	Deep Learning
	Wide Column	Document	Key-Value	Graph	Stream Data Model	TokuDB	Caffe
	Apache Hbase	MongoDB	Redis	Giraph	EventStore	HandlerSocket	Microsoft CNTK
	Apache Cassandra	RethinkDB	Linkedin Voldemort	Neo4J		Akiban Server	Tensor Flow
	Hypertable	ArangoDB	RocksDB	TitanDB		Drizzle	Theano
	Apache Accumulo	CouchDB	OpenTSDB	OrientDB		Haeinsa	Torch
	Apache Kudu	DynamoDB				SenseiDB	Mxnet
	Apache Parquet	Gemfire				Sky	Chainer
						BayesDB	Keras
						InfluxDB	
						VoldDB	
						SAP HANA	
Distributed File System							
	Apache HDFS	Quantcast File Syste	Lustre File System	GridGain			
	Red Hat GlusterFS	Ceph File System	Alluxio	XtreemFS			



MapReduce

The diagram consists of a large green rounded rectangle containing two blue rounded rectangles stacked vertically. The top blue rectangle is labeled 'MapReduce' and the bottom one is labeled 'HDFS – Hadoop Distributed File System'.

HDFS – Hadoop Distributed File System

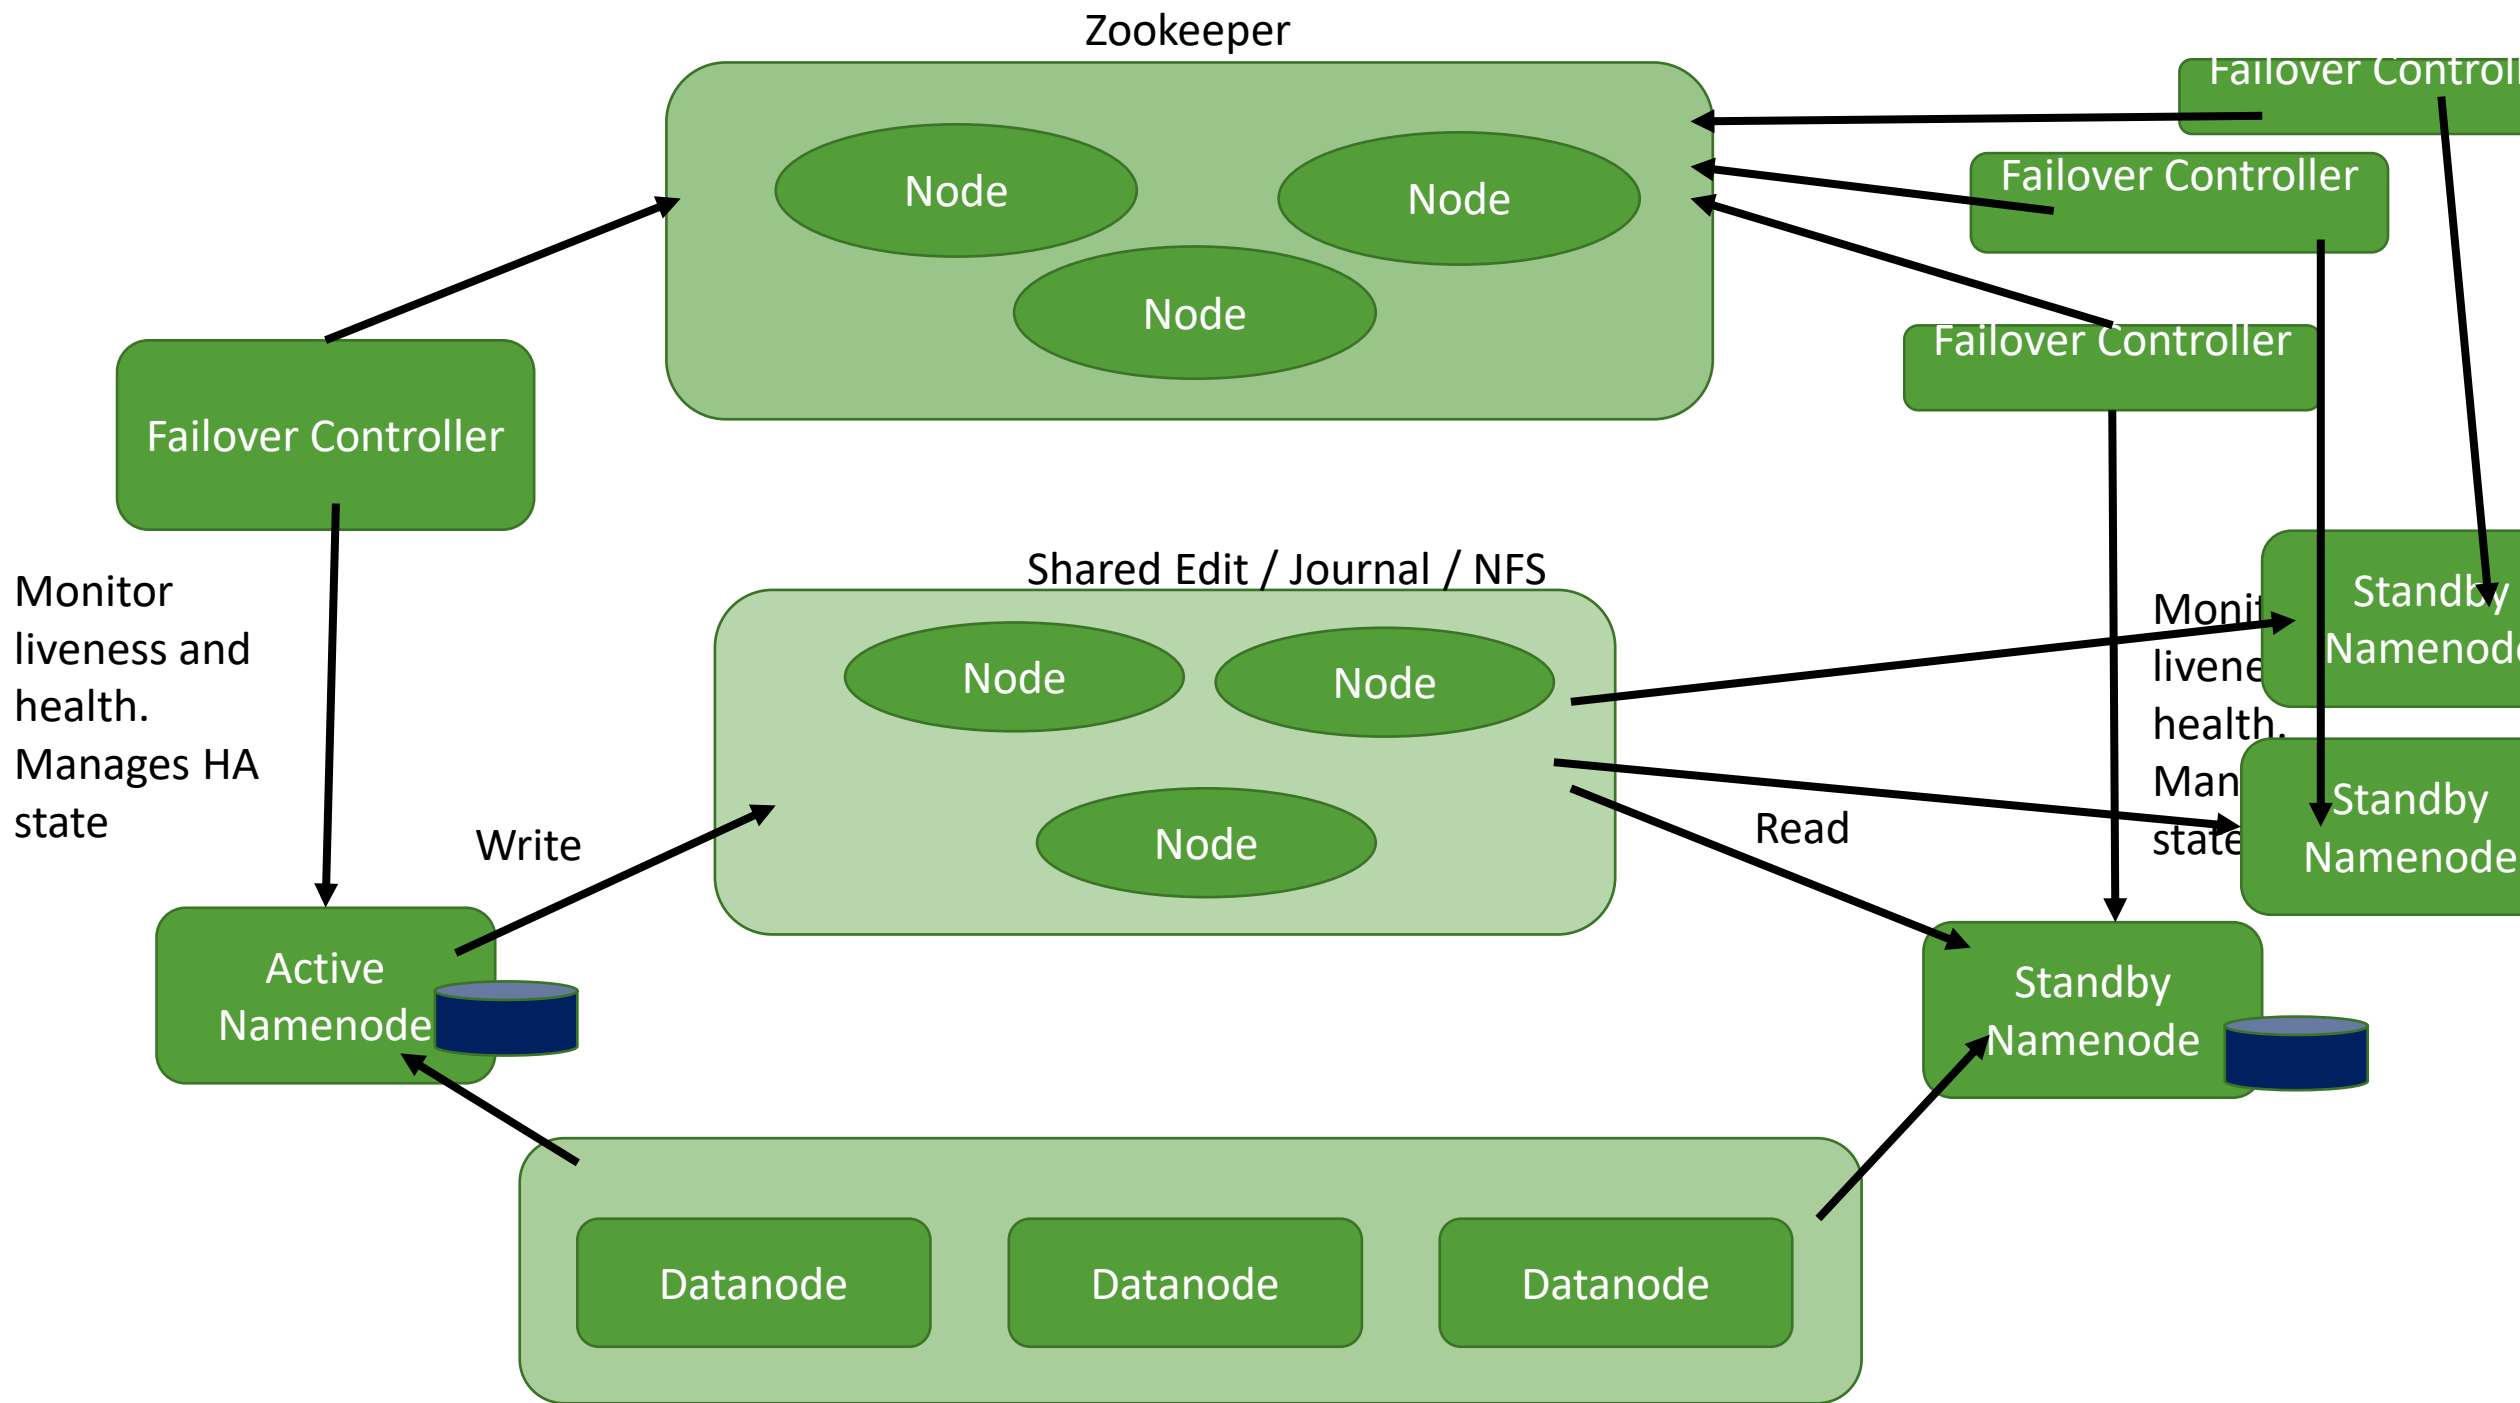
```
graph TD; RM[Resource Manager] --- NM1[Node Manager]; RM --- NM2[Node Manager]; RM --- NM3[Node Manager];
```

Resource Manager

Node Manager

Node Manager

Node Manager





HDFS Roles

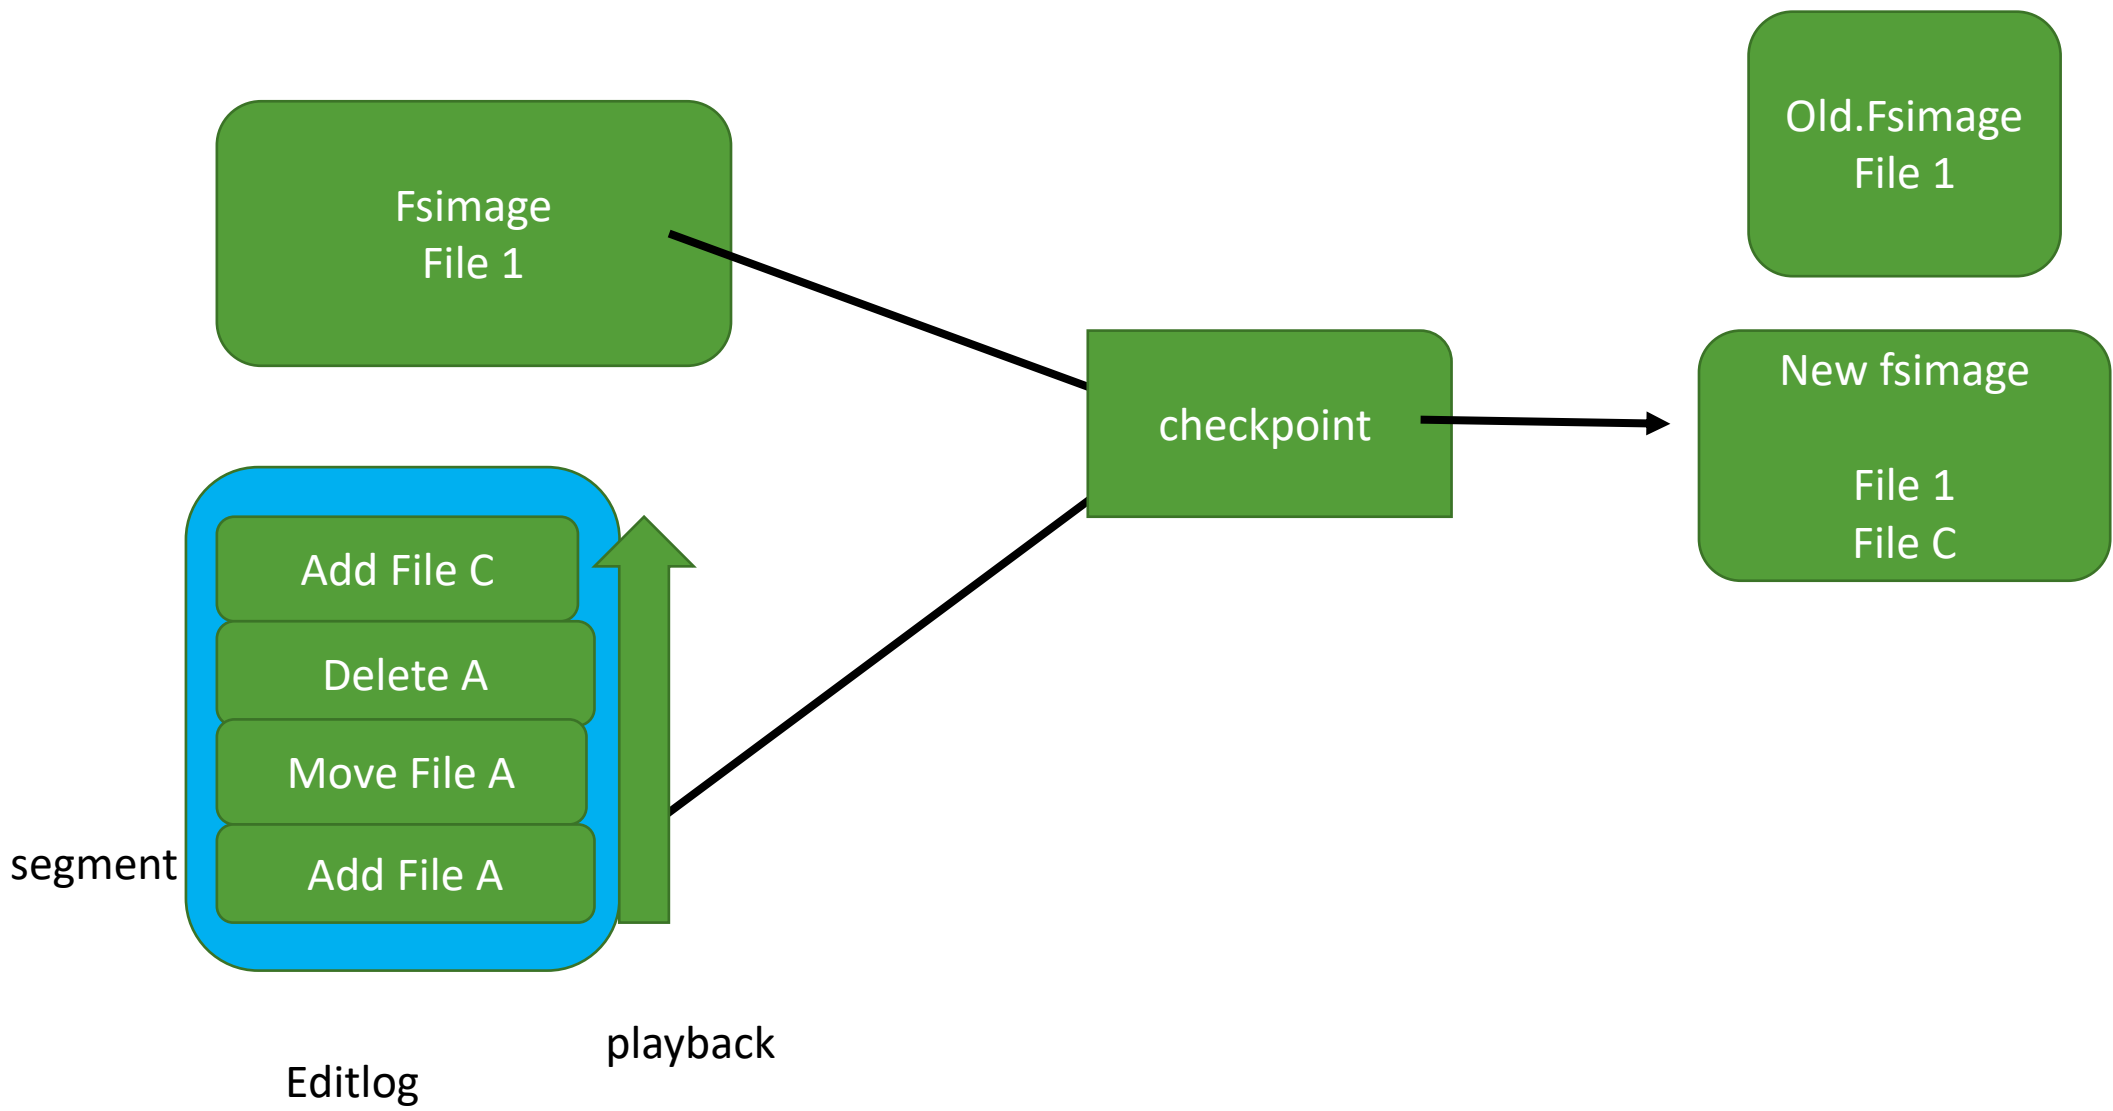
- Namenode
- Secondary Namenode
- Balancer
- HttpFS
- NFS Gateway
- Datanode
- Failover Controller
- Gateway
- JournalNode


Hadoop Ecosystem


datanode storage utilization %











Secondary Namenode

1. Check Precondition

3. Fetch new fsimage
and edits

4. Reload new fsimage
if any

5. Replay new edit logs

6. Create new fsimage



7. Send new fsimage
to NN

Namenode


2. Roll NN's edit log

http calls

8. NN saves
new fsimage
Along with MD5



Standby NN



1. Check precondition

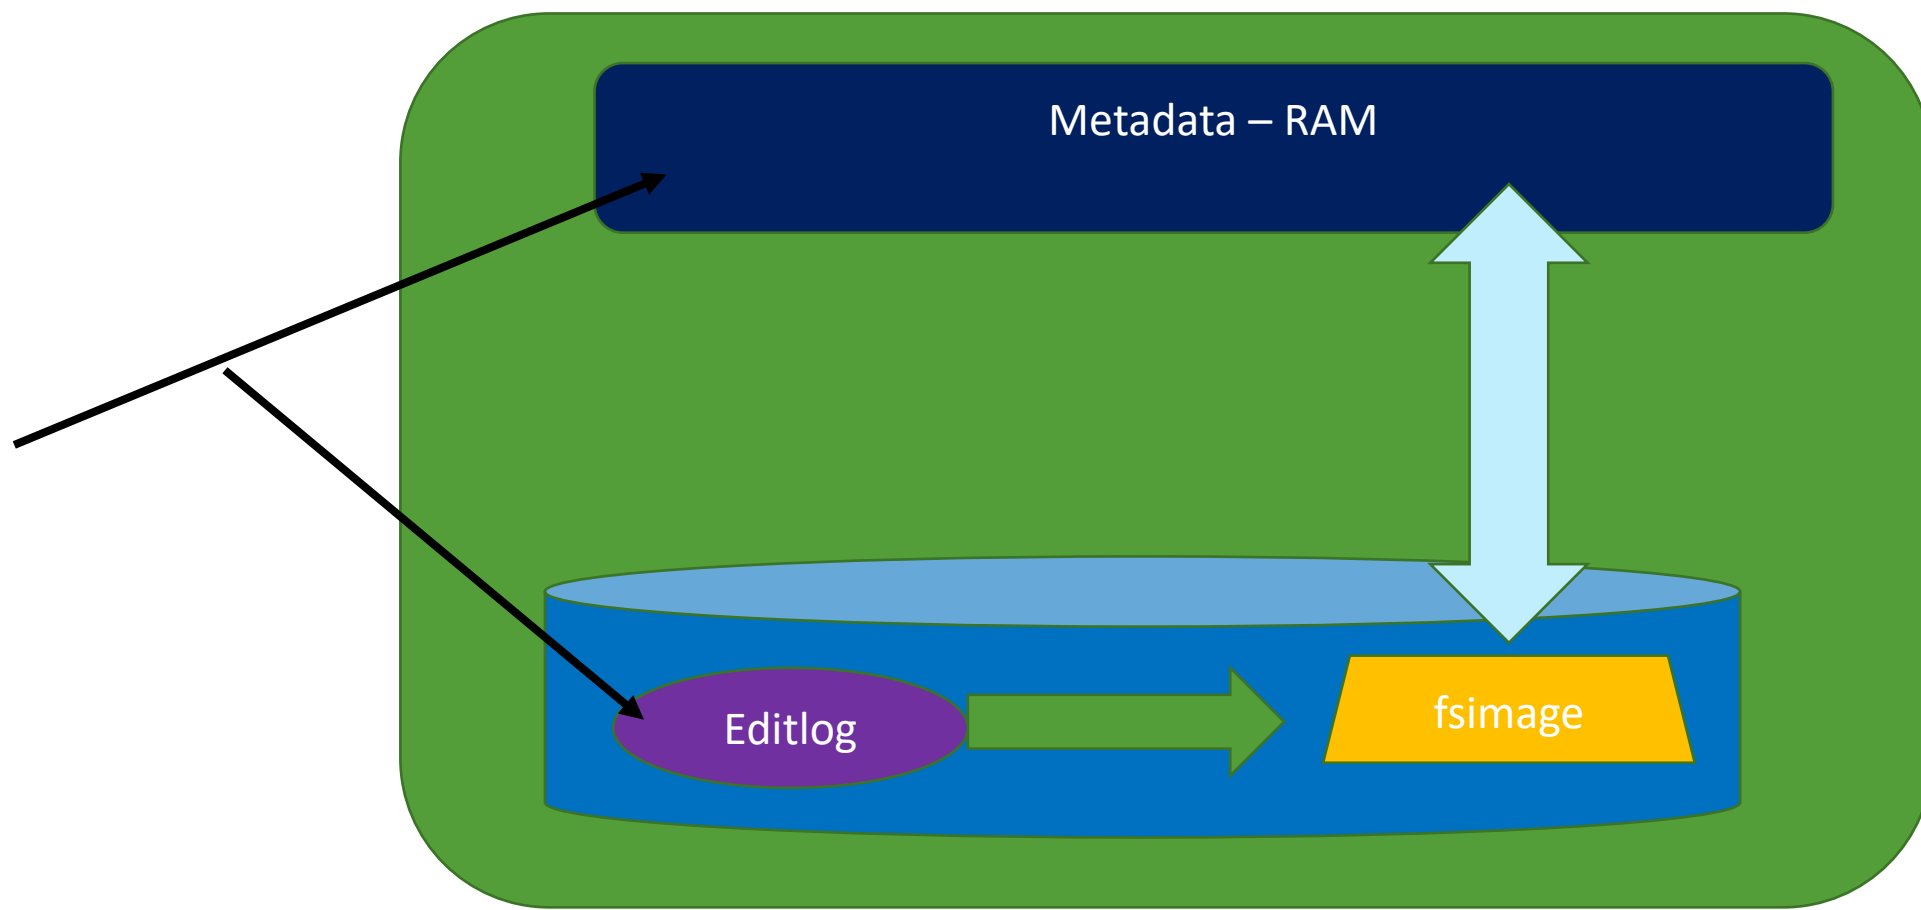
2. Save new fsimage + MD5

3. PUT or send the fsimage
To NN

Active NN








Saves the fsimage and MD5

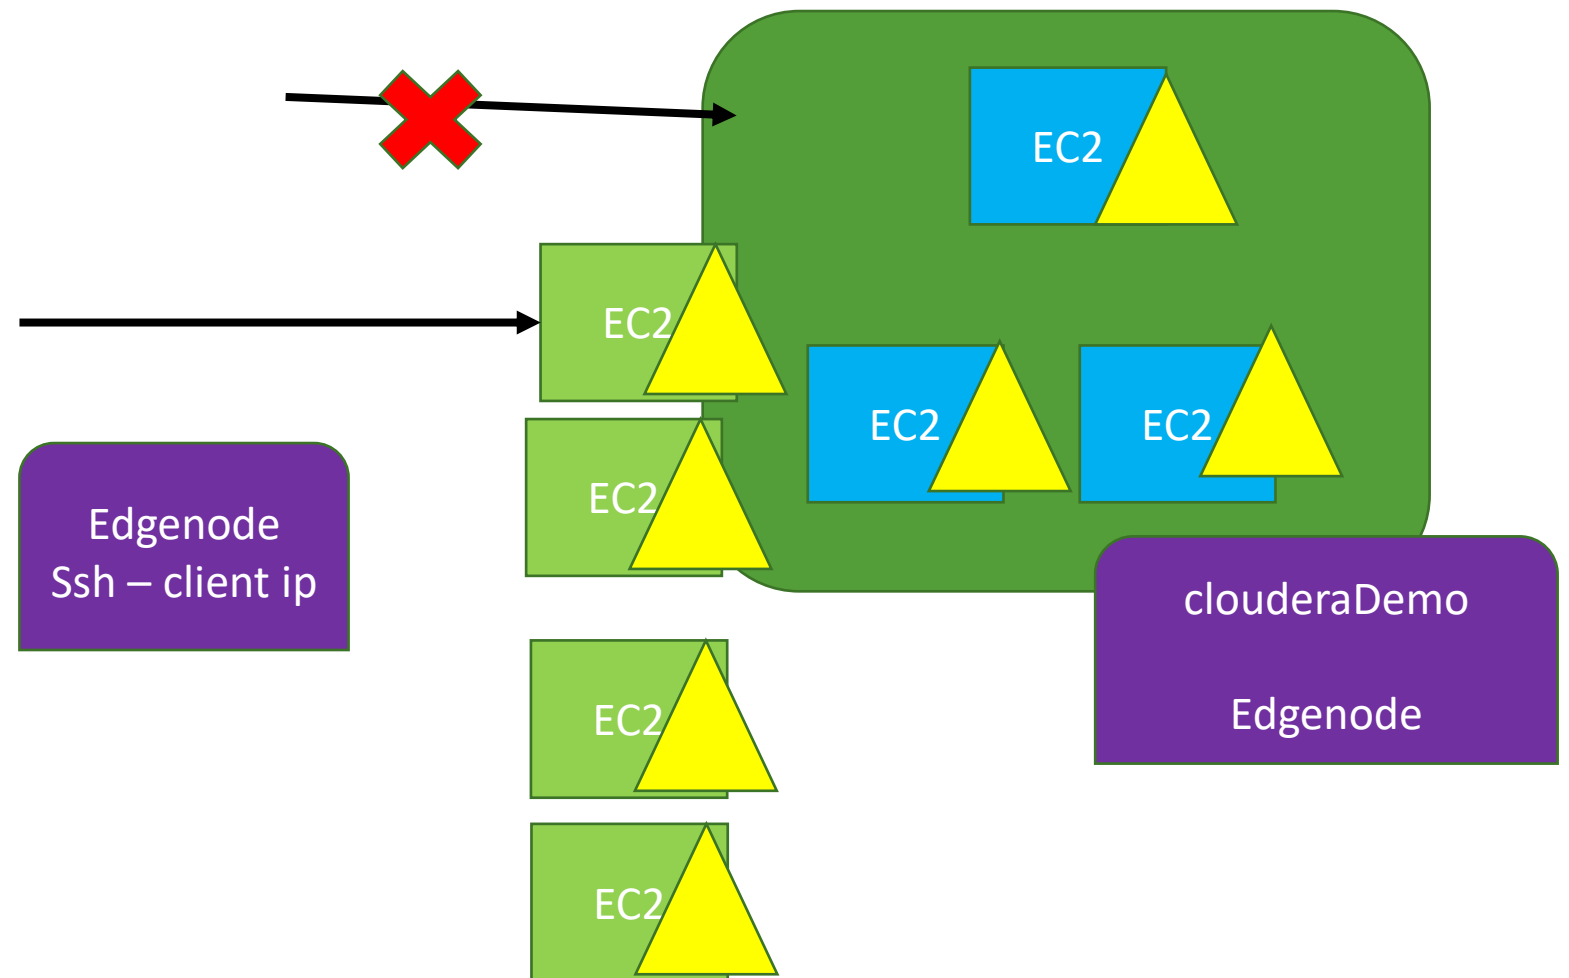



Property that controls checkpoint


```
dfs.namenode.checkpoint.period - 3600  
dfs.namenode.checkpoint.txns - 1000000  
dfs.namenode.checkpoint.check.period - 60  
dfs.namenode.num.checkpoints.retained - 1000000  
dfs.namenode.num.extra.edits.retained - 2  
dfs.namenode.max.extra.edits.segments.retained - 10000
```







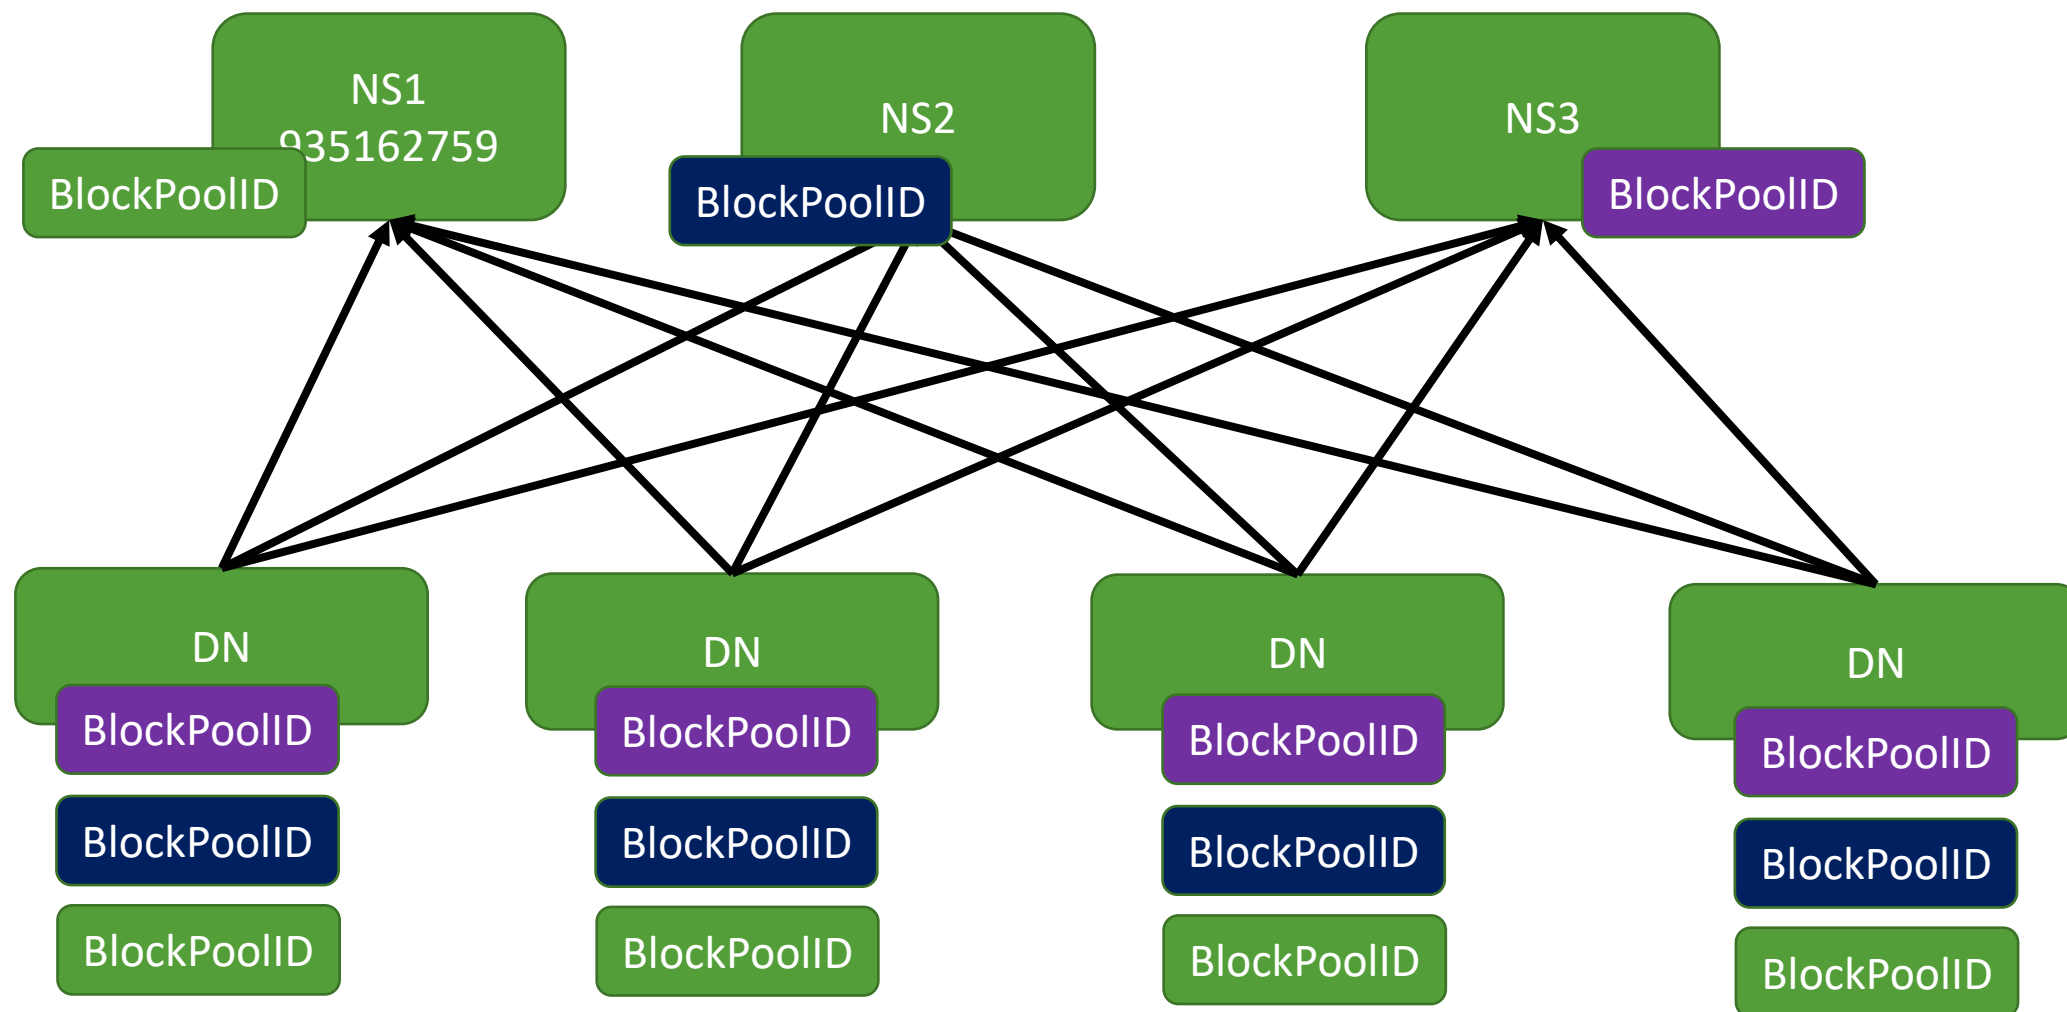


WebHDFS

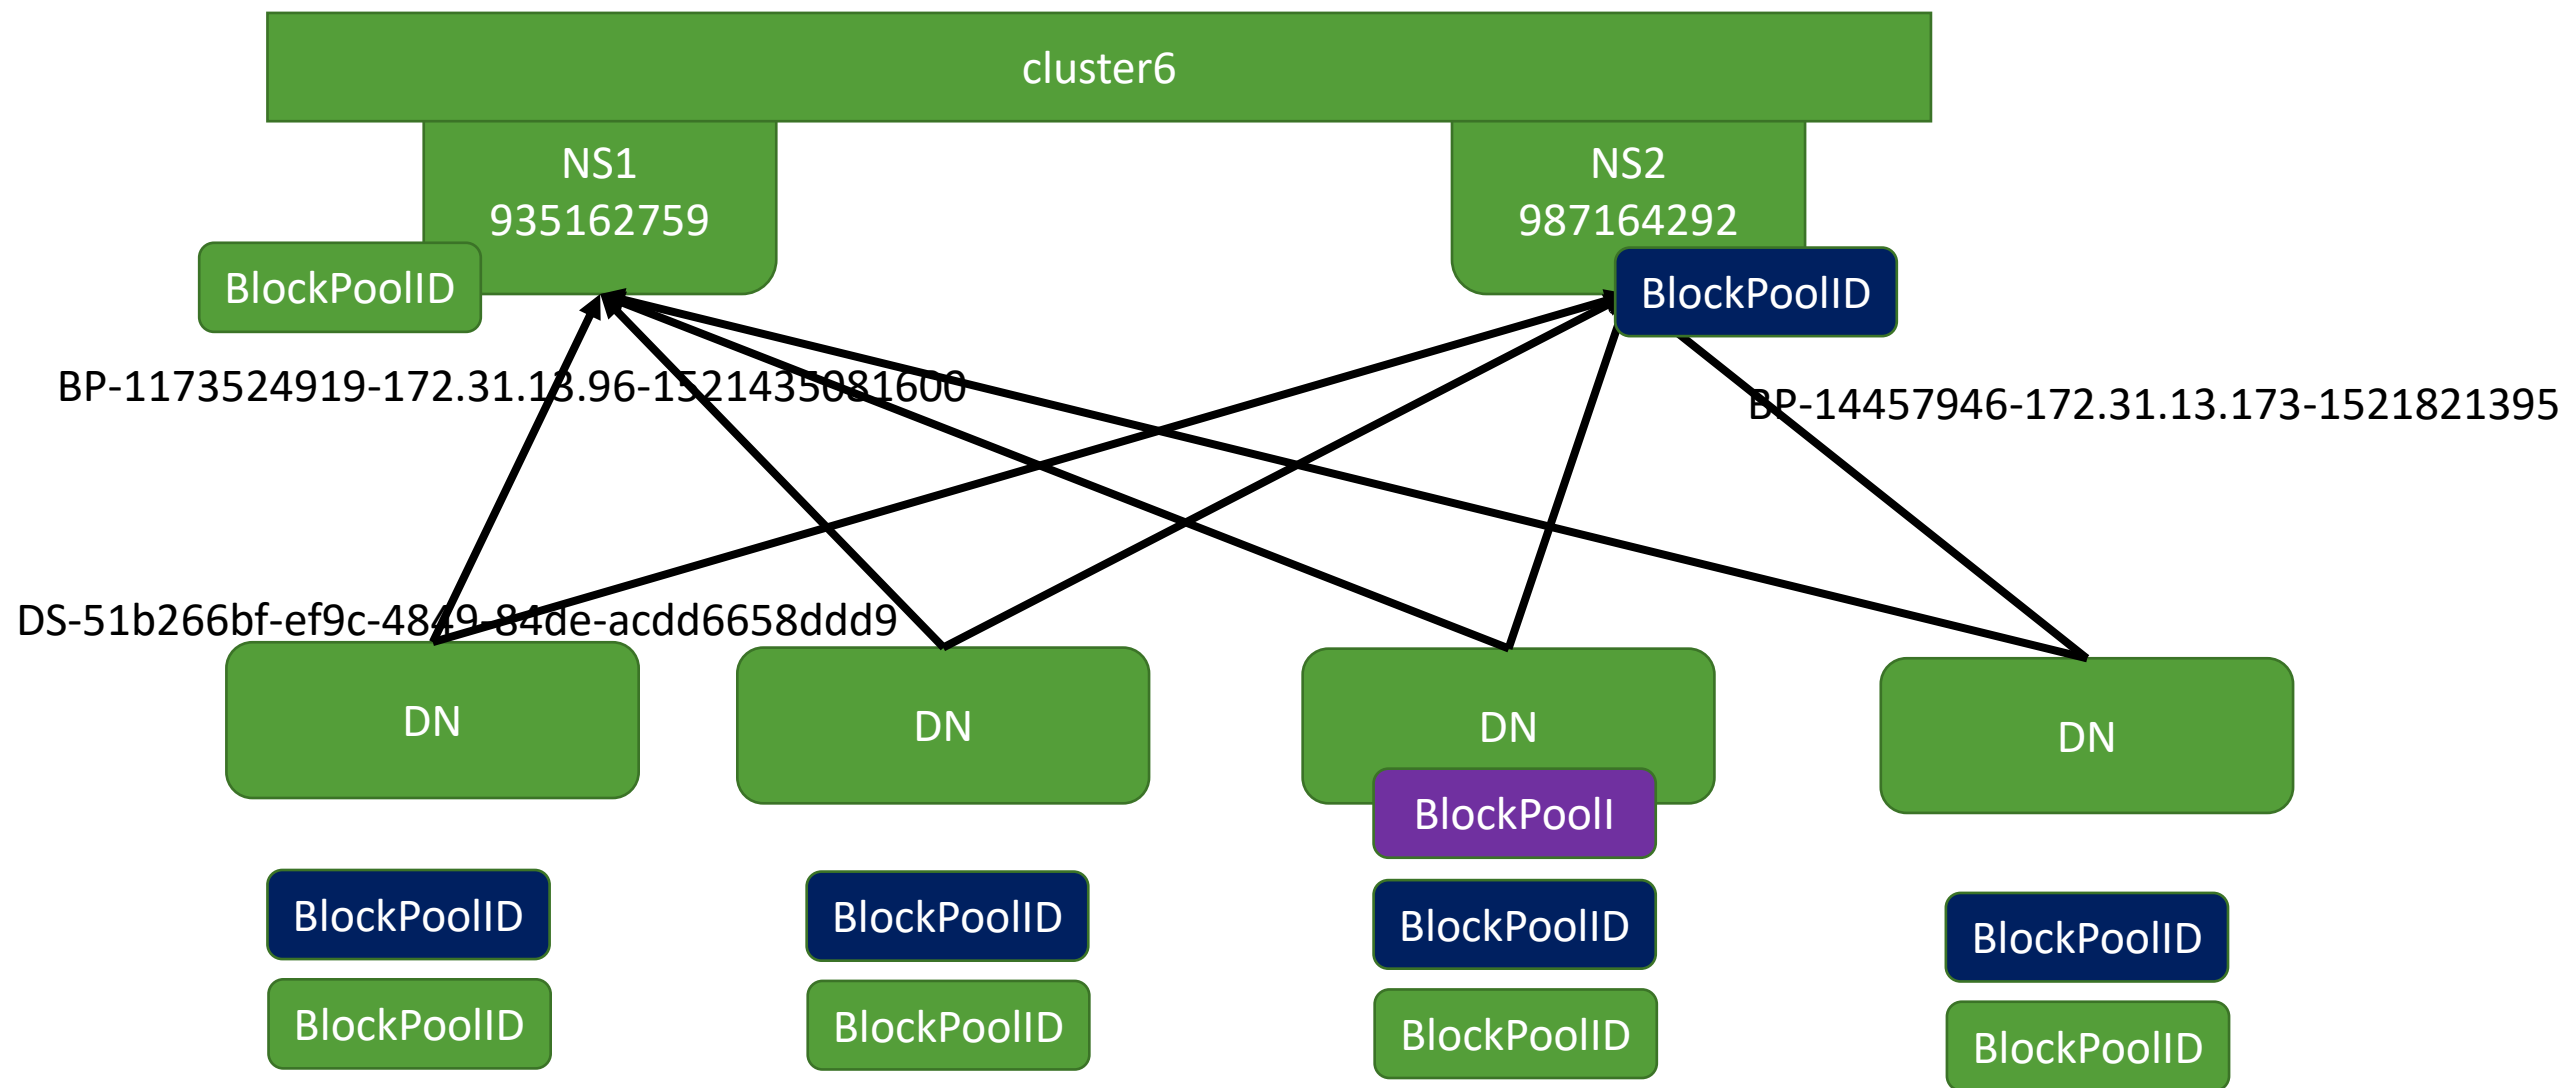
- Complete HDFS Interface (Read, Write, Directory Operations, Permissions)
- HTTP REST API
- Wire Compatibility – Can talk to different version of Hadoop
- Secure Authentication – Can be used along with Kerberos
- Data Locality – Calls will be redirected to Datanode for Read/Write
- HDFS Built-in Component

Federation


BP-1173524919-172.31.13.96-1521435081600



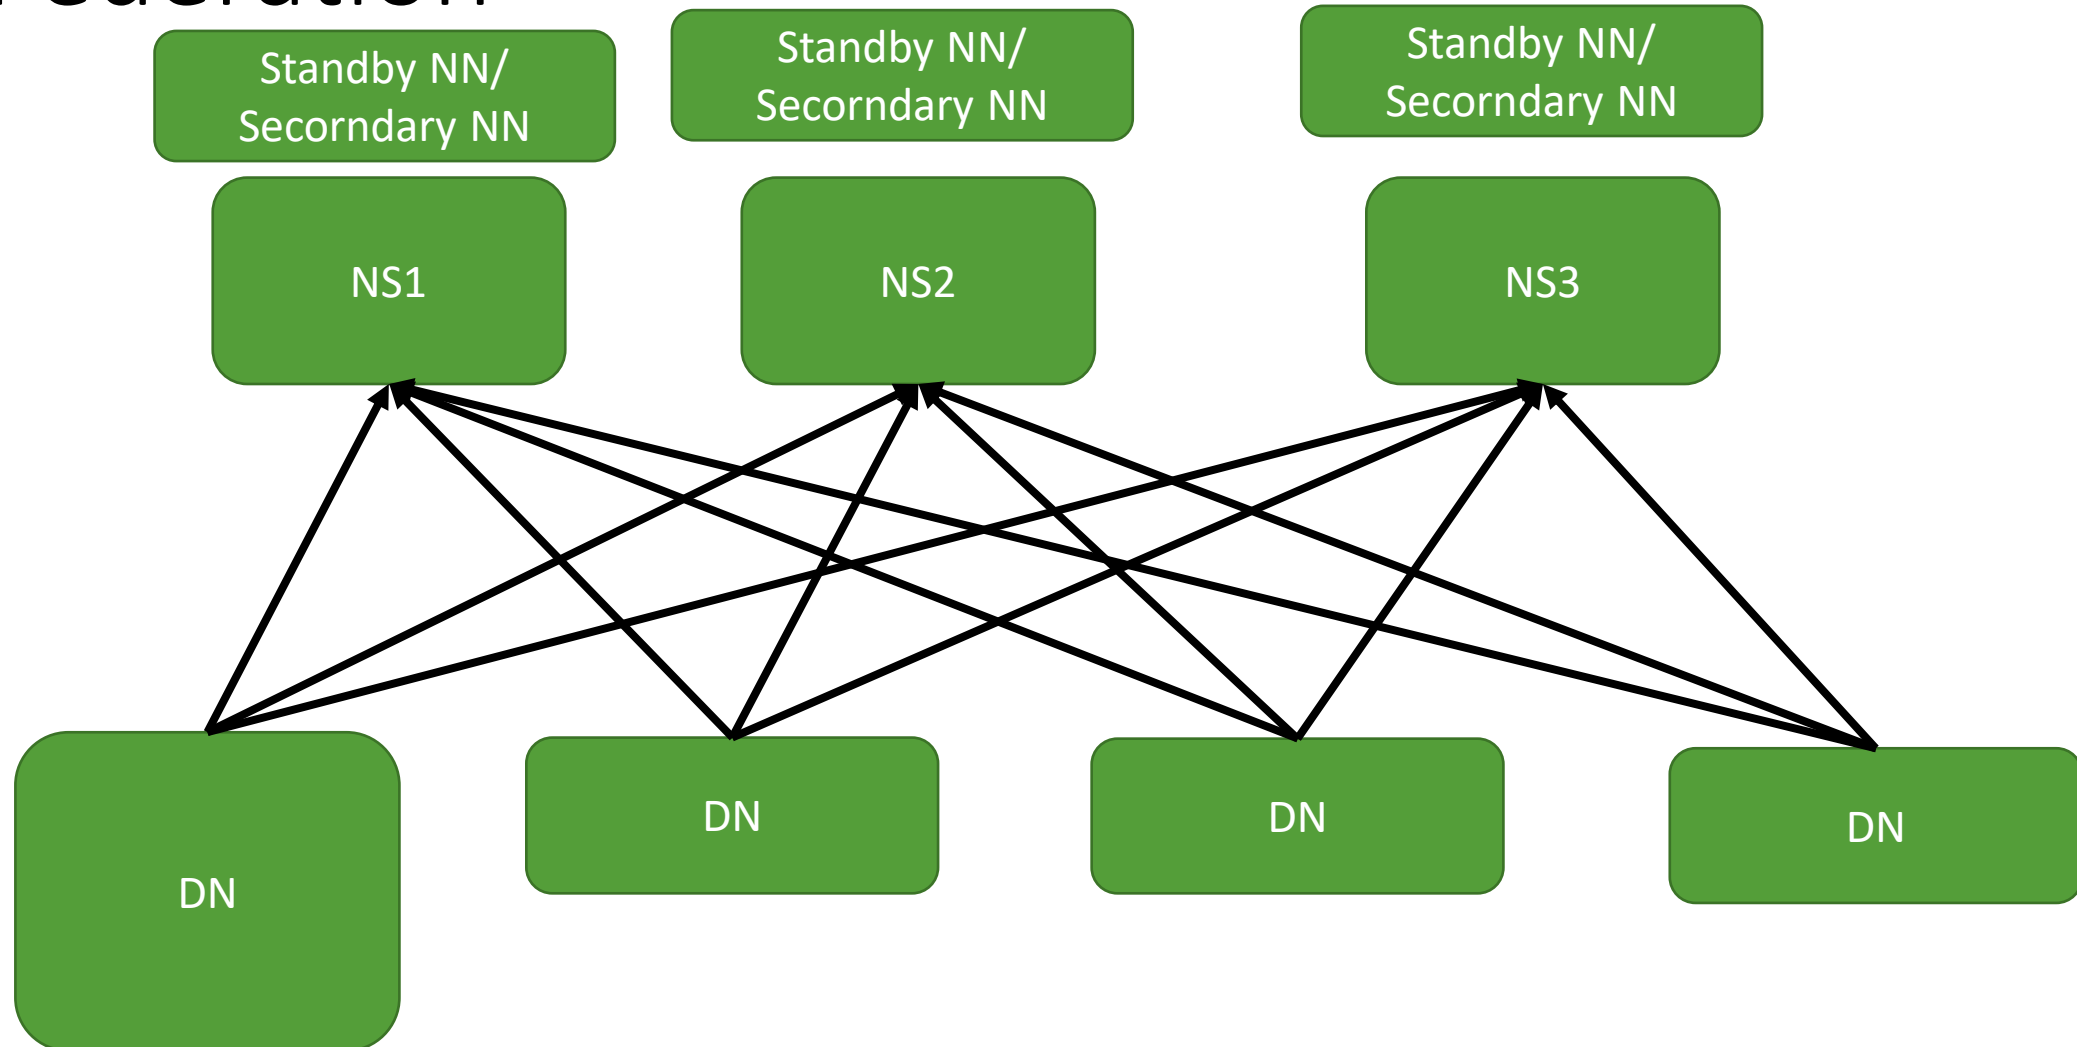
Federation



Federation



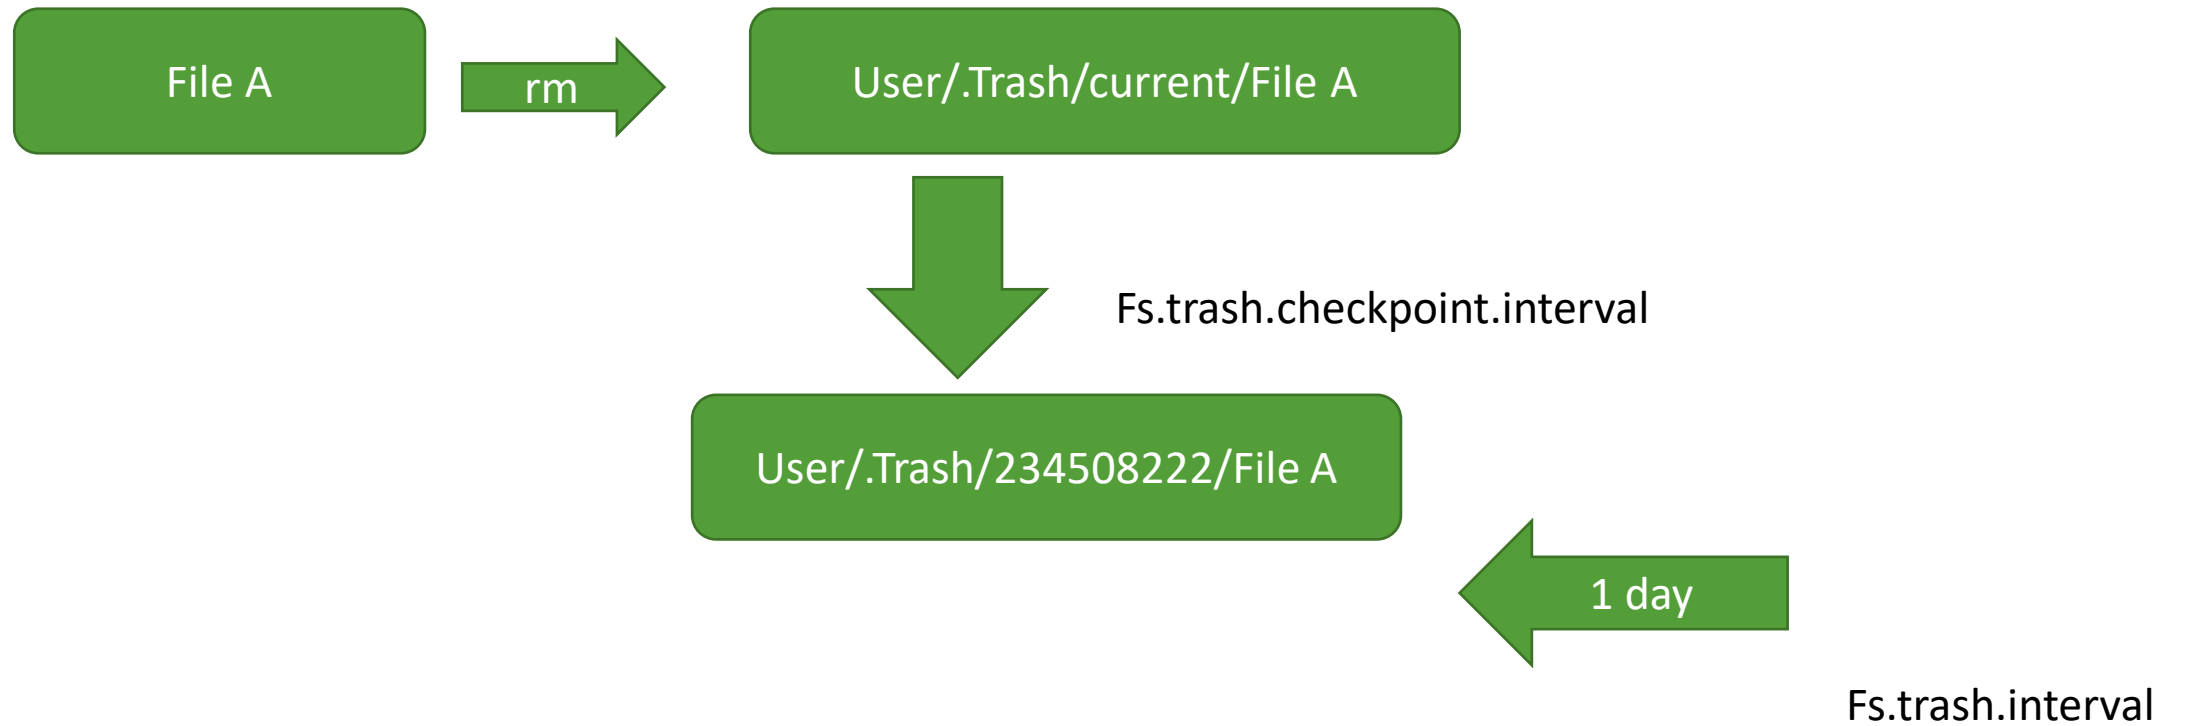
Federation






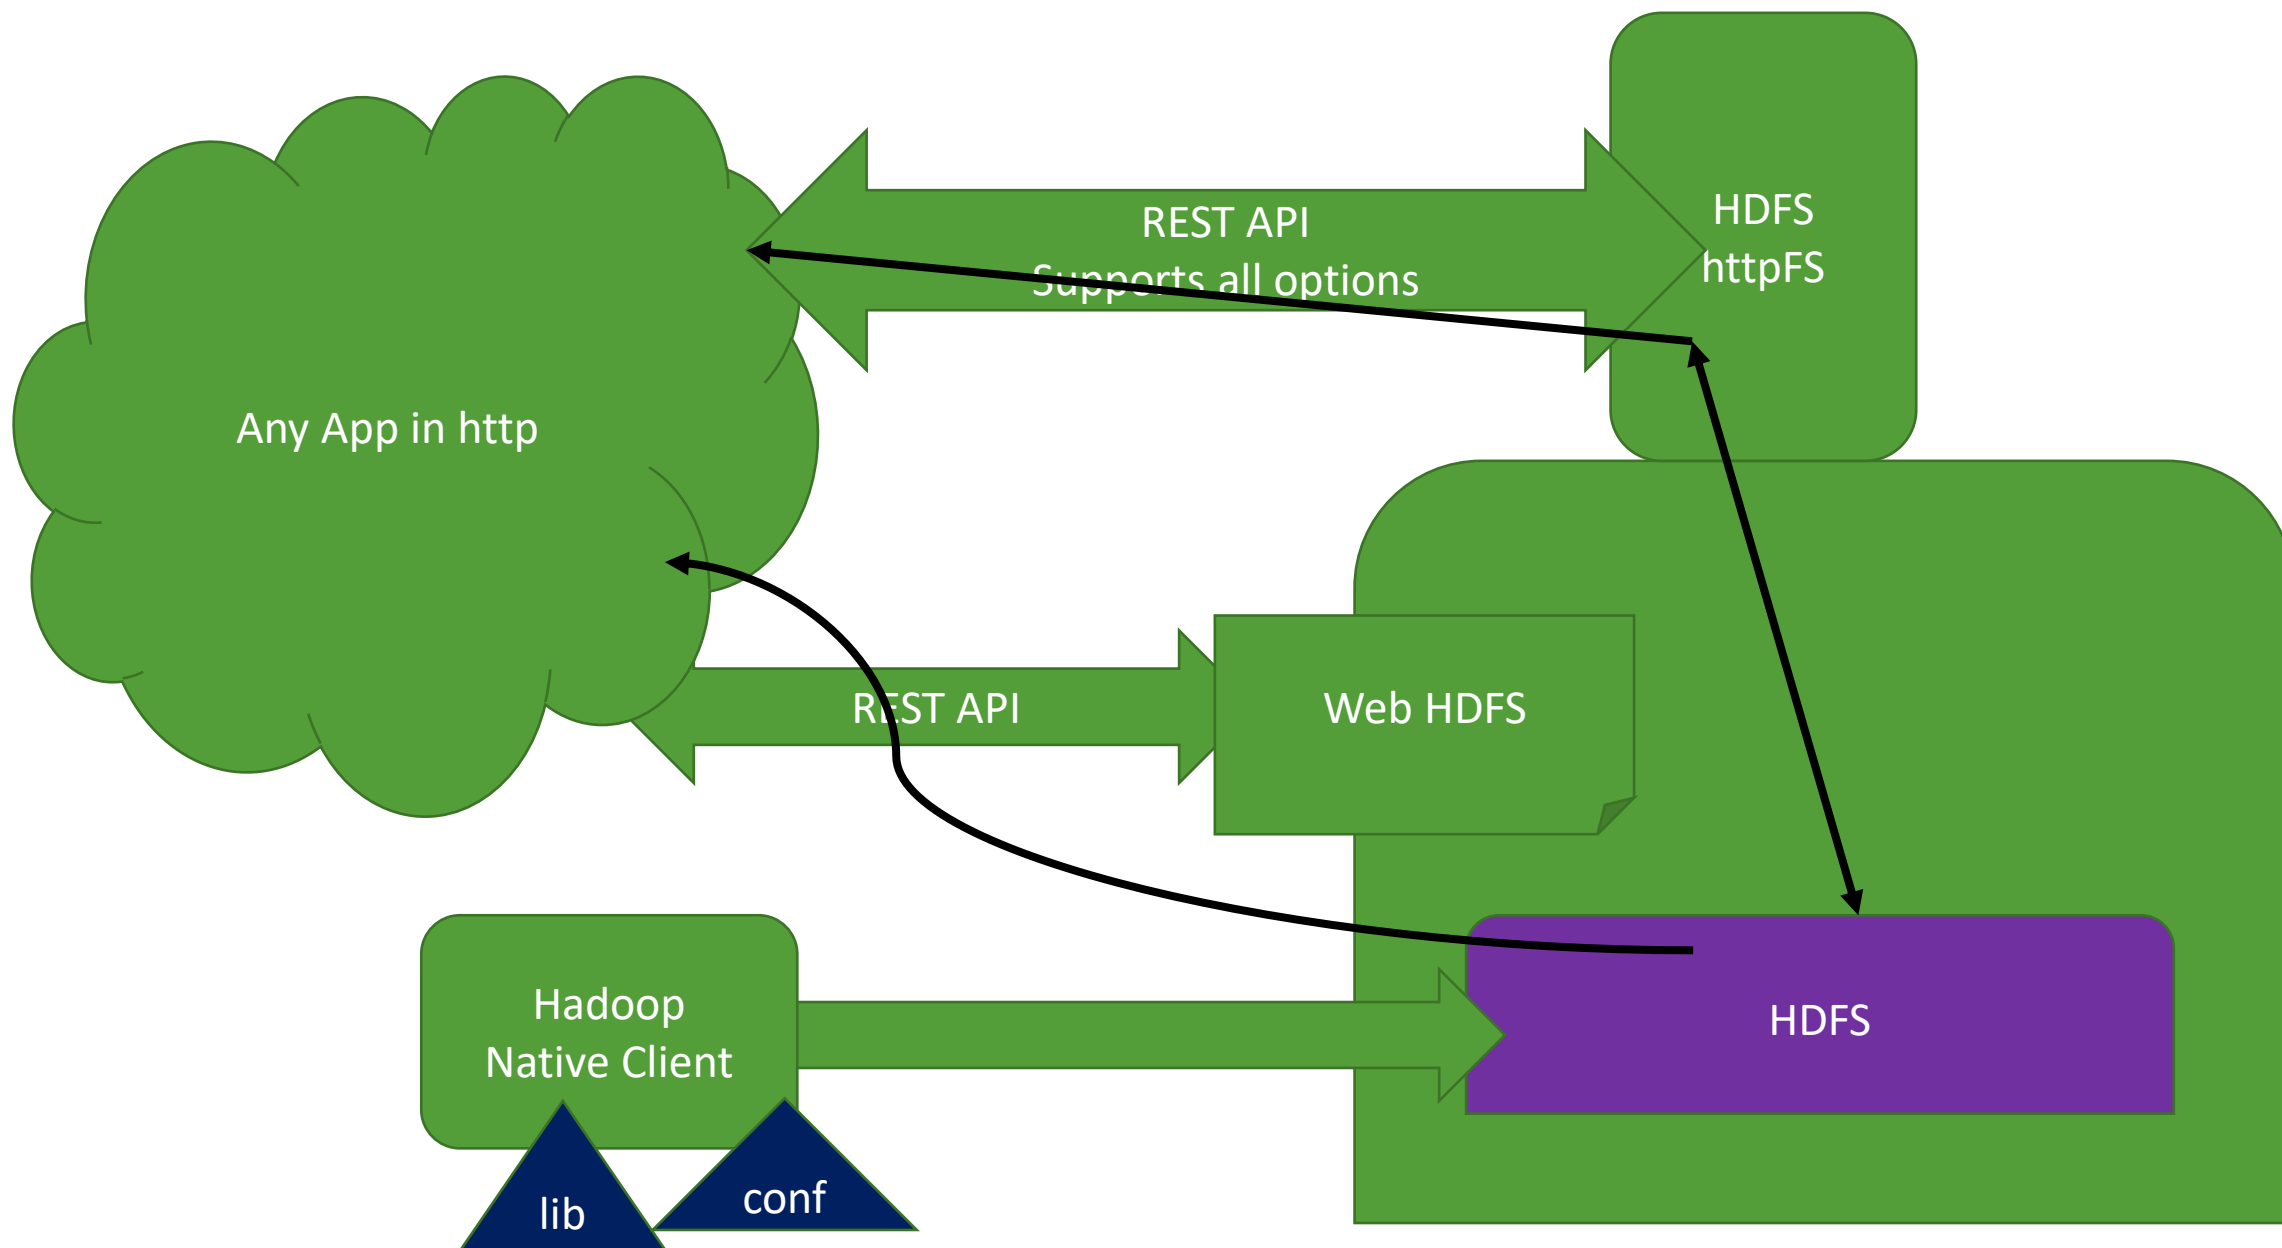
HDFS Trash

- HDFS trash overview
- -skipTrash
- Recover from trash
- Expunge
- Trash interval
- Trash checkpoint interval

Trash Process







Web HDFS

- A Complete HDFS Interface
- HTTP REST API
- Secure Authentication
- Data Locality
- A HDFS Built-in Component
- WebHDFS needs access to all nodes of the cluster
- On read it is transmitted from that node directly

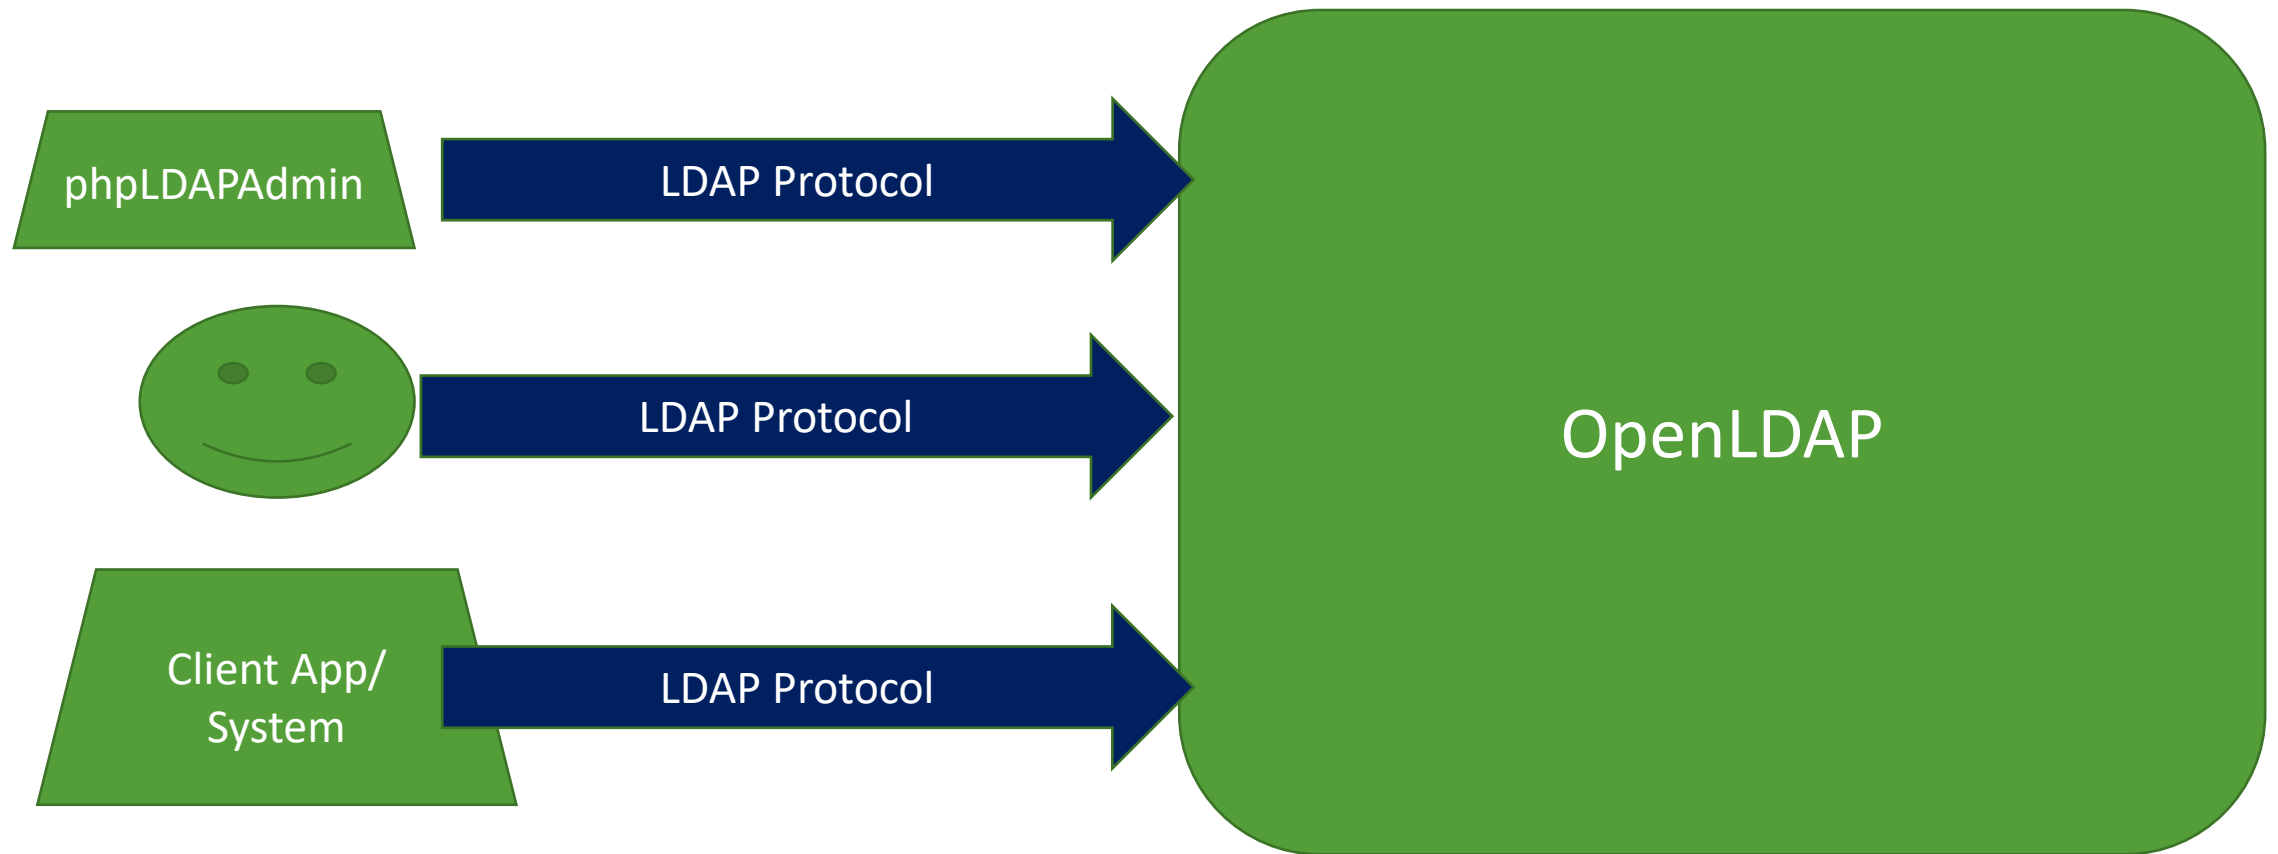
httpFS

- A single node will act similar to a "gateway"
- A single point of data transfer to the client node
- Single Point of Failure
- Could be Over loaded during a large file transfer
- Reduces the footprint required to access HDFS
- Increases the security
- Preconfigured Tomcat bundled with HttpFS

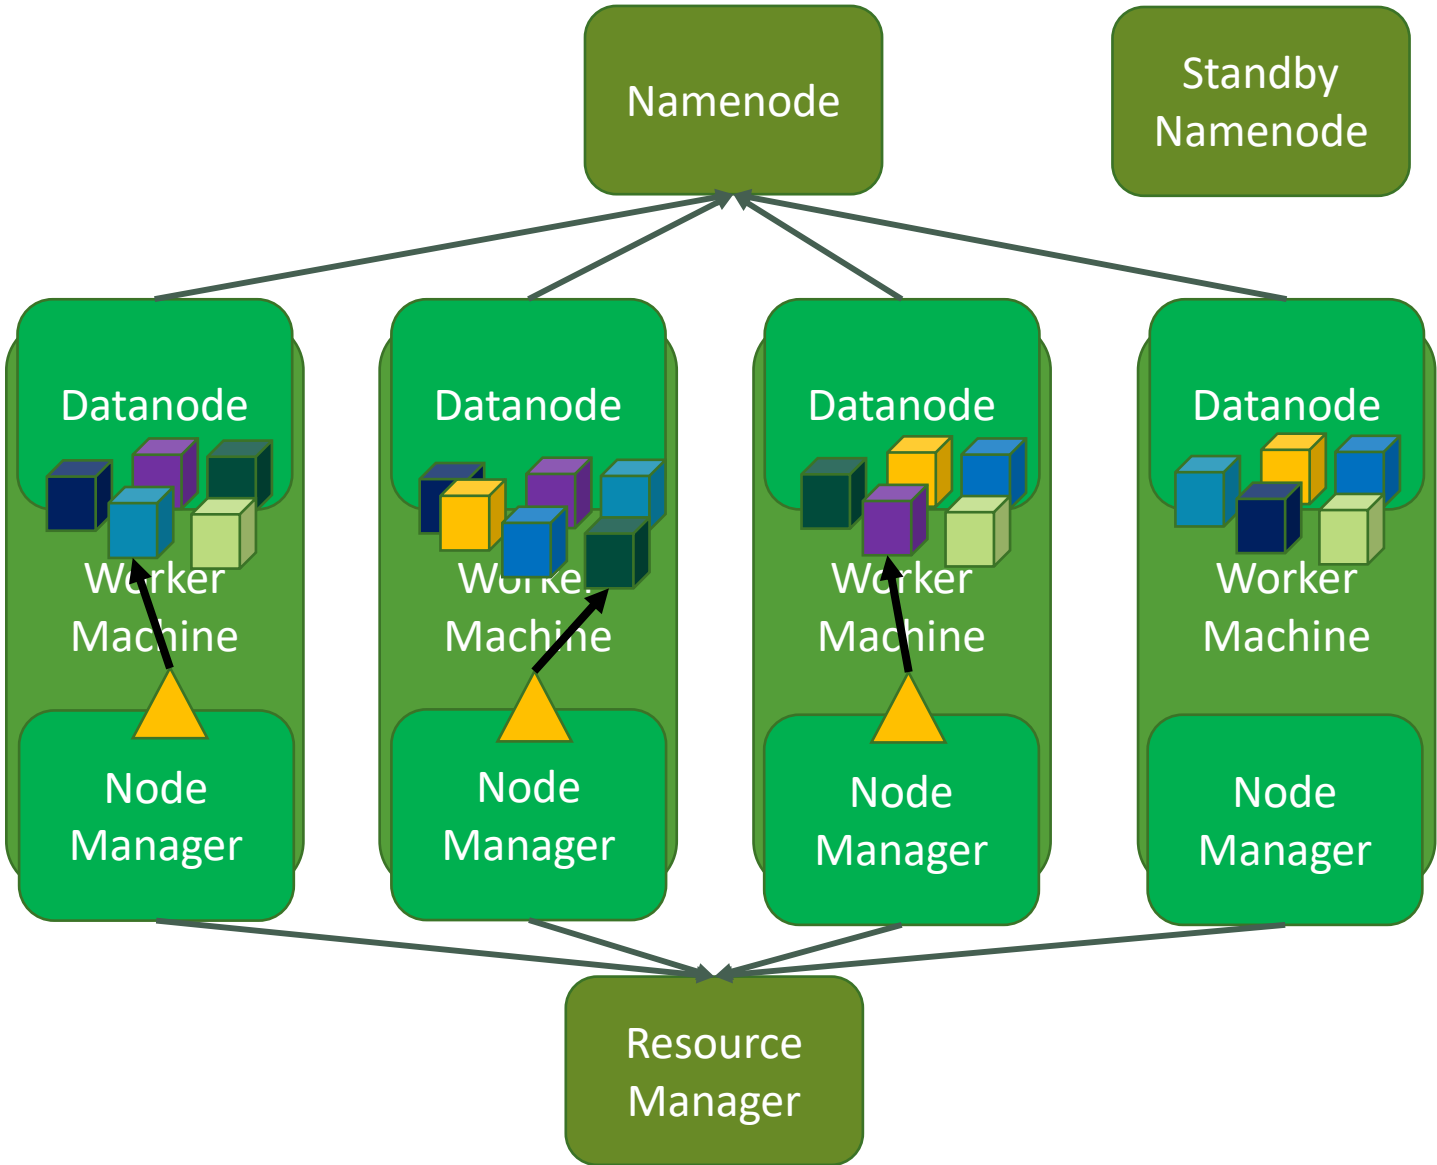
OpenLDAP

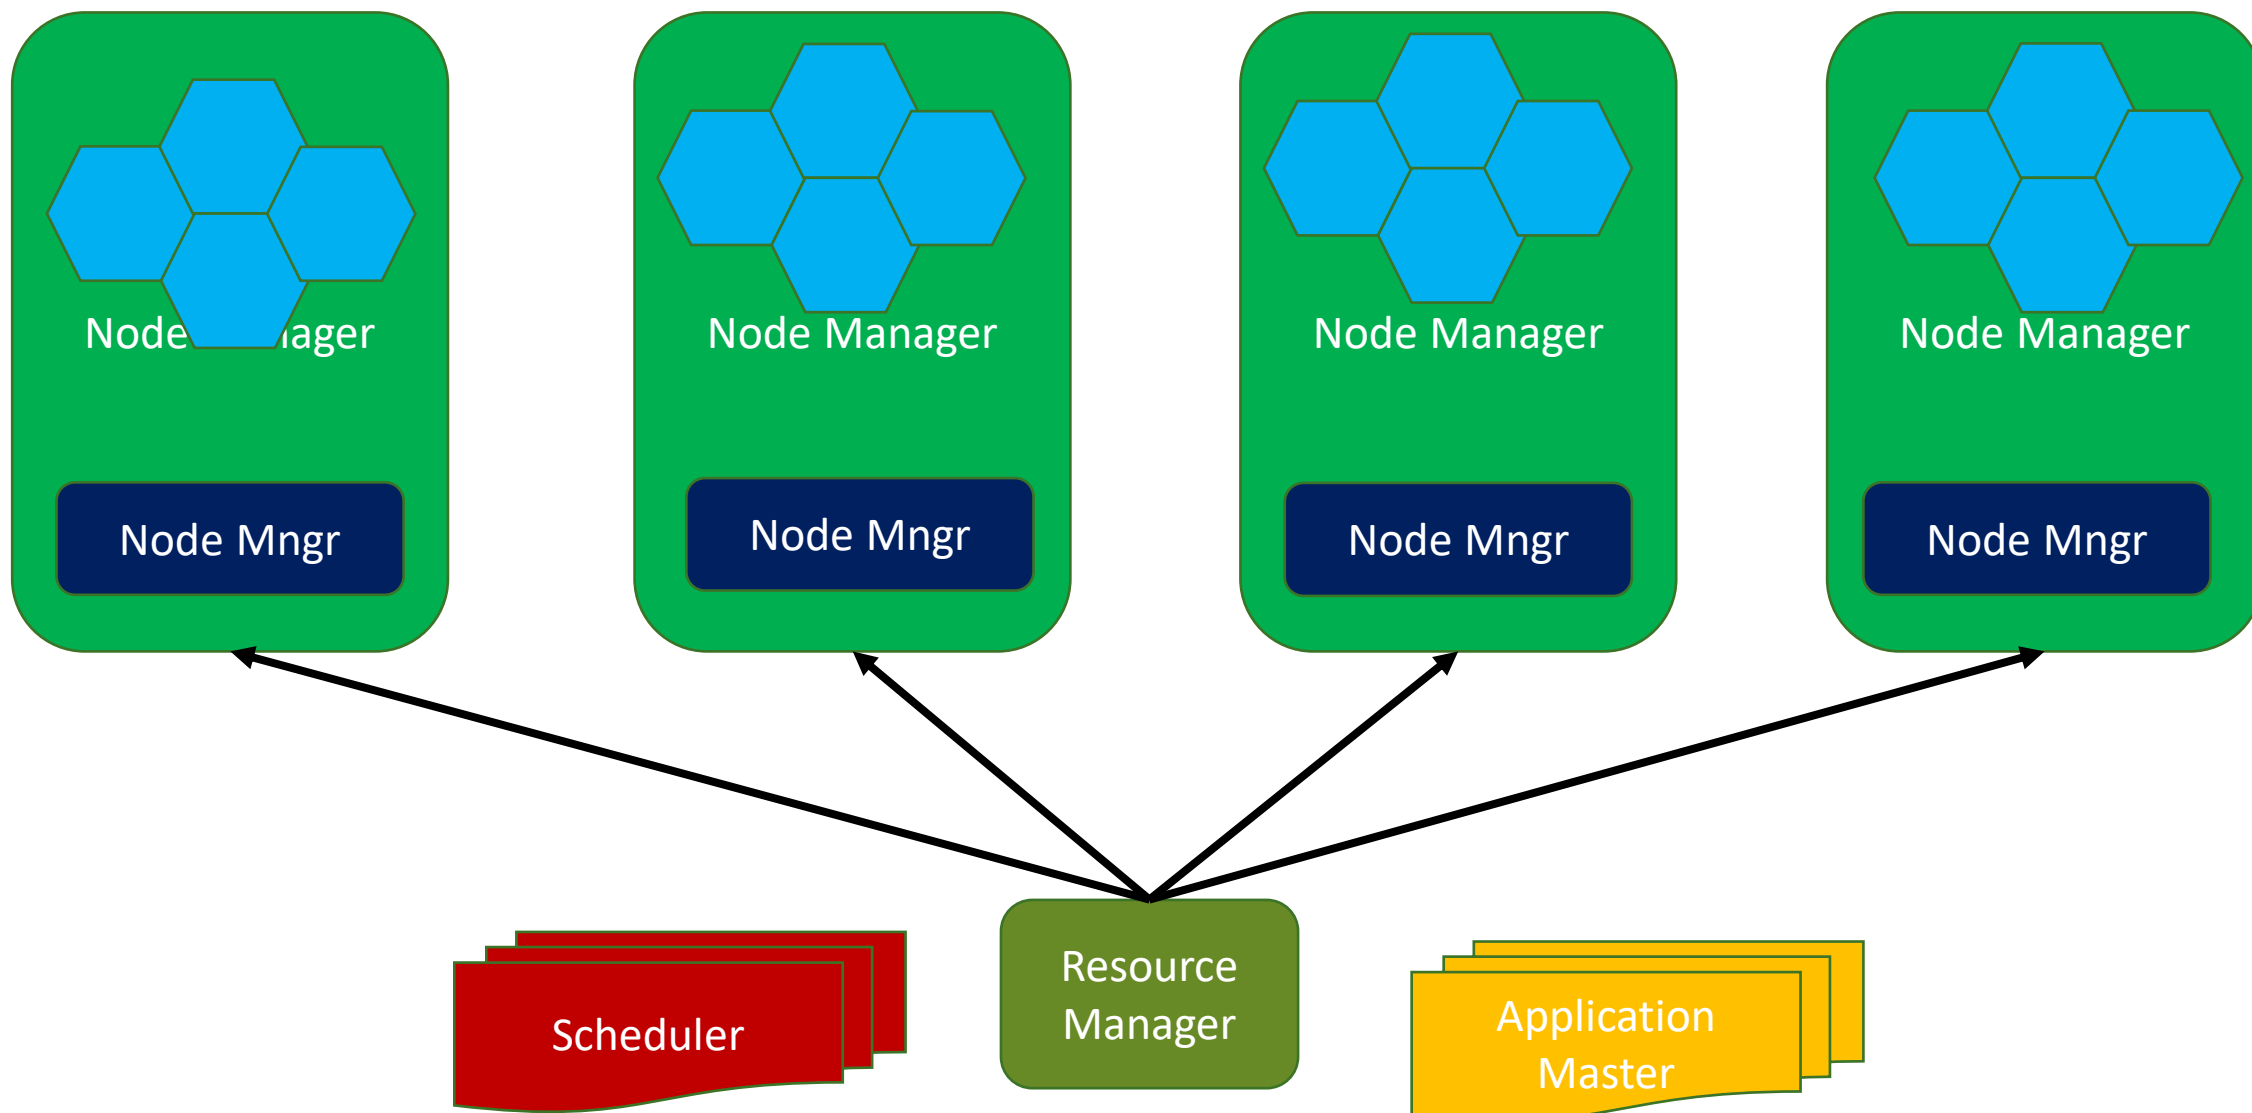
- Open-source and directory server
- organizational information stored on the server which can be used by client to search
- access, modify and manipulate entries in the directory.
- Client applications connect to OpenLDAP server using the Lightweight Directory Access Protocol (LDAP)
- User permissions can be controlled

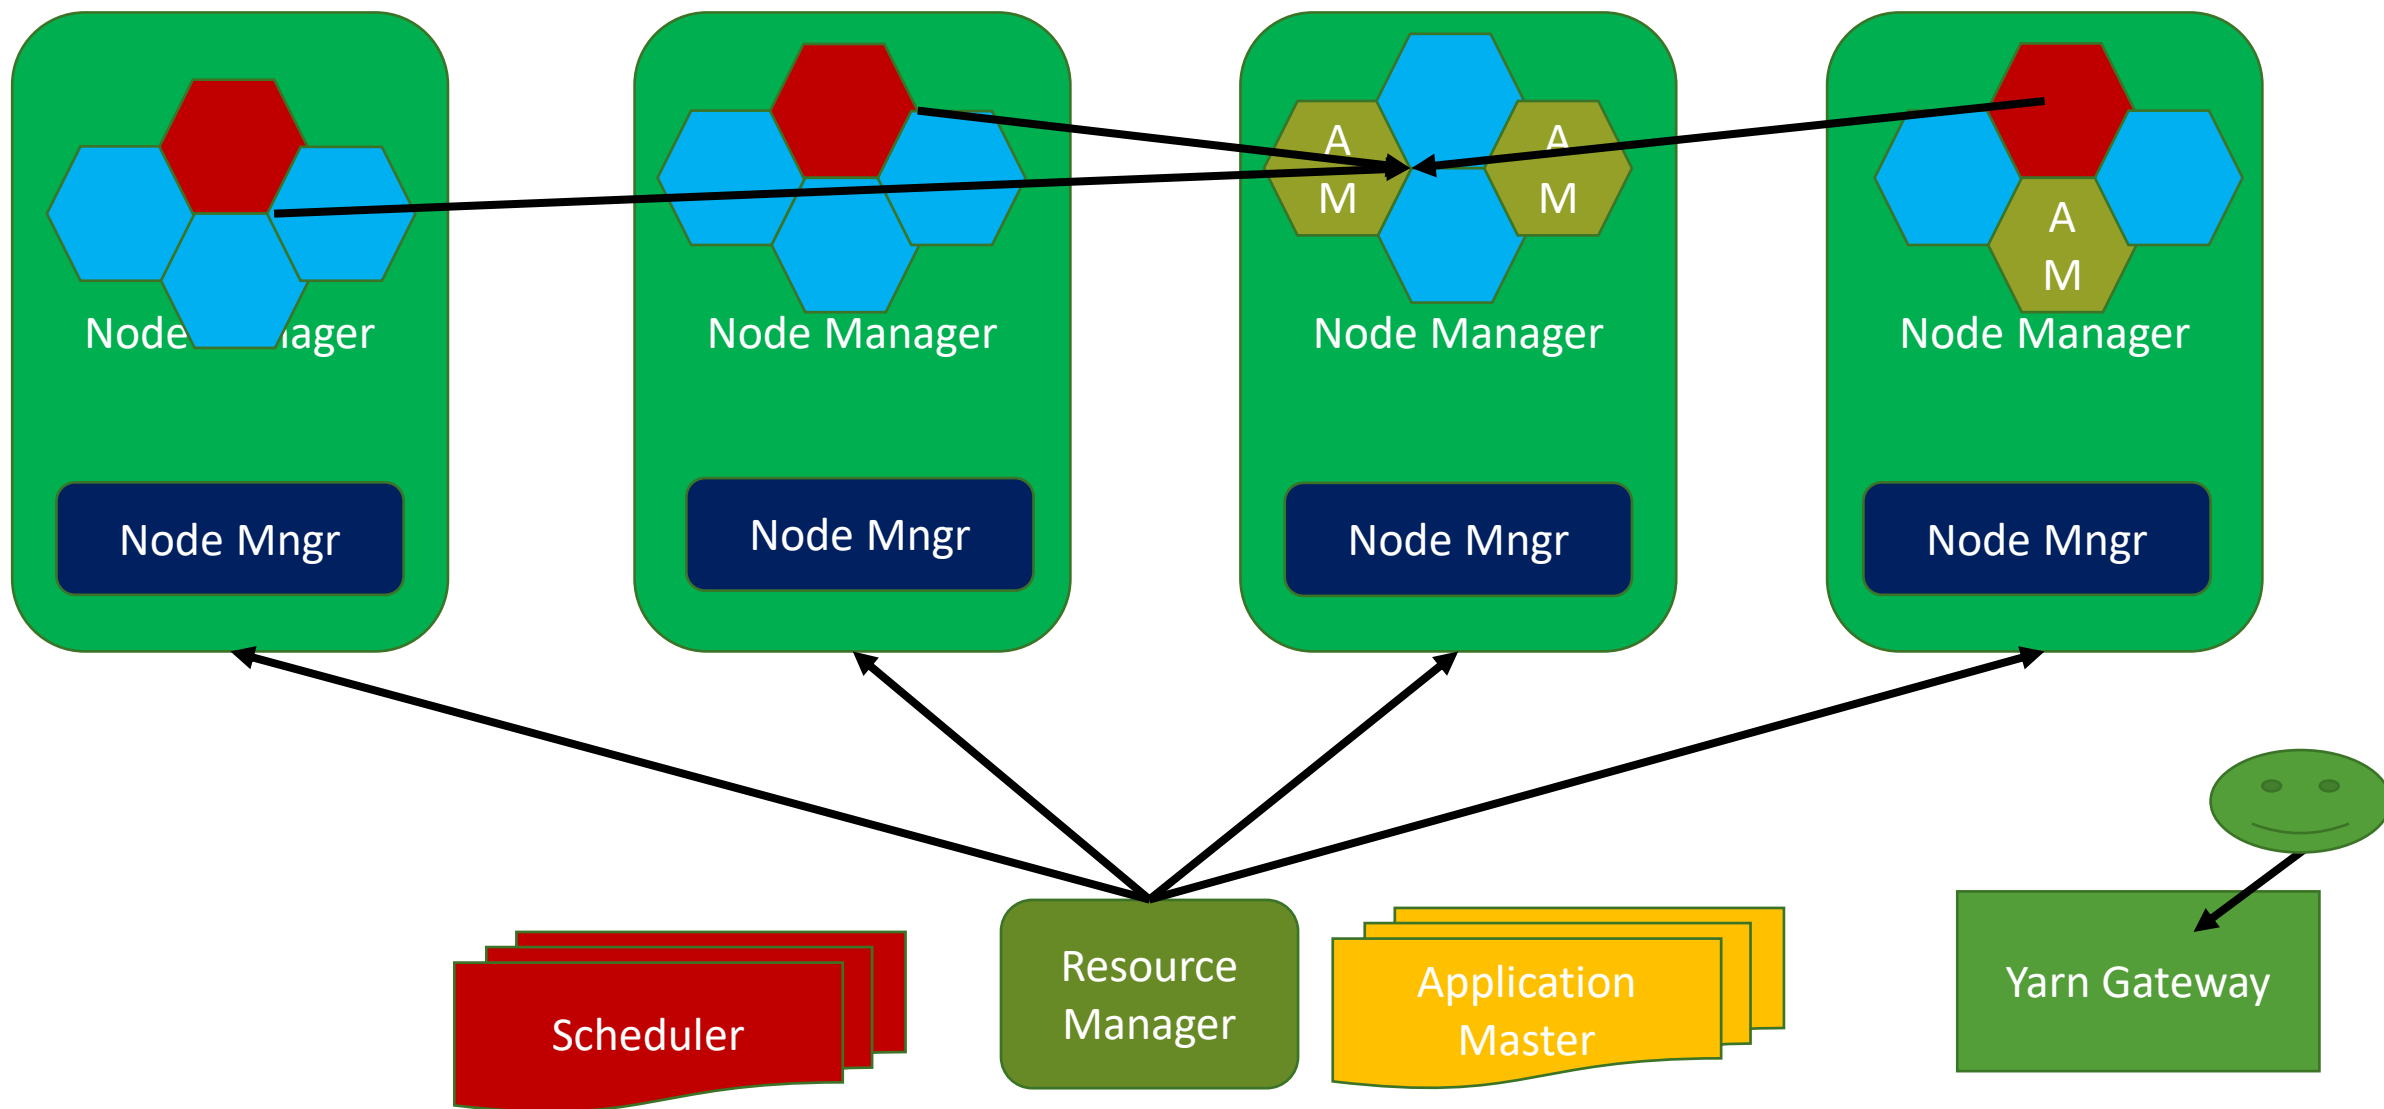
OpenLDAP

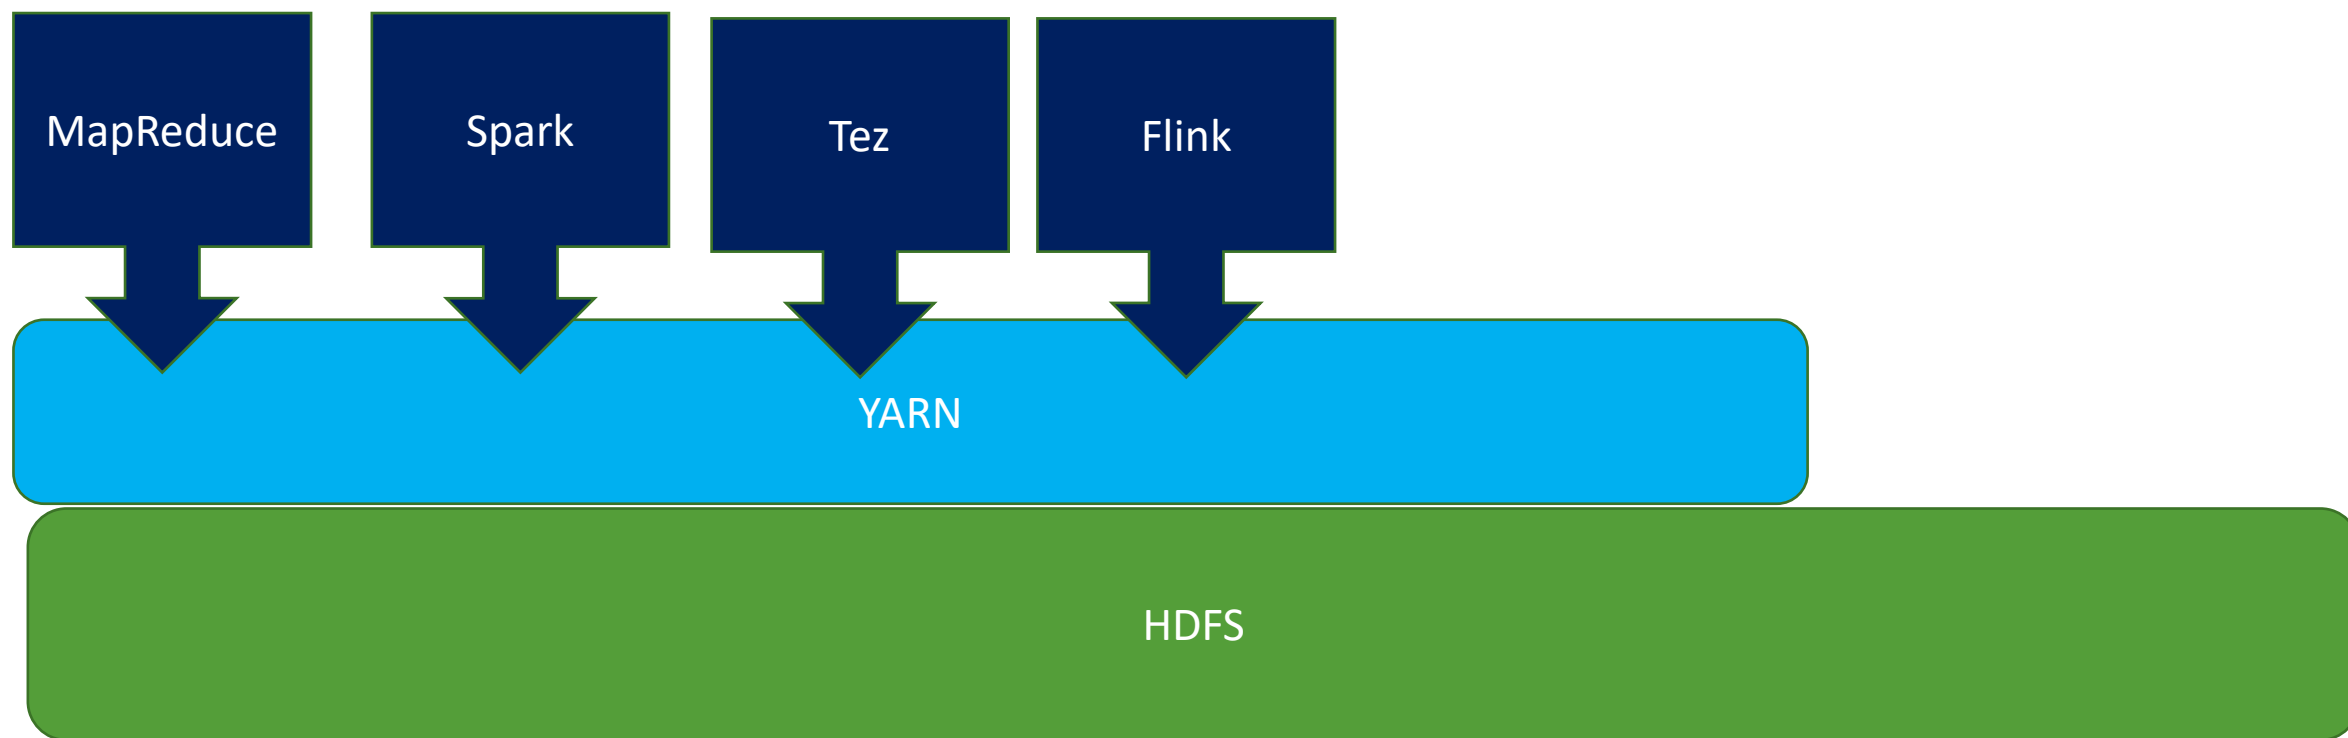



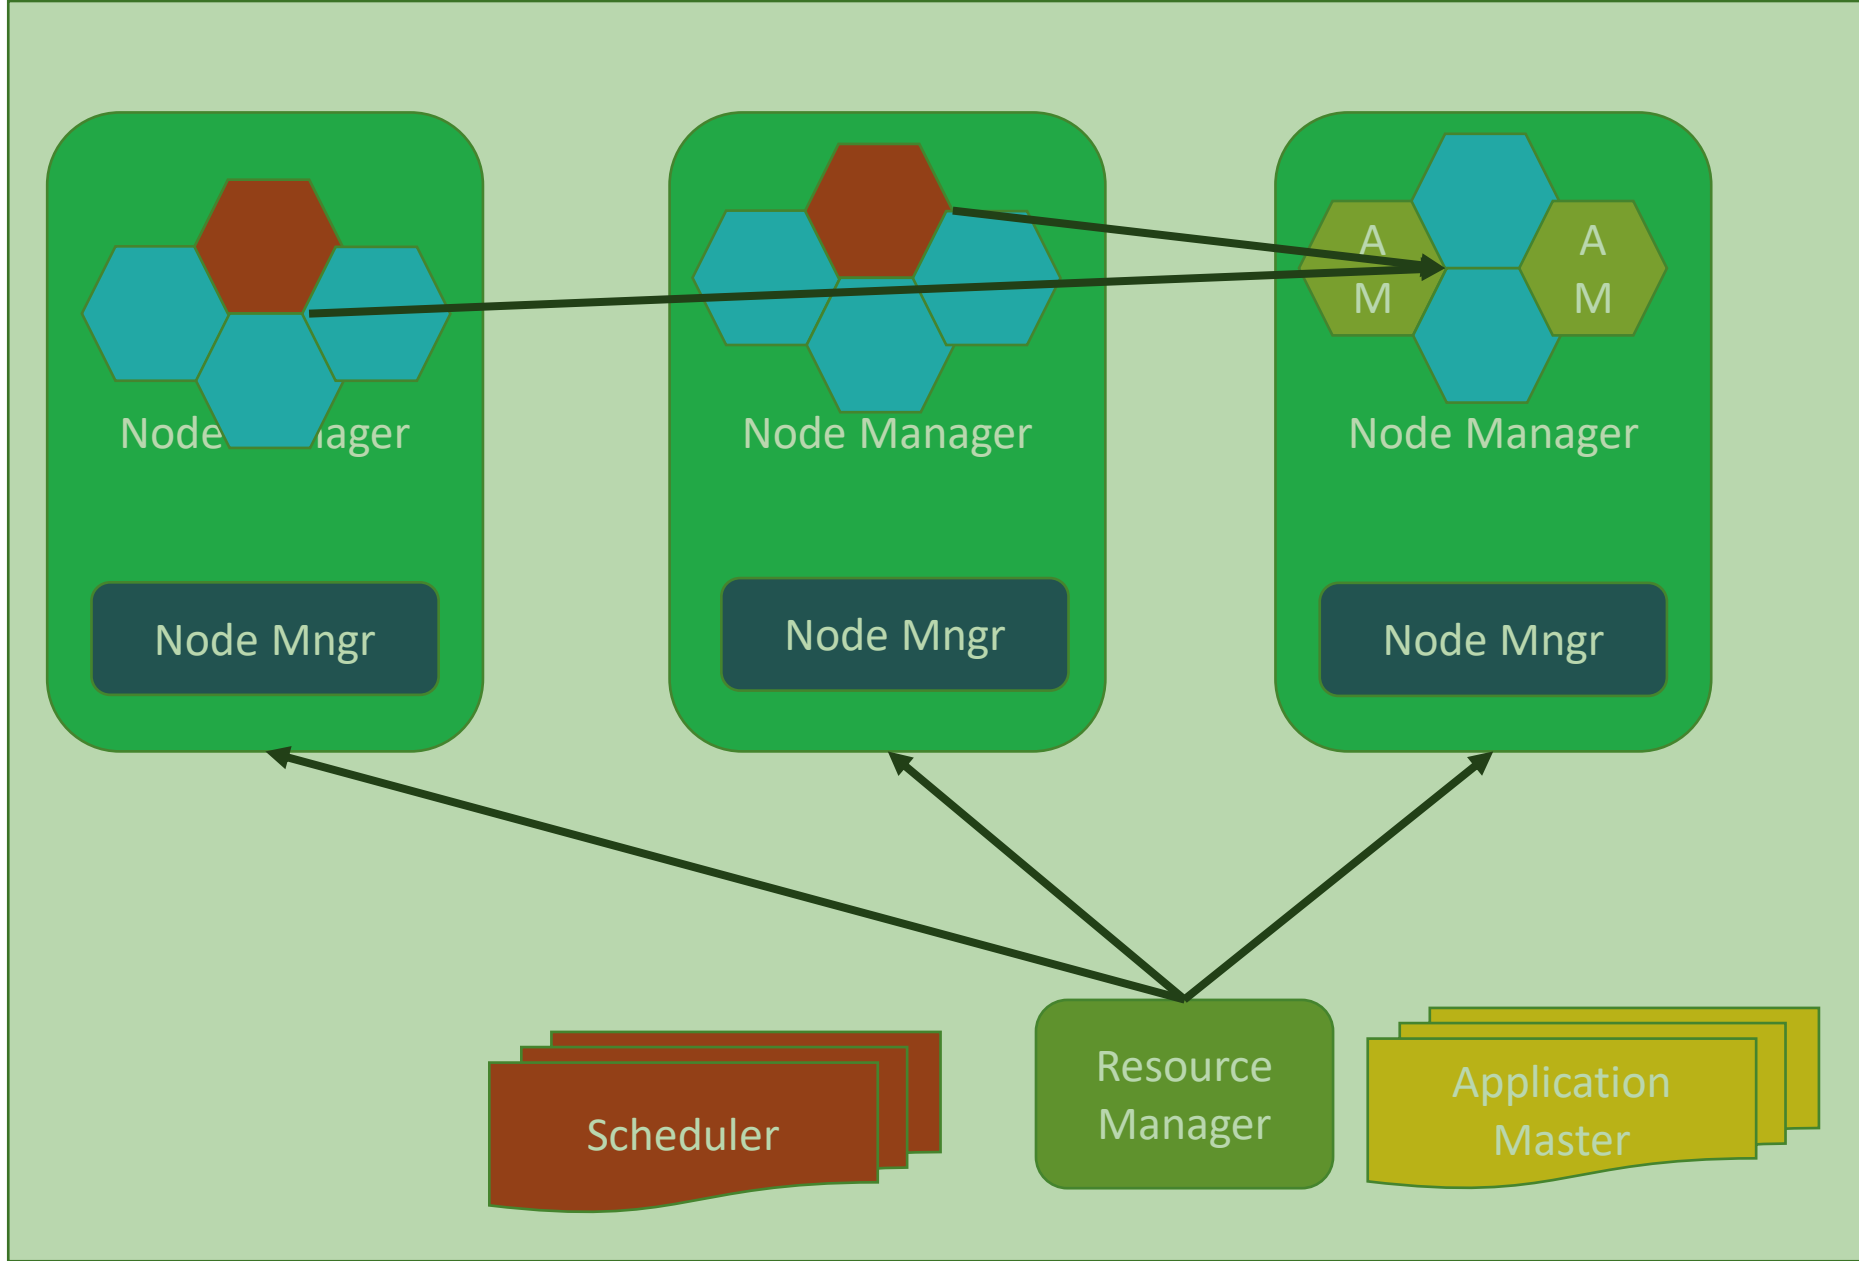
HDFS and YARN












Resource Management

- Static Service Pools
- Dynamic Resource Pools
- Yarn (MRv2 and MapReduce (MRv1) Schedulers
- Resource Management for Impala (*Specific to Cloudera)

Static Service Pool




Schedulers

- FIFO Schedulers (Apache Hadoop)
- Capacity Schedulers (Hortonworks)
- Fair Share Schedulers (Cloudera Default)

FIFO Scheduler

- Simple and original scheduling algorithm from MRv1
- Initial Scheduler by Apache Hadoop
- Pulls the oldest jobs first from the queue
- Had no concept of the priority or size of the job
- Suitable for small size cluster


FIFO Scheduler



Fair Share Scheduler

- Jobs get average, equal share of resources over time.
- When single job running in entire cluster, it consumes all resources.
- Any new job(s) submitted, the free up container will allocated.
- Pools and Priority can be set to Jobs.
- If minimum share is not met for some period of time, the scheduler optionally supports *preemption* (Terminating any running jobs)
- Configurations in yarn-site.xml and fair-scheduler.xml

Fair Share




Yarn-site.xml – For every container request

- Yarn.scheduler.minimum-allocation-mb - 1024
 - Yarn.scheduler.maximum-allocation-mb - 8192
 - Yarn.scheduler.minimum-allocation-vcores - 1
 - Yarn.scheduler.maximum-allocation-vcores – 32
-
- yarn.scheduler.capacity.<queue-path>.maximum-capacity
 - yarn.scheduler.capacity.<queue-path>.capacity

Capacity Scheduler

- Elasticity over queues to share capacity
- Access control
- Max capacity per queue to enforce capacity
- User limits within queue
- Queues with priority

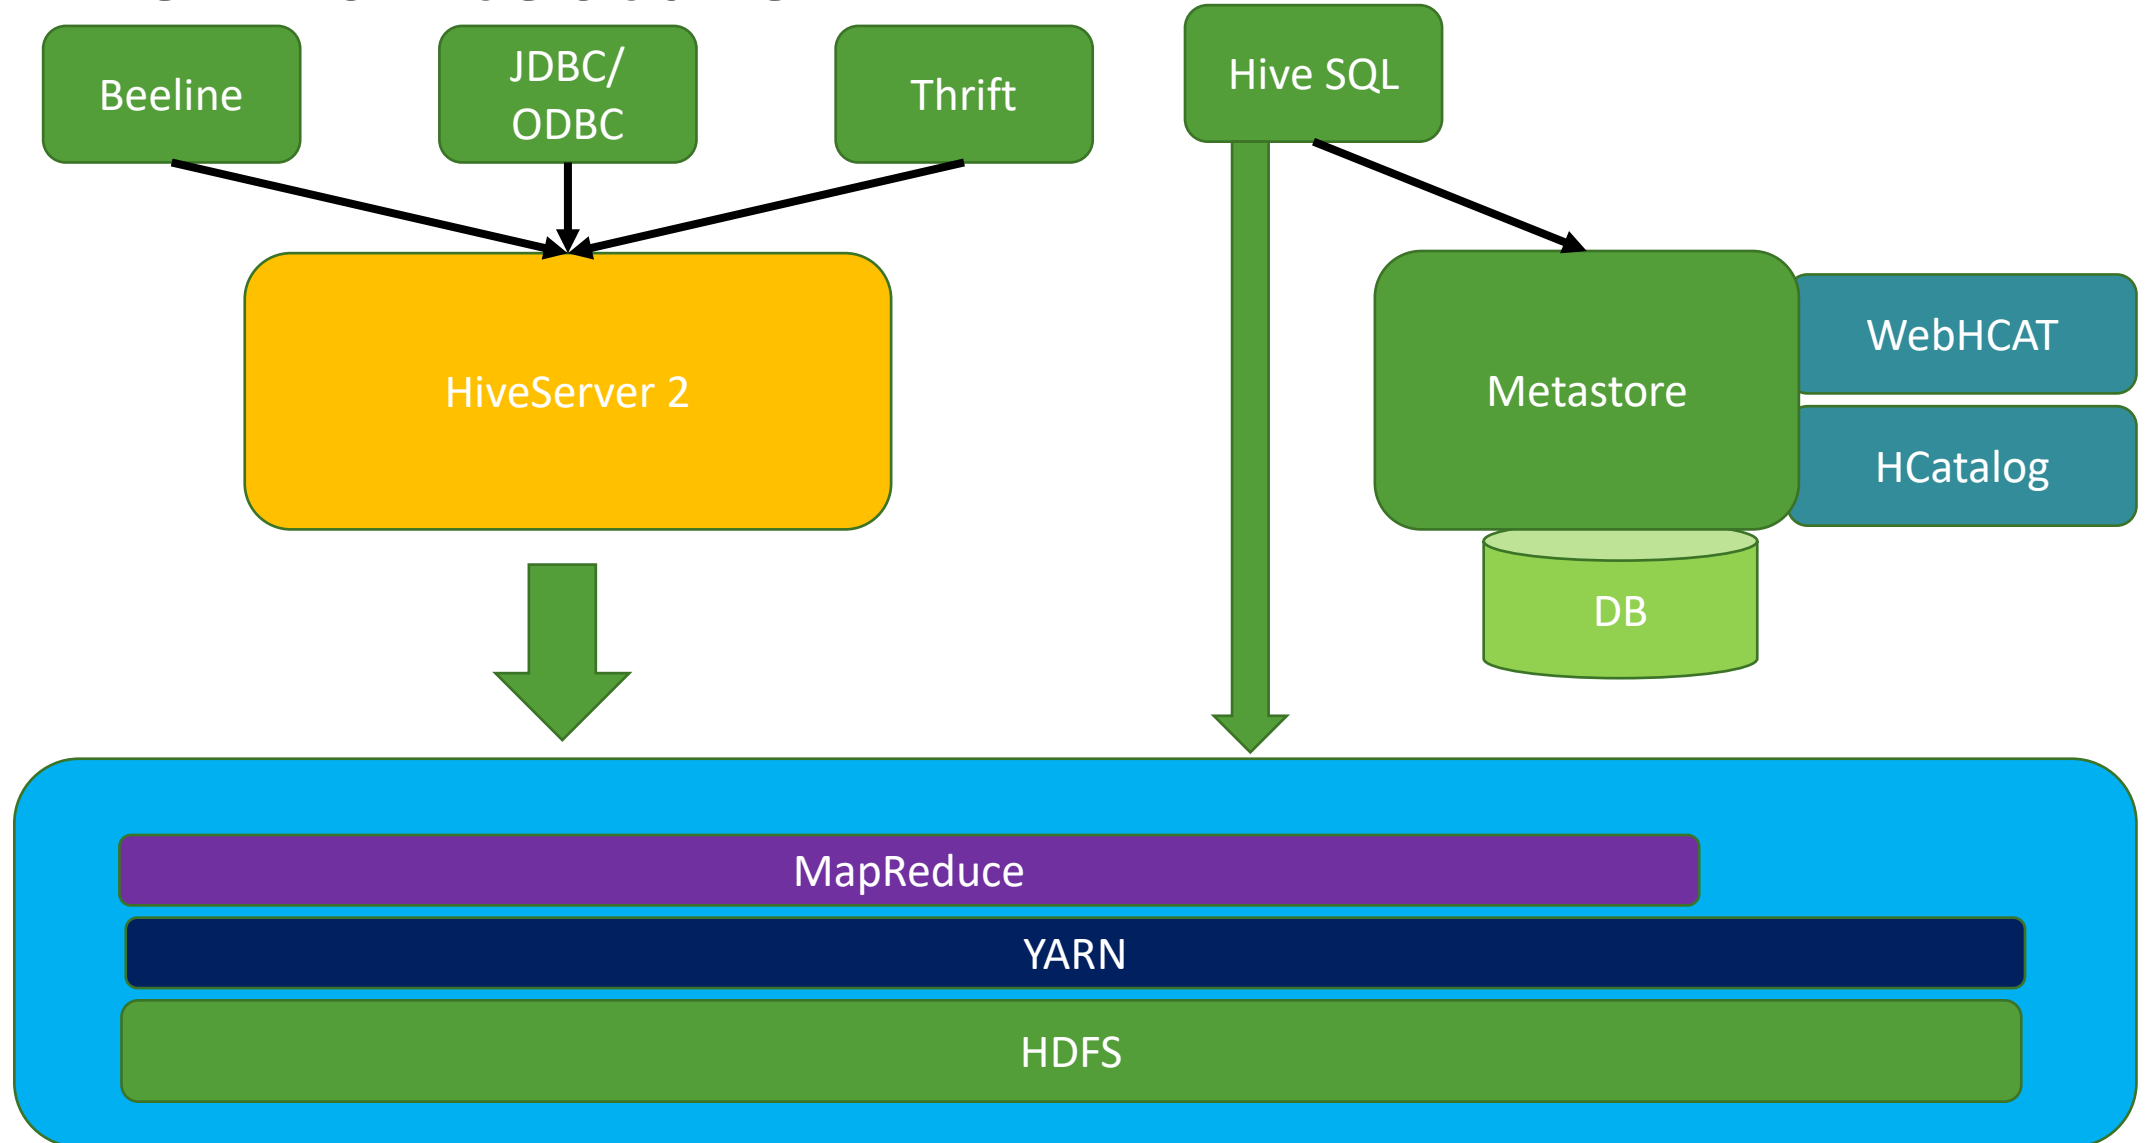
Capacity Scheduler



Hive

- Petabyte-scale data warehouse offering
- Enables reading, writing and managing large datasets in distributed environment
- Uses HiveQL, similar to SQL
- HiveQL transformed into sequence of MR jobs and executed on Hadoop Cluster through MapReduce and Spark
- Works with structured data only

Hive Architecture



Metastore database

- Holds the metadata about databases, tables, columns, partitions and Hadoop specific details like files and block location in HDFS
- Shared by other components like Impala, Pig, Hcatalog
- Runs on its own JVM
- Metadata gets stored in a dedicated RDBMS

HiveServer2

- Server interface which facilitates remote clients to submit queries
- Supports multi-client concurrency, capacity planning, security, authorization, etc.,
- Container for Hive execution engine.
- Supports JDBC clients such as Beeline CLI, ODBC clients.
- Client connection is established with Thrift API based Hive service.
- Can leverage and use Spark and HBase

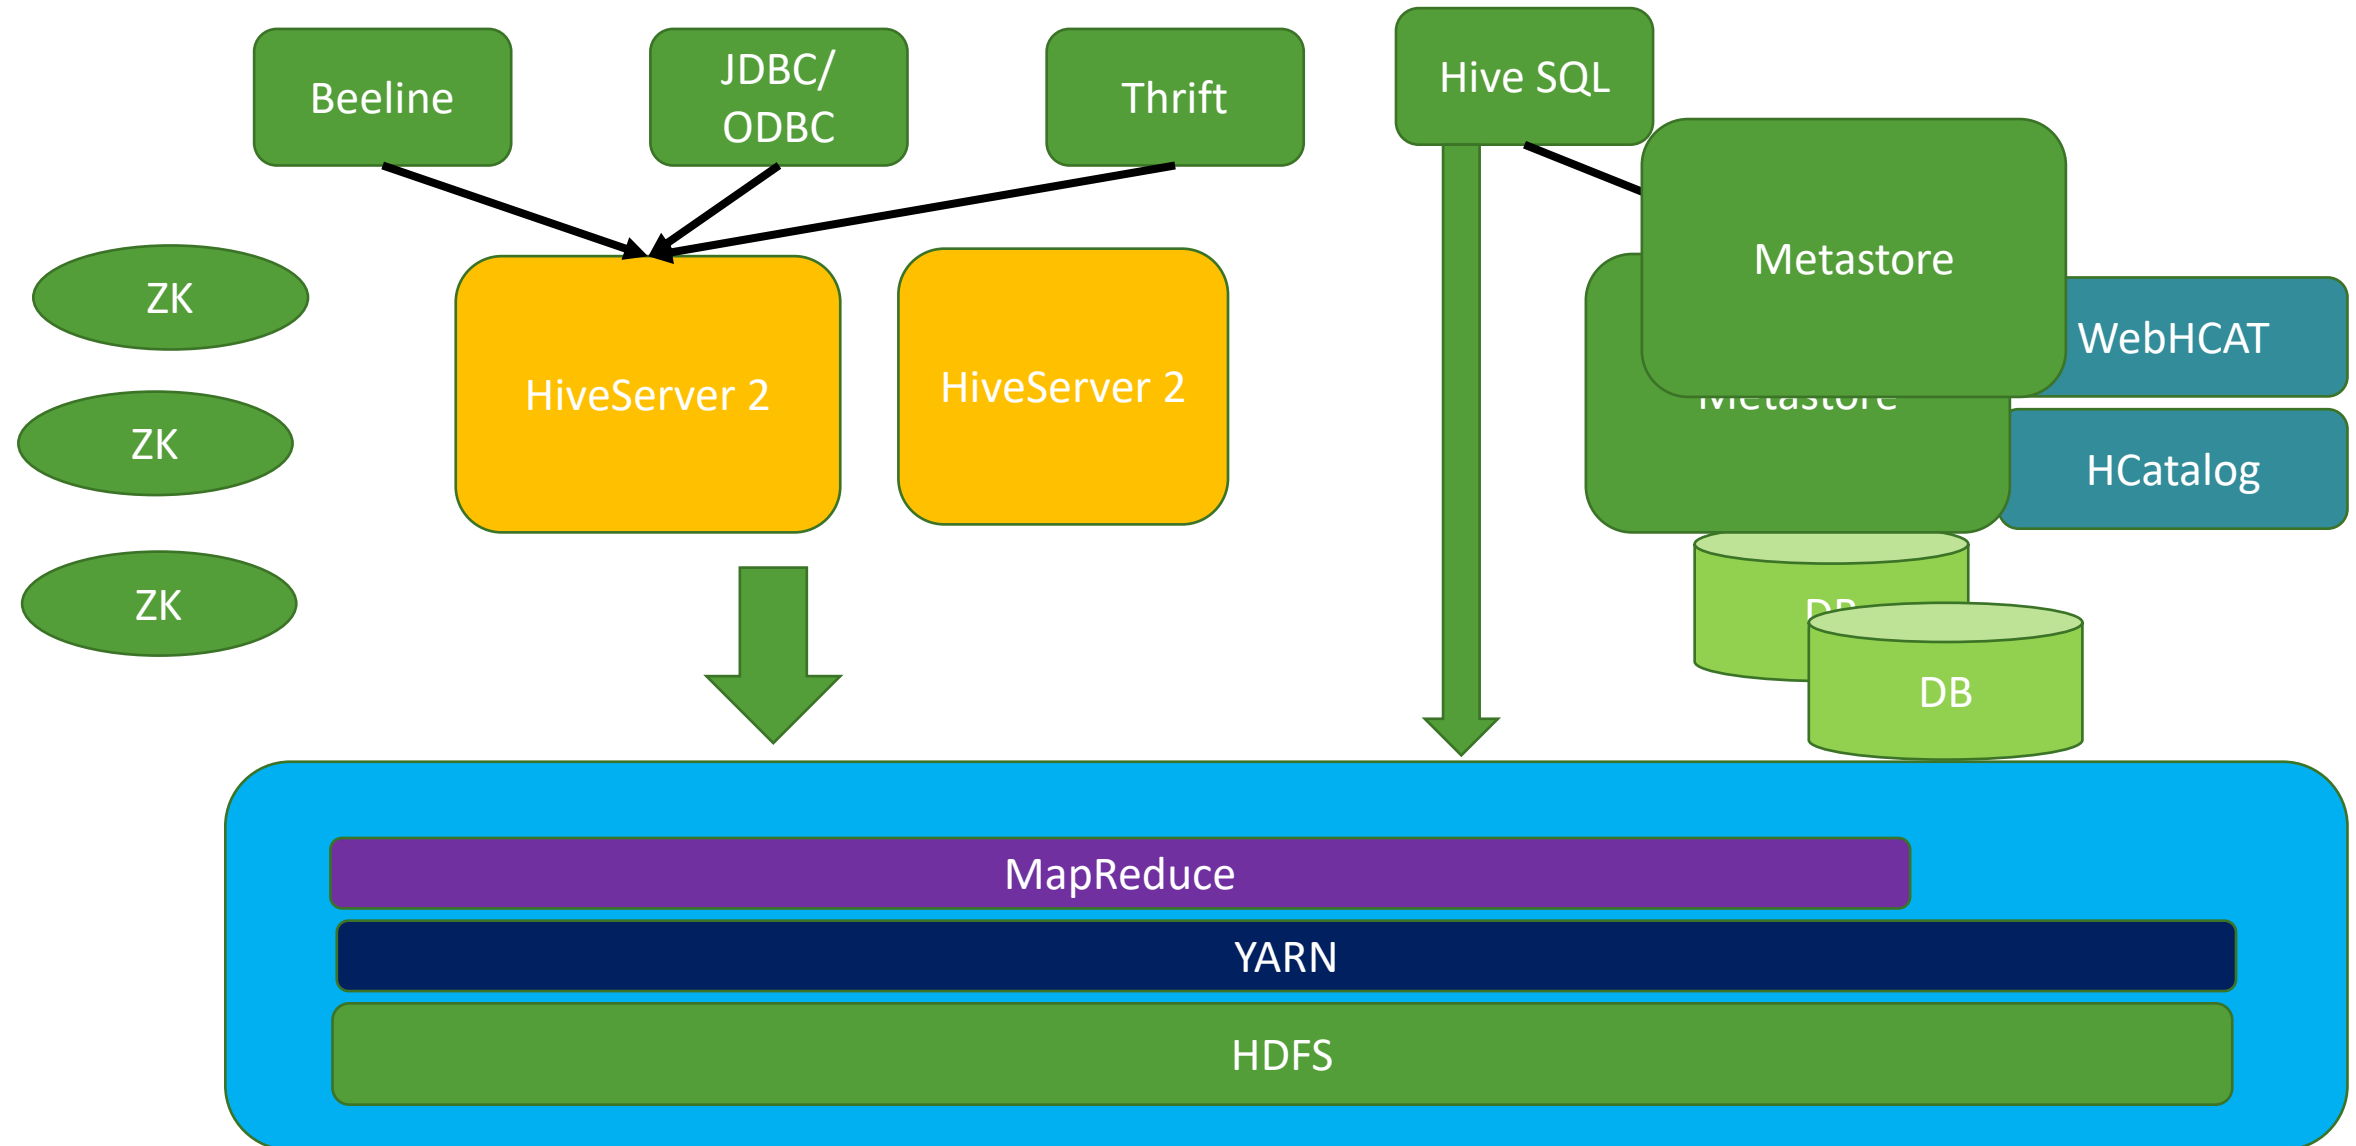
Hue (Hadoop User Experience)

- Open source web based interactive query editor
- Developed by Cloudera
- Browser, Editor, Workflow Editor, Shell, Solr Searches

Oozie

- Scheduler system for Hadoop Jobs from Apache
- Can create and manage collection dependency jobs as workflow
- Scheduler can trigger the workflow at regular interval
- Workflows are managed as DAG cycle

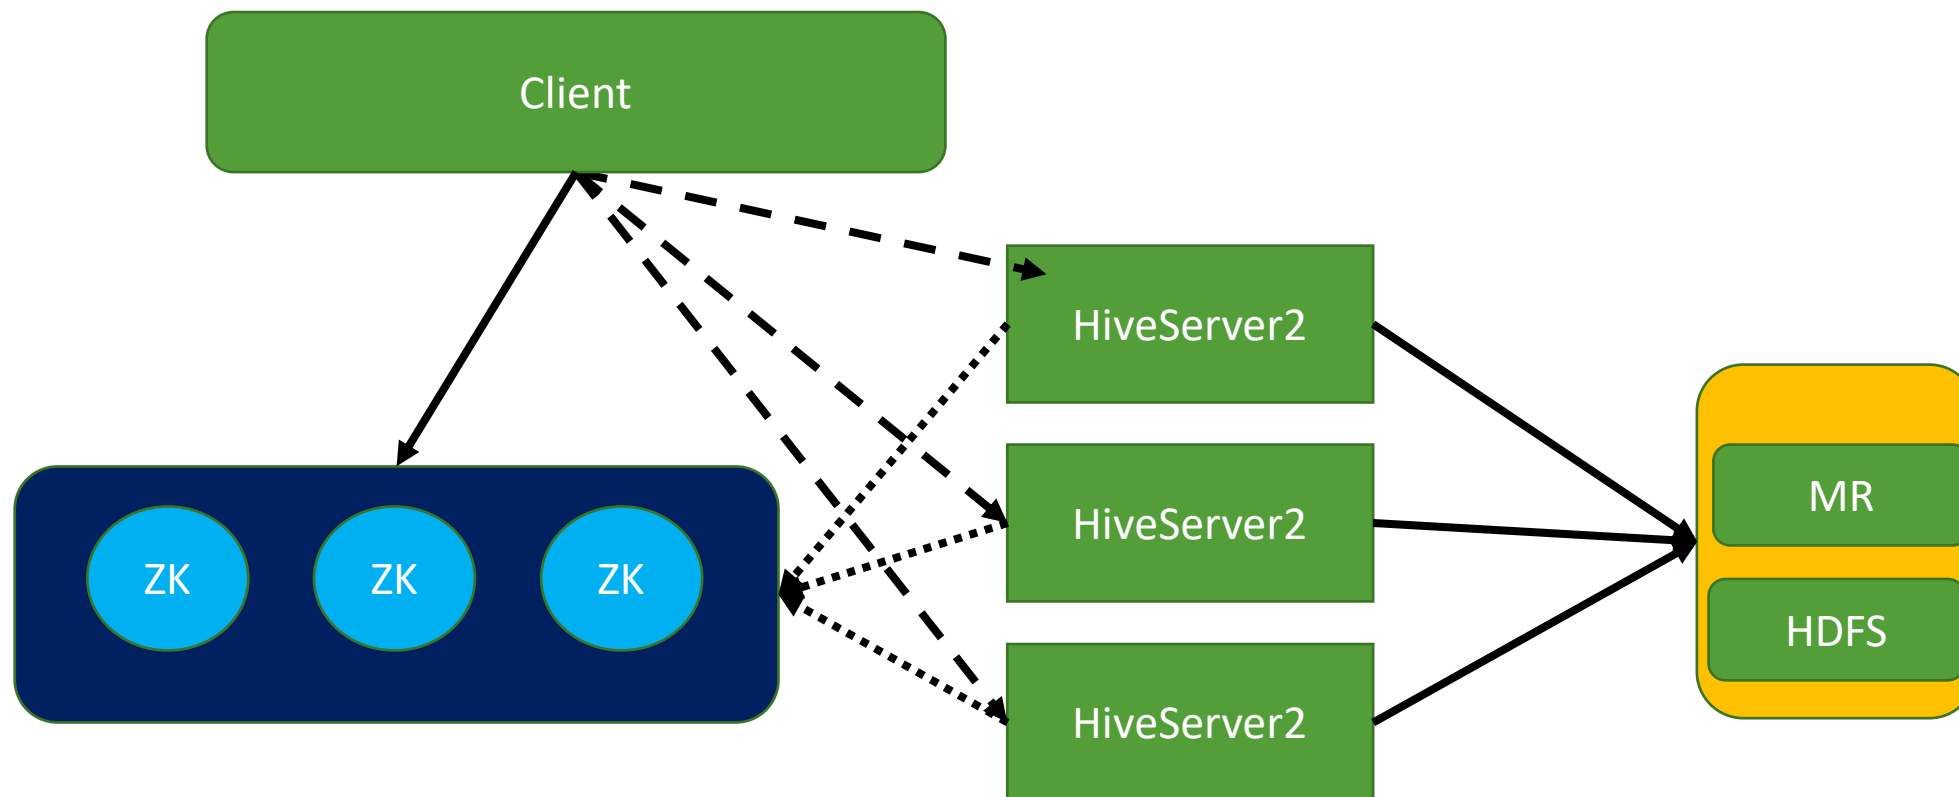
Hive High Availability



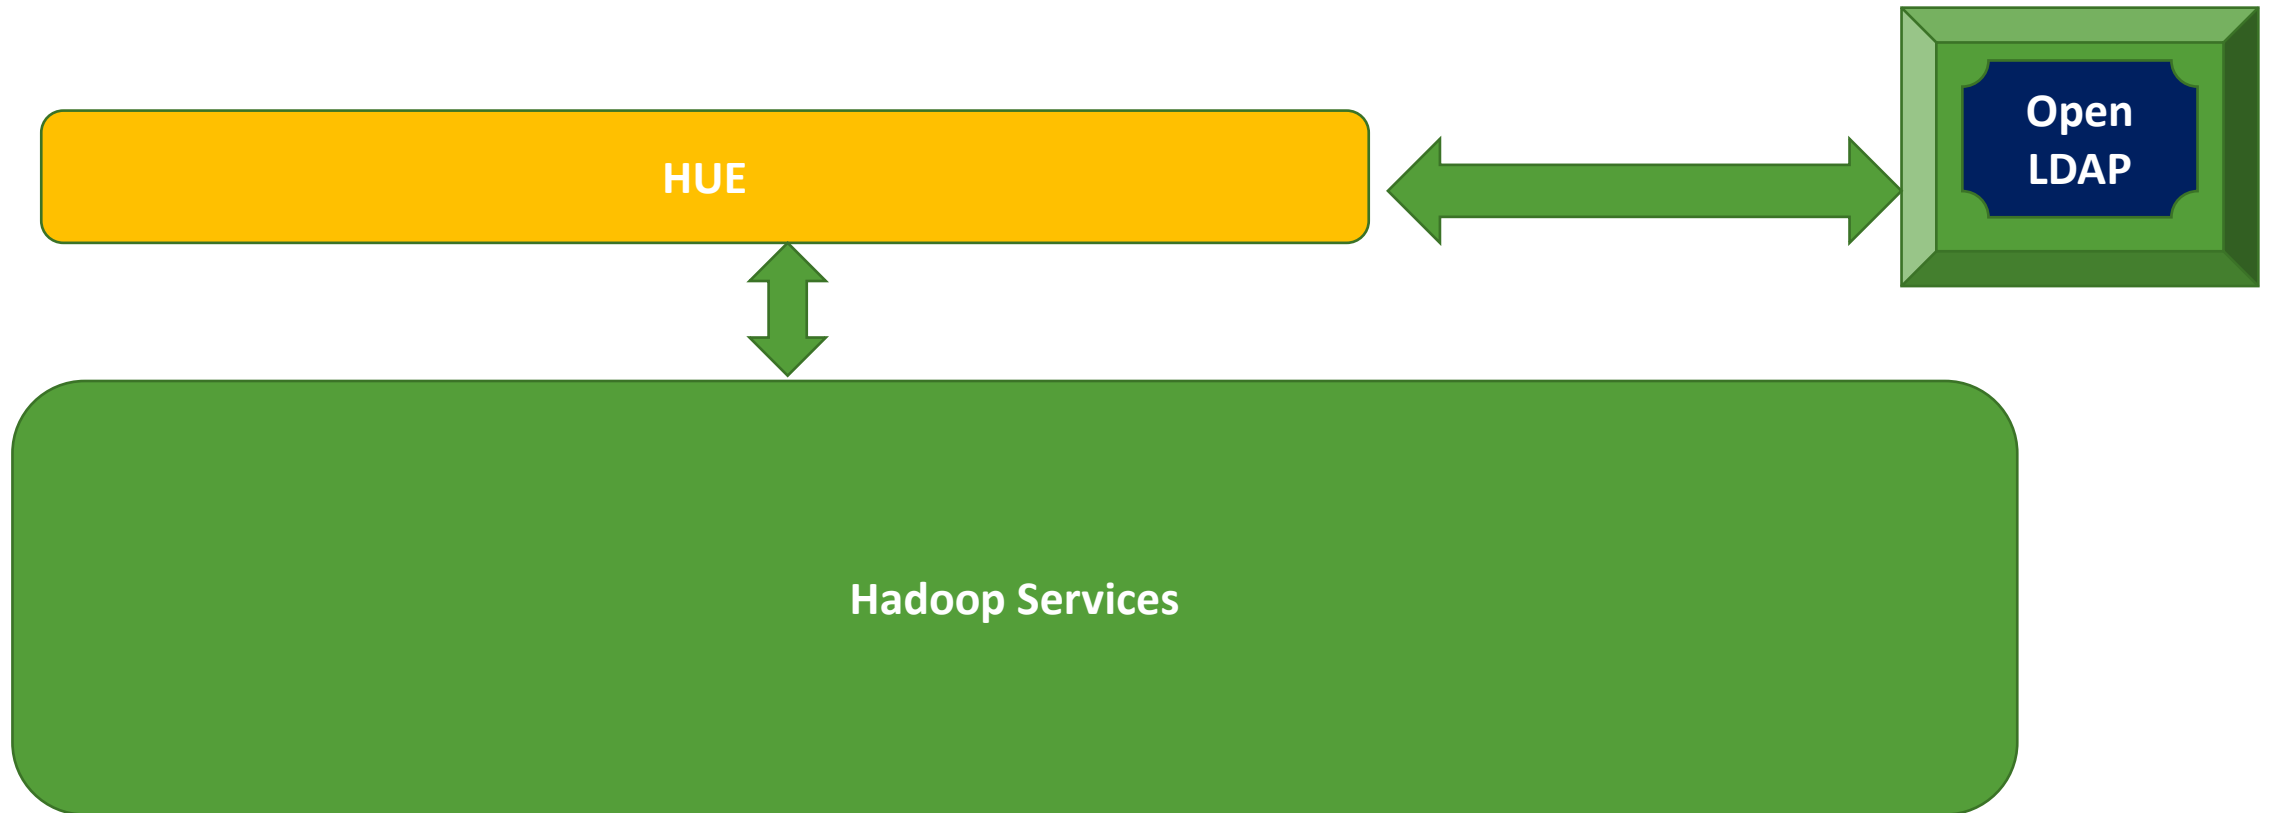
Point of Failure

- RDBMS holding the metastore
- **Hive Metaserver**
- **HiveServer2**
- Zookeeper to coordinate high availability

HiveServer2 HA



HUE Open LDAP Configuration



HUE Configuration Property

- Authentication Backend -> desktop.auth.backend.LdapBackend
- ldap_url -> ldap://<<hostname>>:389
- ldap_username_pattern -> uid=<username>,**ou=users,dc=muthu4all,dc=com**
- search_bind_authentication -> Select (True)
- use_start_tls -> True
- create_users_on_logon -> True
- base_dn -> **dc=muthu4all,dc=com**
- bind_dn -> **cn=admin,dc=muthu4all,dc=com**
- bind_password -> Provide as per LDAP configuration
- user_filter -> **objectClass=***
- user_name_attr -> **uid**
- group_filter -> **objectClass=posixGroup**

HDFS ACL

- Linux Permissions

Owner	Group	Other
rwx	r-x	r-x
4+2+1	4+0+1	4+0+1
7	5	5

- What if User1 and User2 needs to be owner of the file?

Linux Permissions

Permissions user group

```
-rw-rw-r-- 1 hdfs hdfs 322930496 Mar 23 14:52 20newsgroup.txt
```

user other group

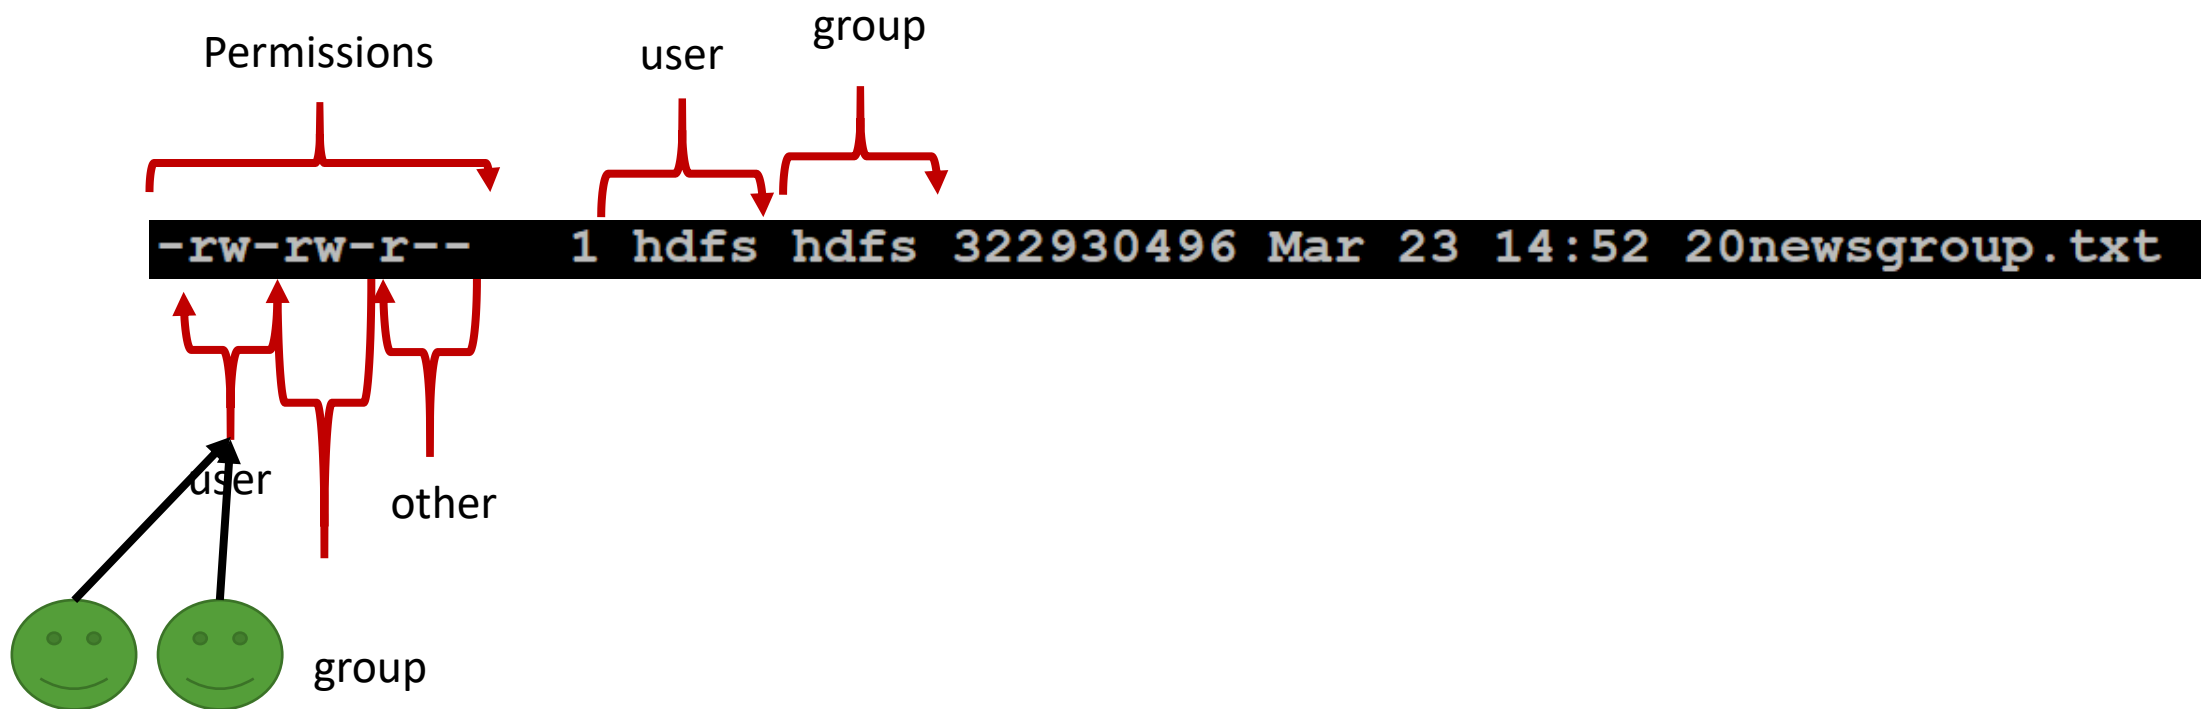
The diagram illustrates the Linux permissions for the file `20newsgroup.txt`. The permissions are `-rw-rw-r--`. Red arrows and brackets indicate the following breakdown:

- Permissions:** A bracket above the first three characters (`-rw-rw-r--`) points to the label "Permissions".
- user:** A bracket above the first two characters (`-rw`) points to the label "user".
- group:** A bracket above the next two characters (`-rw`) points to the label "group".
- other:** A bracket above the last character (`-r--`) points to the label "other".
- user:** A bracket below the first two characters (`-rw`) points to the label "user".
- group:** A bracket below the next two characters (`-rw`) points to the label "group".
- other:** A bracket below the last character (`-r--`) points to the label "other".

Scenario


- User Kumar should have read and write permission
 - Kumar part of marketing team.
 - All others in marketing team will have read permission.
 - Others outside marketing team should not have neither read nor write.
-
- RW_+R__+___ Owner will be Kumar and group will be marketing
 - What if another user in marketing needs write permission?

Linux Permissions



Properties - ACL

- `dfs.permissions.enabled = true`
- `dfs.permissions.superusergroup = supergroup`
- `dfs.namenode.acls.enabled = true`



ACL Options

- -R: List ACLs recursively.
- -b: Revoke all permissions except the base ACLs for user, groups and others.
- -k: Remove the default ACL.
- -m: Add new permissions to the ACL.
- -x: Remove only the ACL specified.
- <acl_spec>: Comma-separated list of ACL permissions.
- --set: Completely replace the existing ACL. Previous ACL entries will no longer apply.

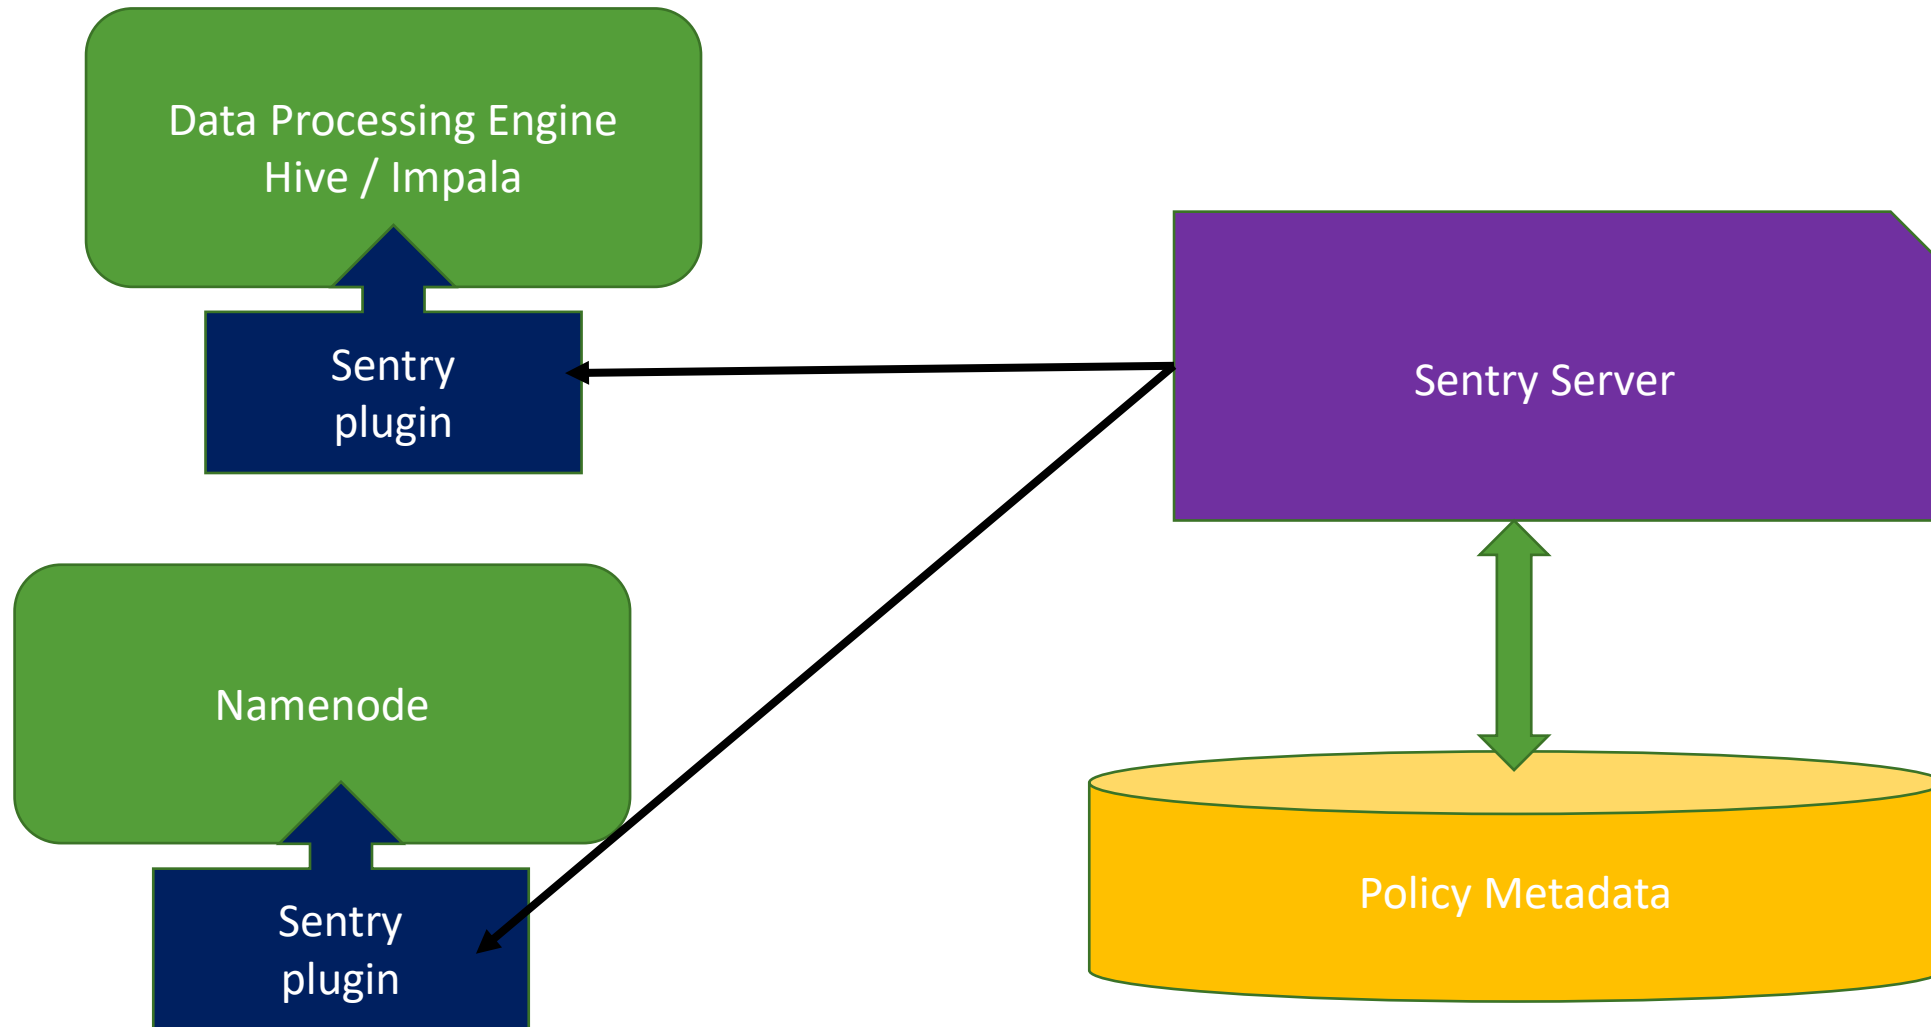
Order of evaluation of ACL Entries

- User is file owner – Owner permission bits are enforced
- Named user ACL entry.
- Member of file's group
- Named group in ACL entry (Union of previous entry)
- If none then other permission bits are enforced

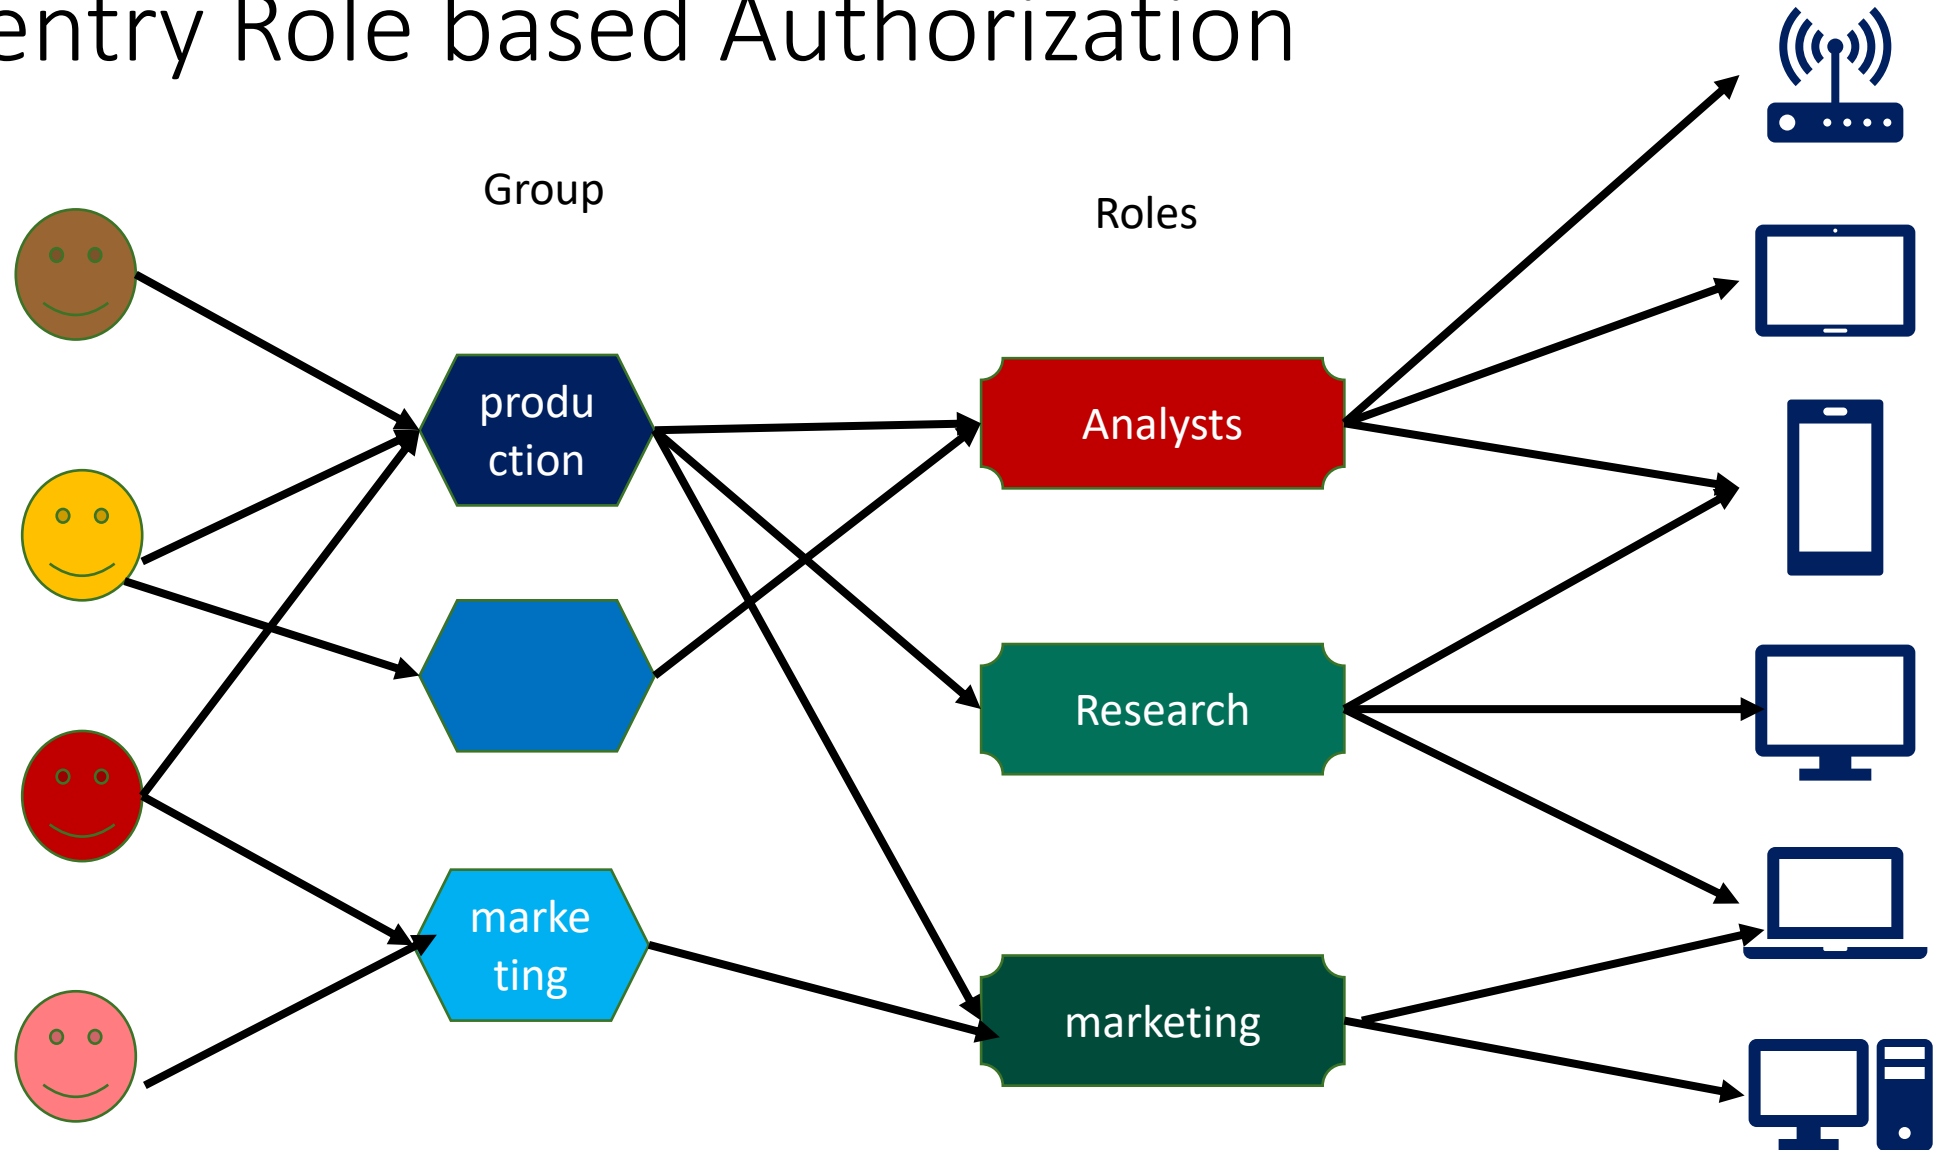
Apache Sentry

- Enforces fine grained authorization
- Role based authorization to data and metadata
- Integrates with Hive and HDFS
- Supports Impala and many more components
- Developed by Cloudera
- Collection level, document level authorization could be provided

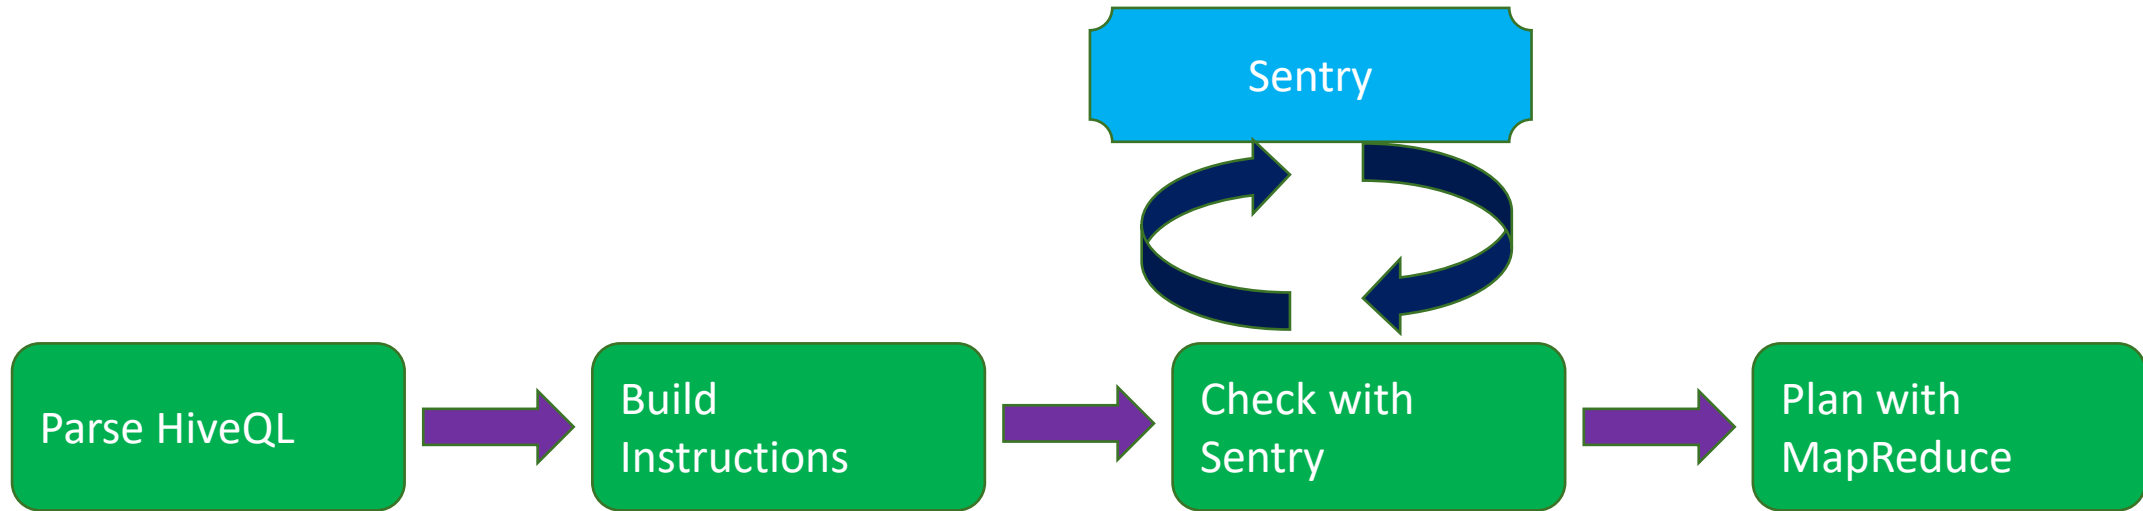
Sentry Components



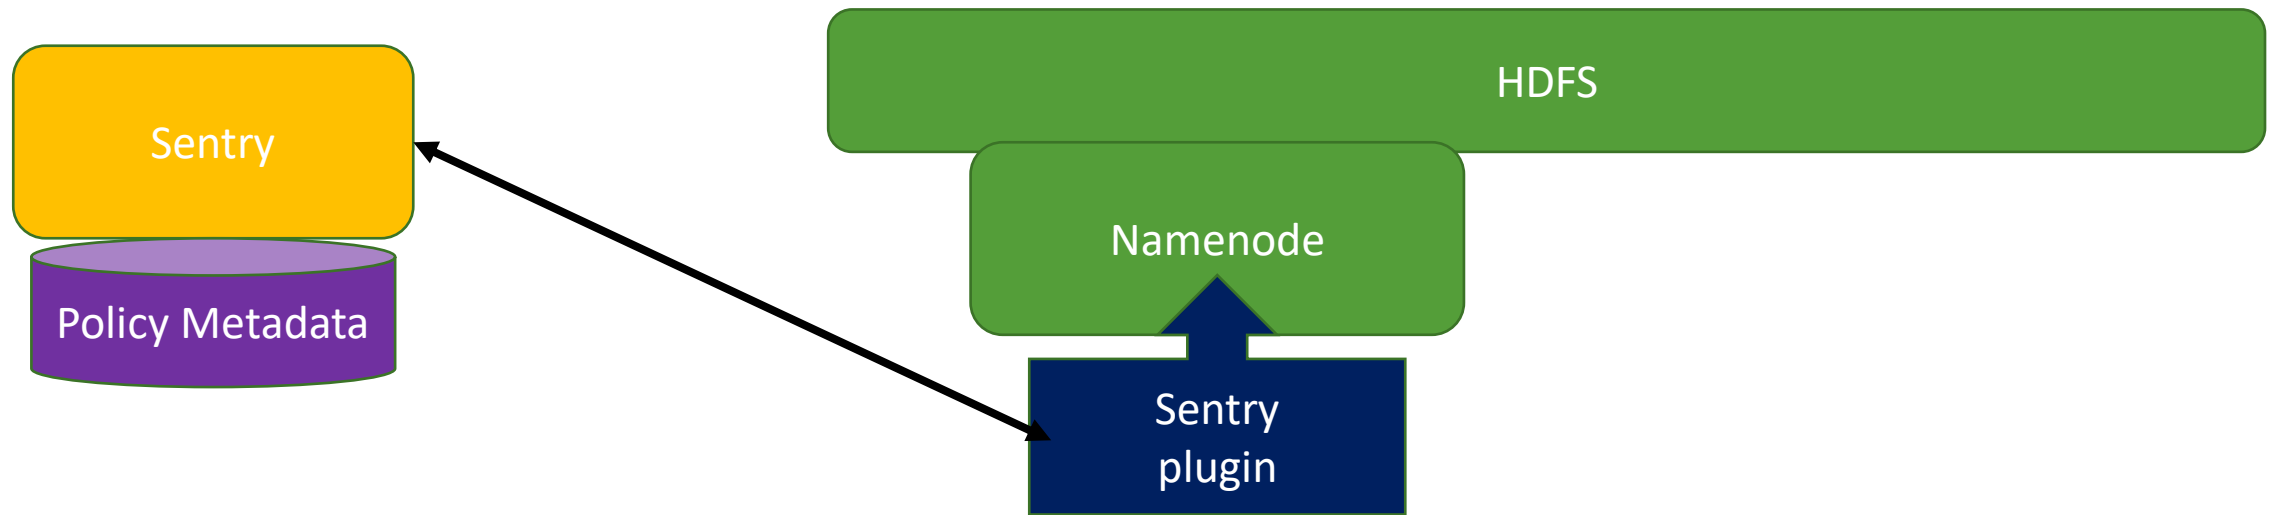
Sentry Role based Authorization



Hive and Sentry



HDFS and Sentry



Impala

- Distributed Massive Parallel Processing (MPP) database engine
- Doesn't build on MapReduce
- Works with its own execution engine
- Native analytic database
- Developed by Cloudera

Impala Vs Hive

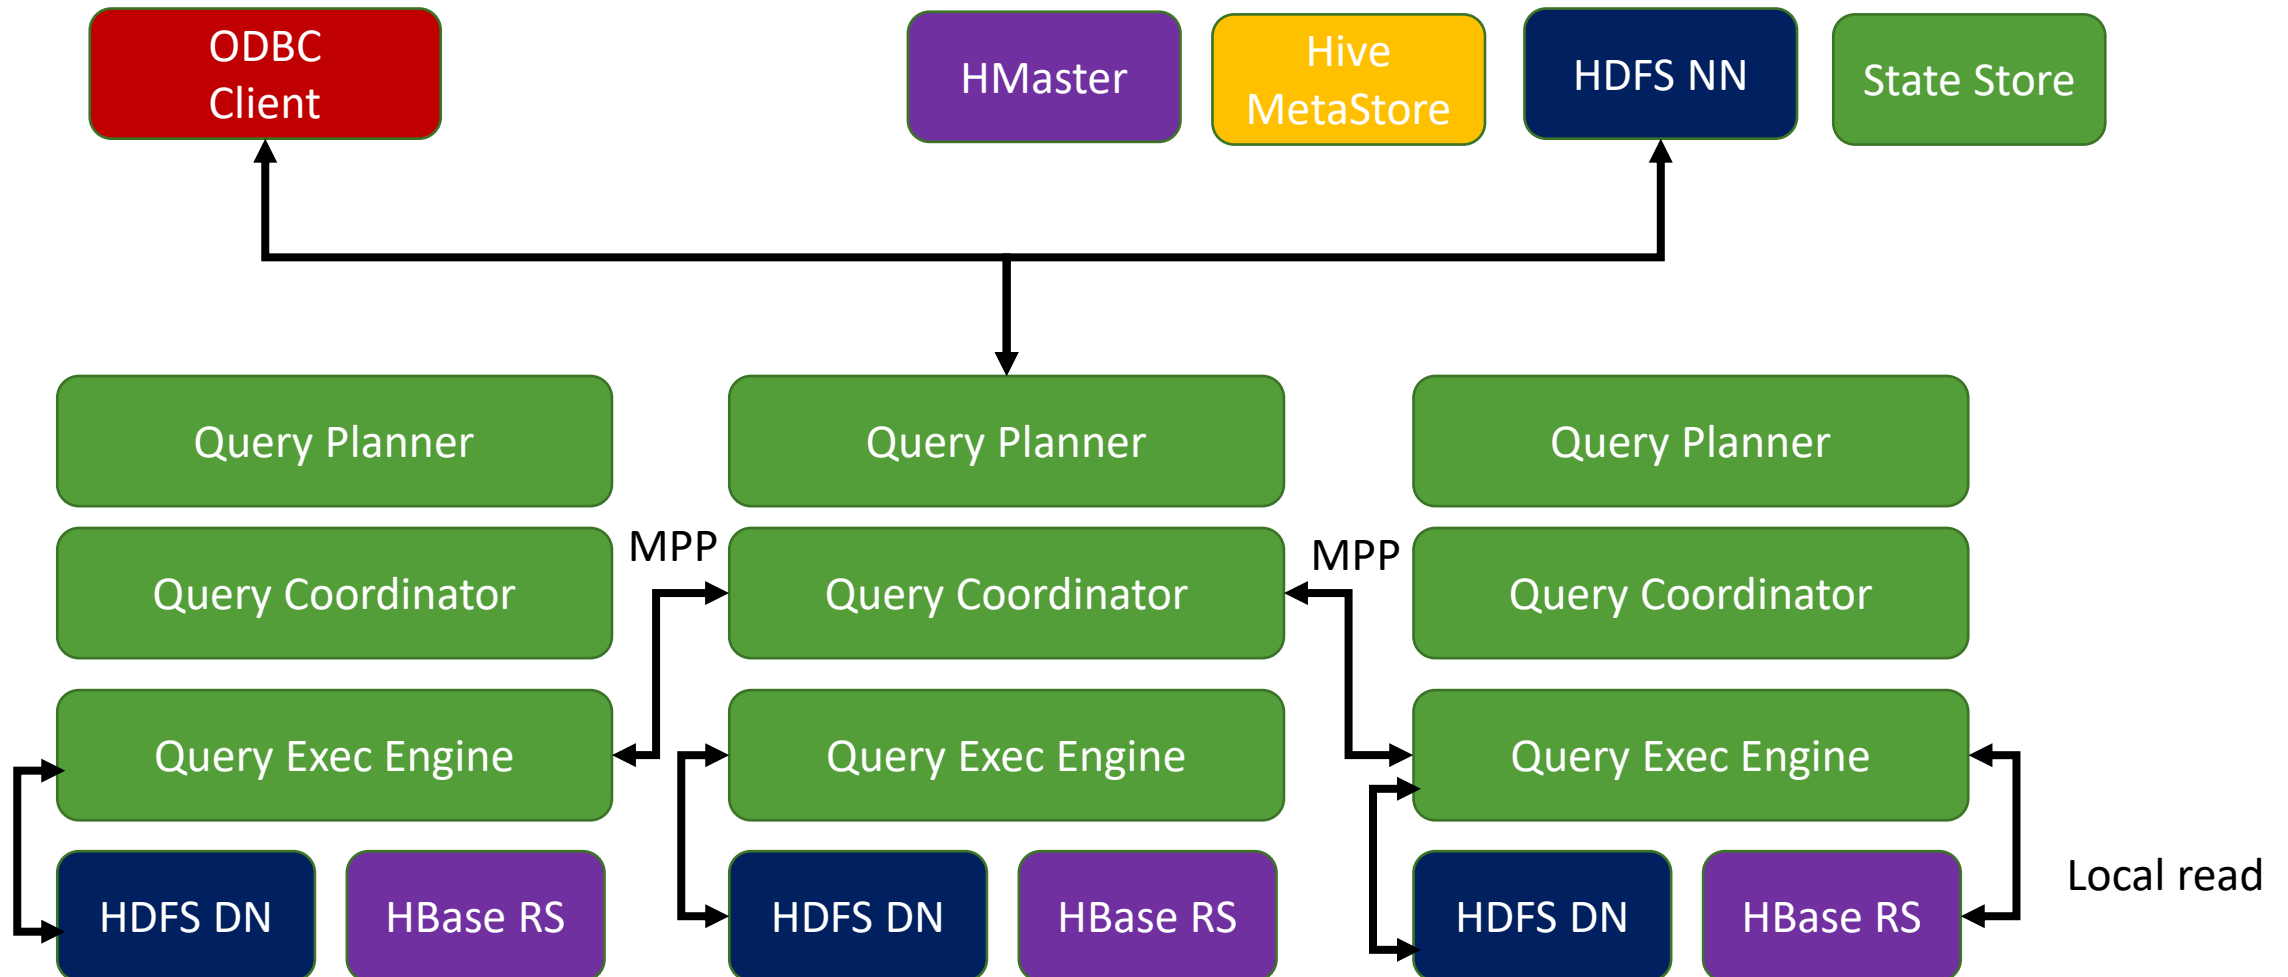
Impala

- Runtime code generation
- Daemons already running. Startup overhead is less
- Integrates with Sentry easily
- Latency is less
- Relies on Massive Parallel Processing

Hive

- Generates at compile time
- Cold start – On execution of query
- Relies on Hive metastore for security
- Latency is high
- Relies on MapReduce

Impala Architecture



Impala Introduction

- Unified Storage, Metastore, Security, Client Interface
- Real time SQL queries, natively distributed query engine

Impala Roles

- Impala Daemon (impalad)
 - Runs on each Datanode
 - Reads and writes to data files
 - Accepts queries transmitted from the shell
- Impala StateStore (statestored)
 - Need only one and not mandatory
 - Required when any impala daemon goes down to resume its functions.
 - Main function is to provide high availability and load balancing of impalad
- Impala Catalog Server (catalogd)
 - Mostly runs along with statestored
 - Avoids the need to refresh or invalidate the metadata when it gets changed

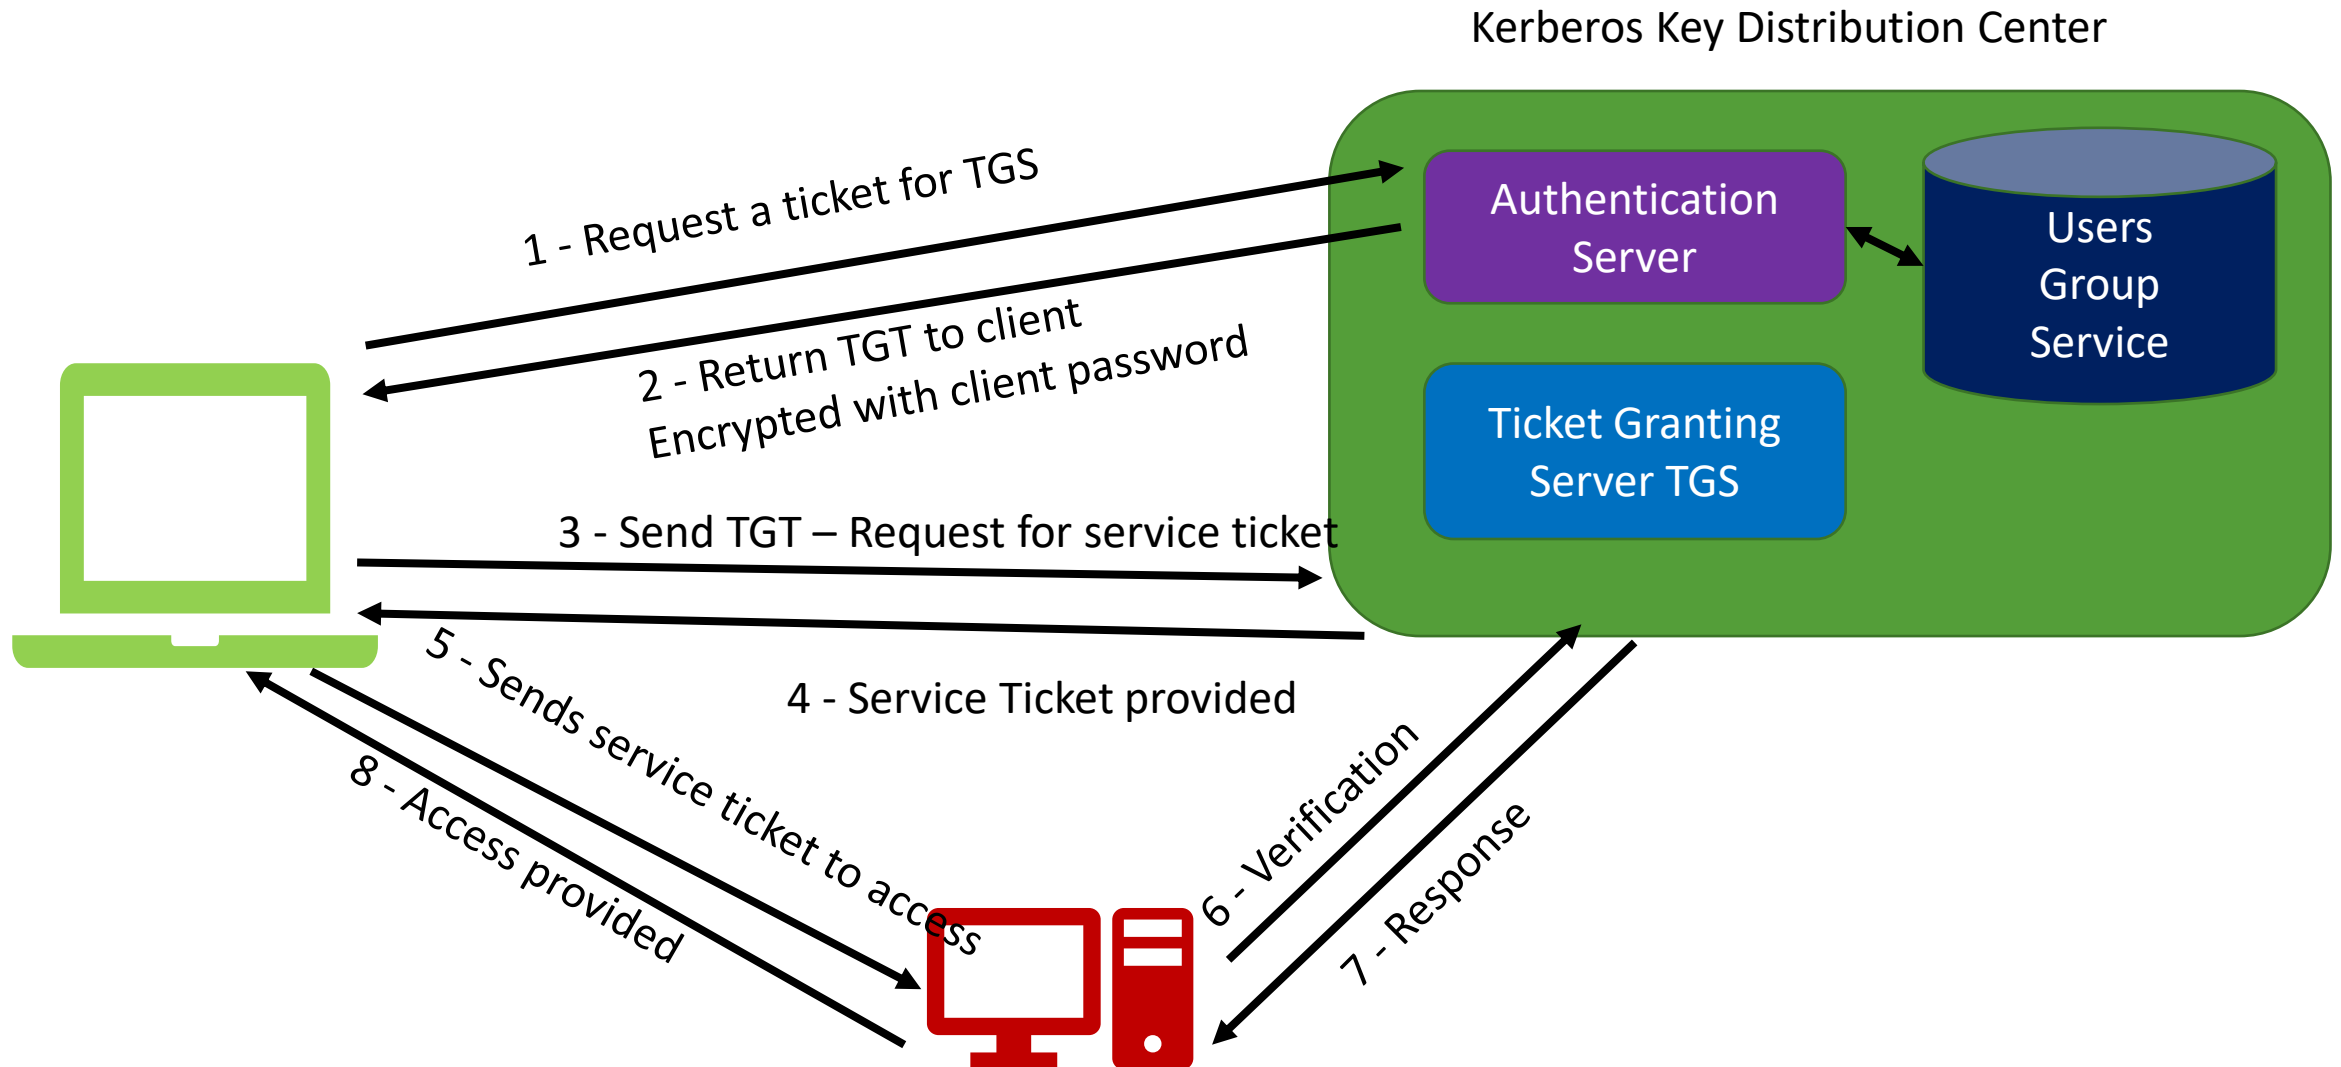
Cloudera Manager – OpenLDAP Integration

- Authentication Backend Order → External Then Database
- LDAP url ldap://<OpenLDAP hostname>:389
- LDAP bind user dn → cn=admin,dc=muthu4all,dc=com
- ldap bind password
- ldap user search filter → (uid={0})
- ldap user search base → ou=clouderausers, dc=muthu4all,dc=com
- ldap group search filter → objectClass=posixGroup
- group search base → ou=clouderausers,dc=muthu4all,dc=com
- ldap distinguished name pattern → uid={0},ou=clouderausers,dc=muthu4all,dc=com

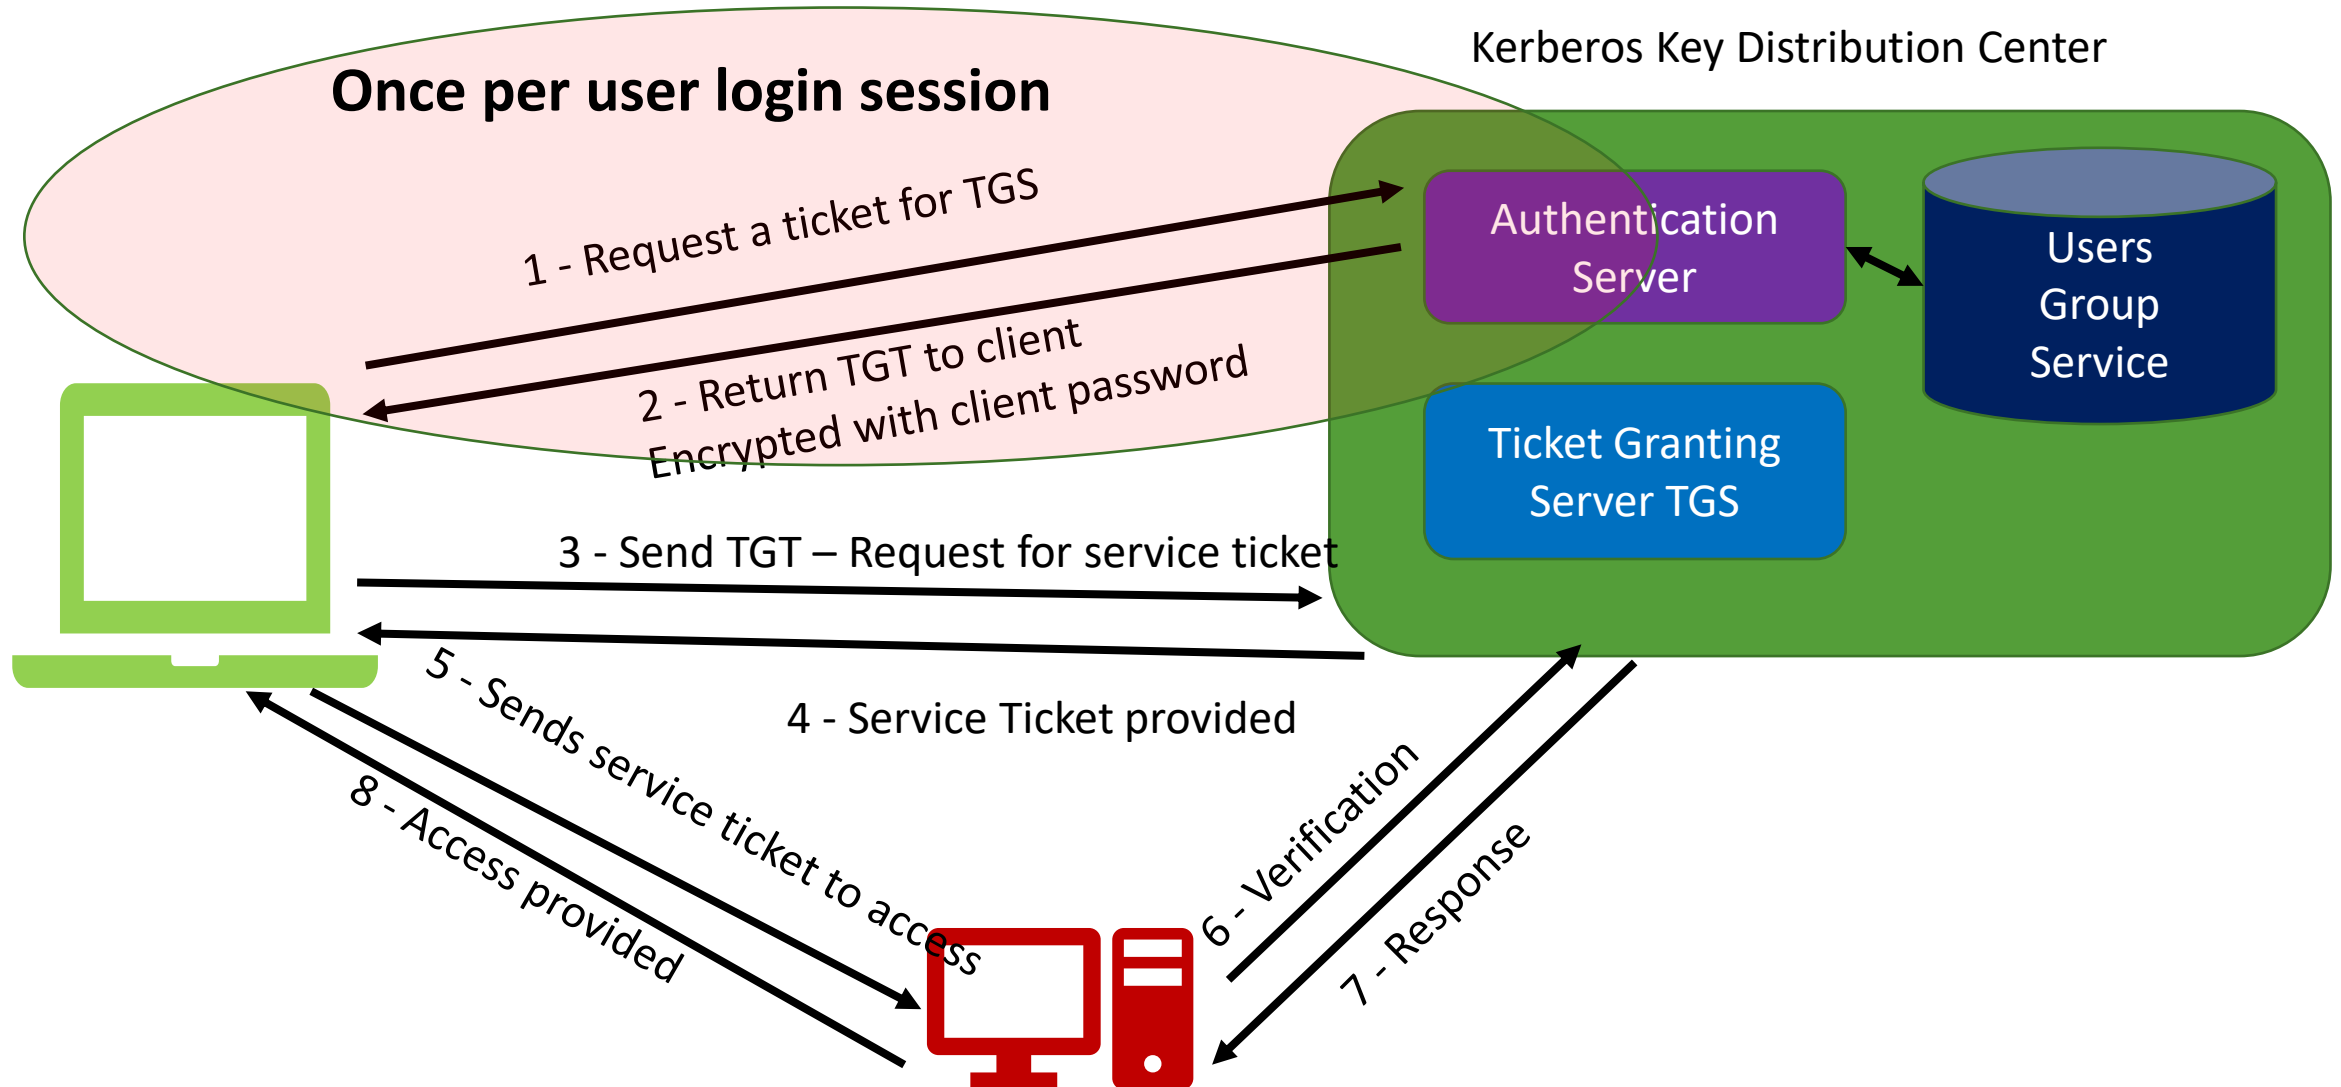
Kerberos Introduction

- Network authentication protocol
- Works based on tickets and key exchange mechanism
- Provides strong security on a non secure network
- Developed by MIT
- Only provides protocol.
- Should be integrated with external LDAP Servers like AD or OpenLDAP

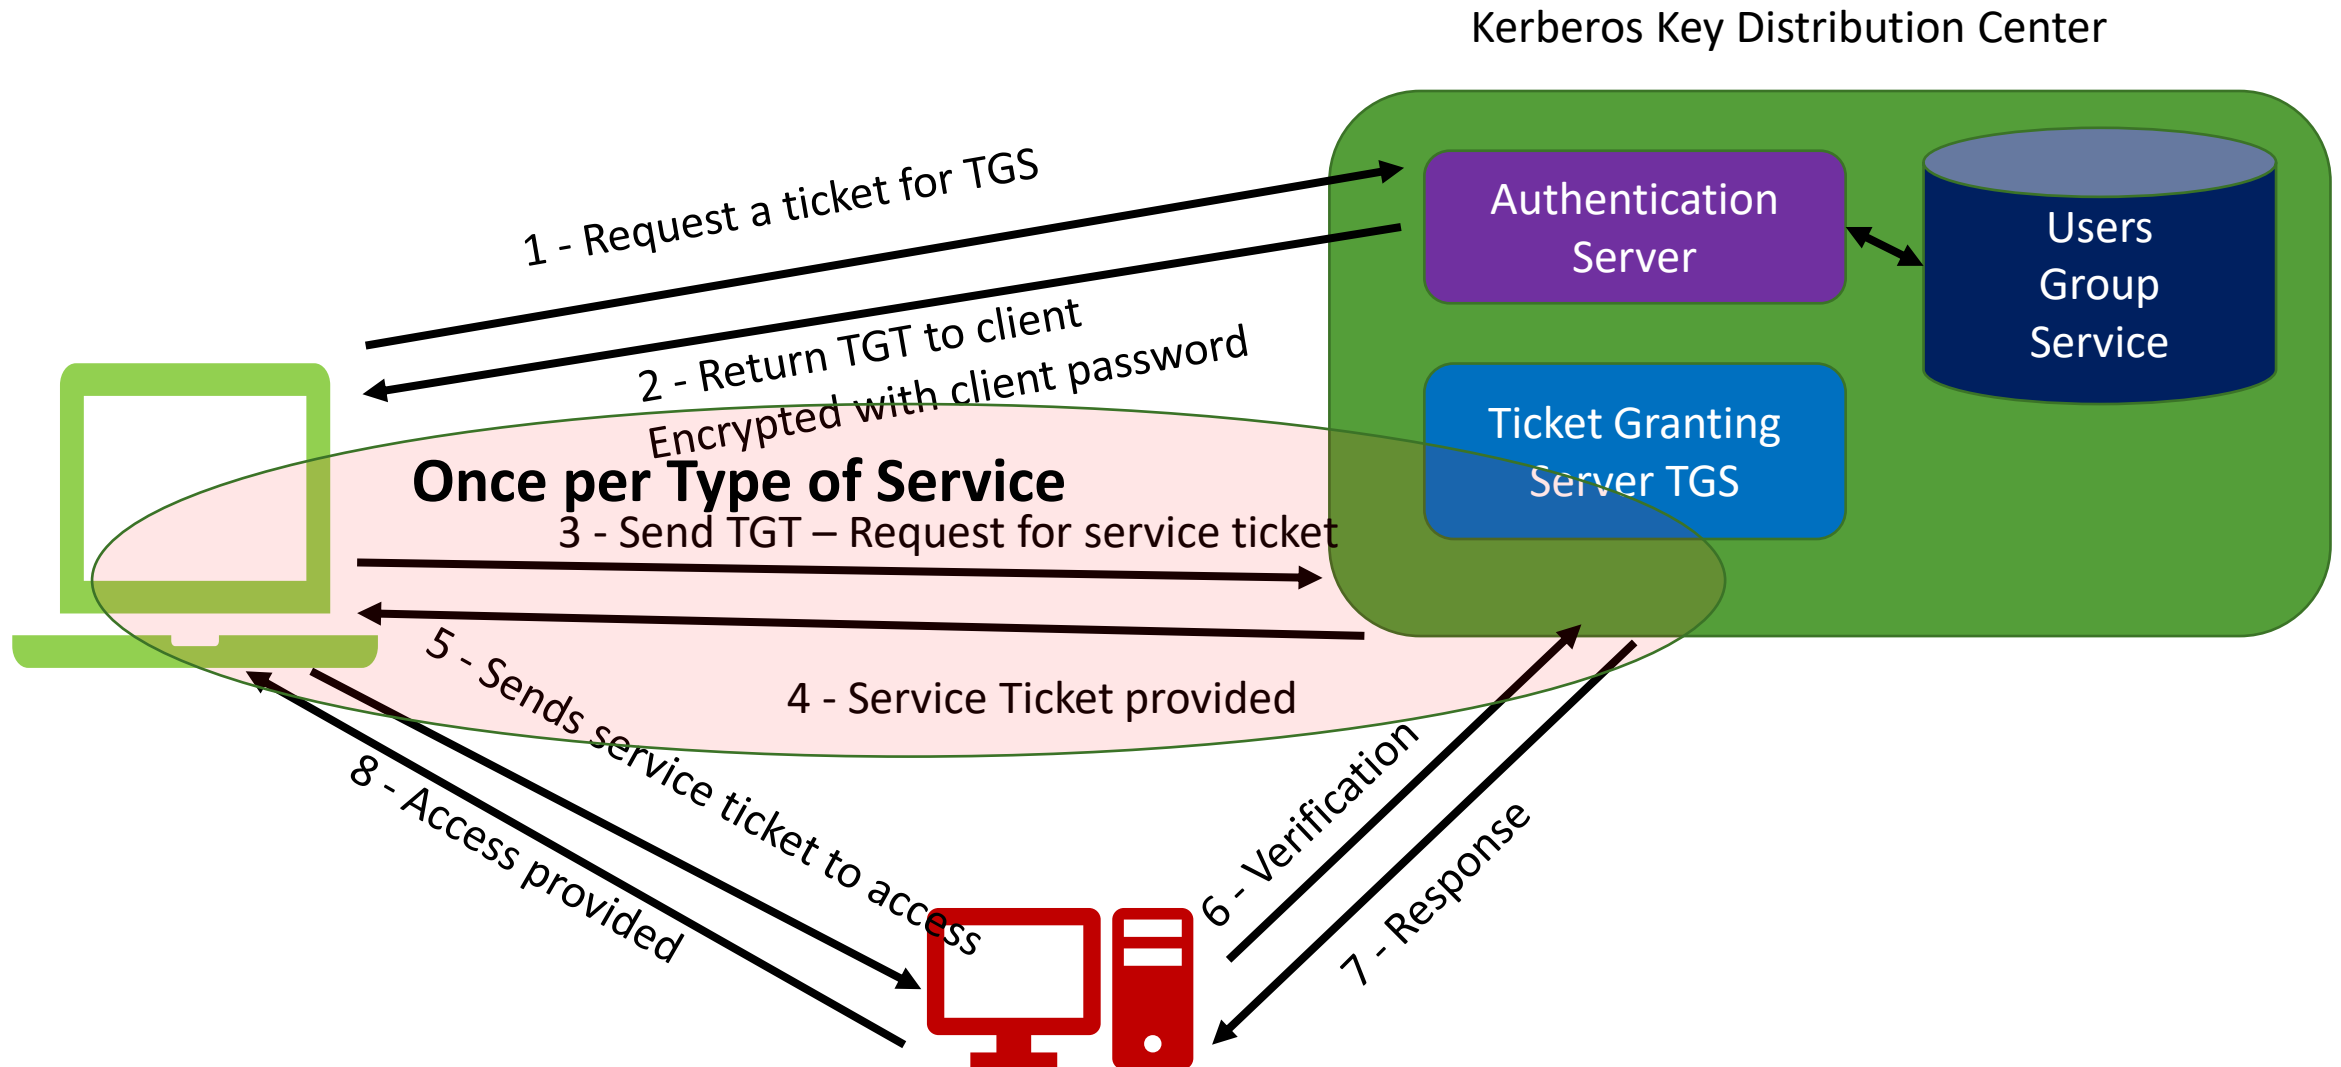
Kerberos Architecture



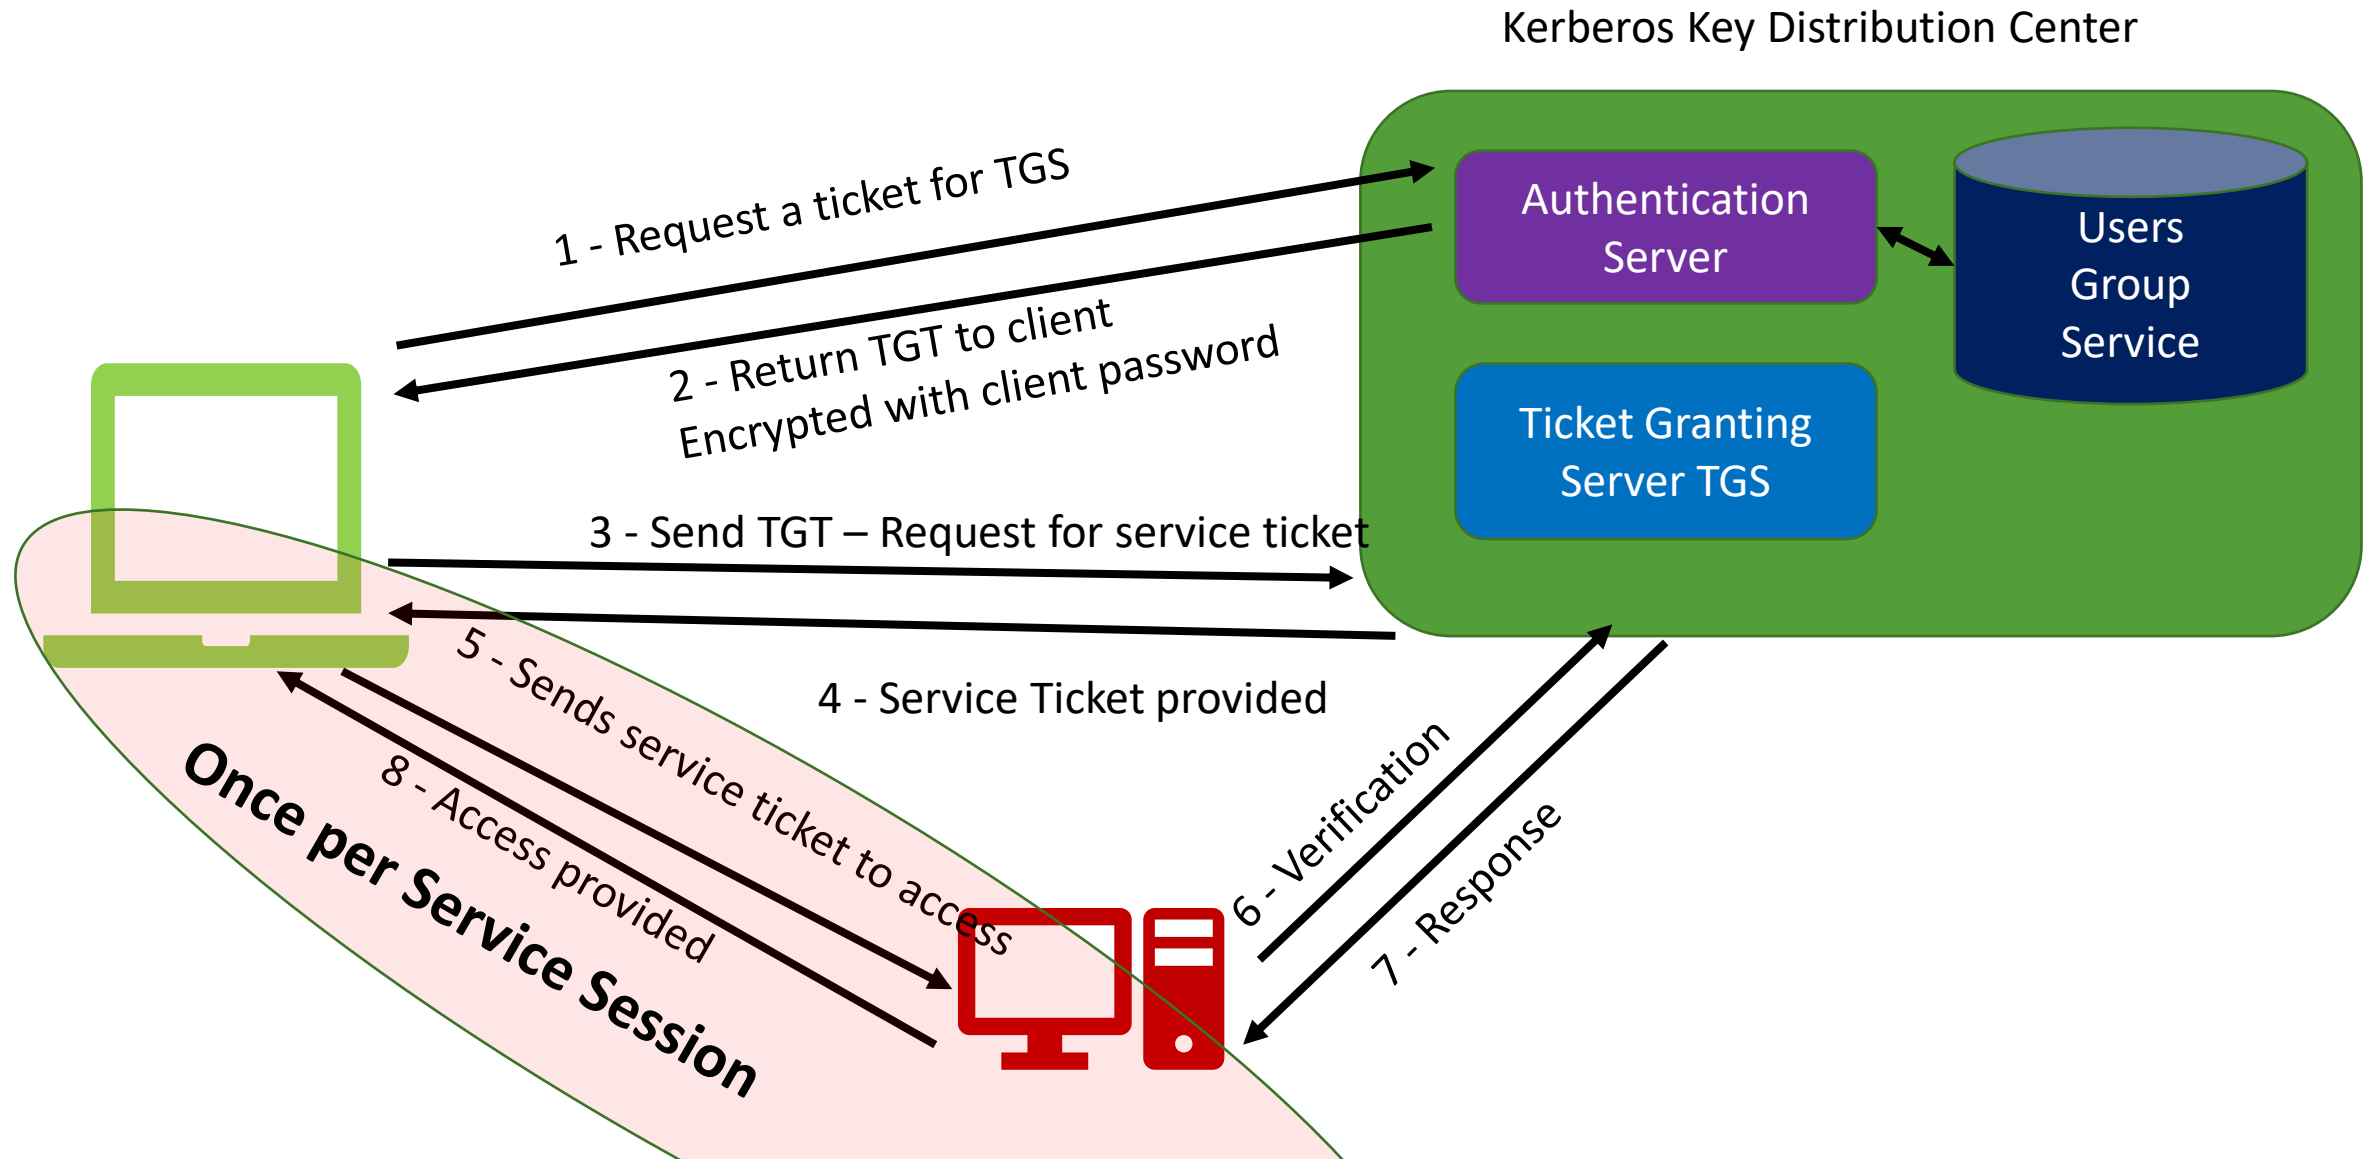
Kerberos Architecture



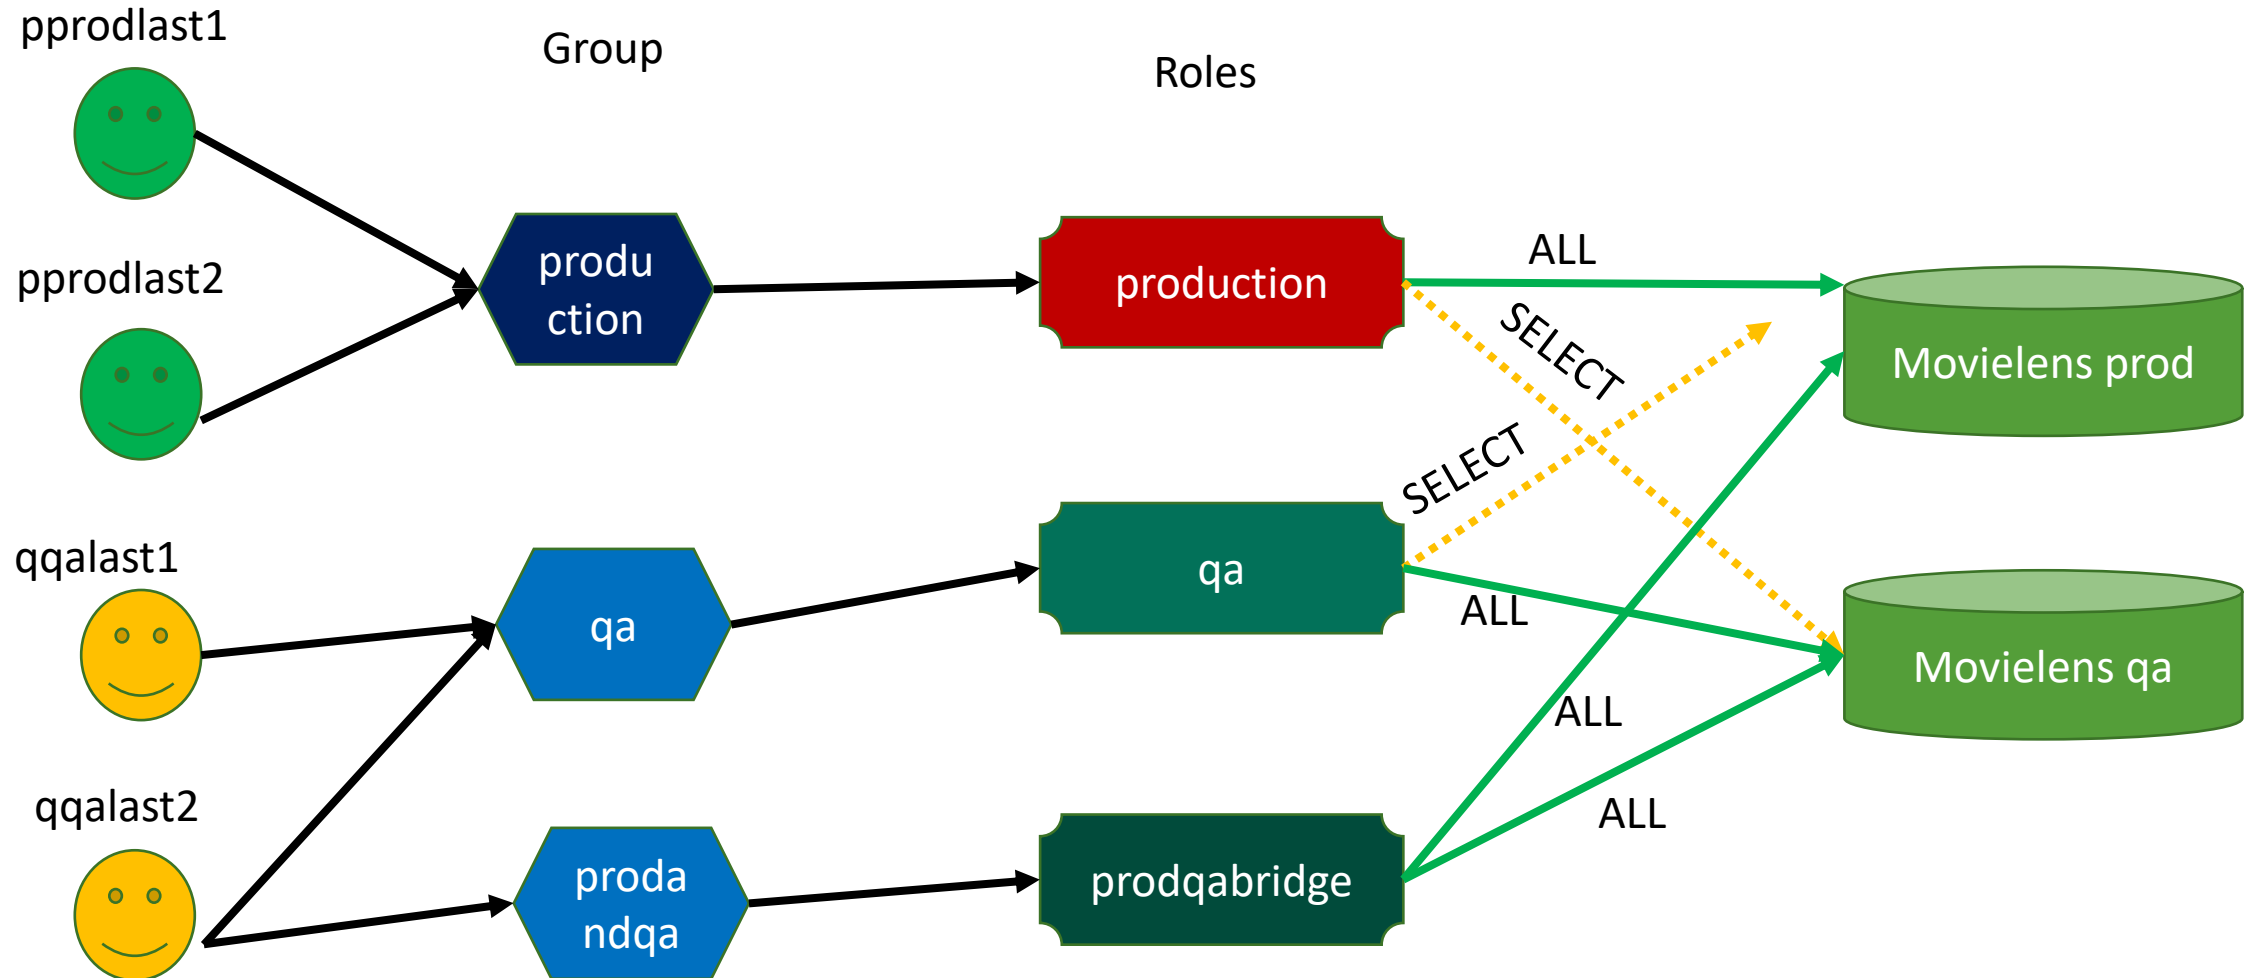
Kerberos Architecture



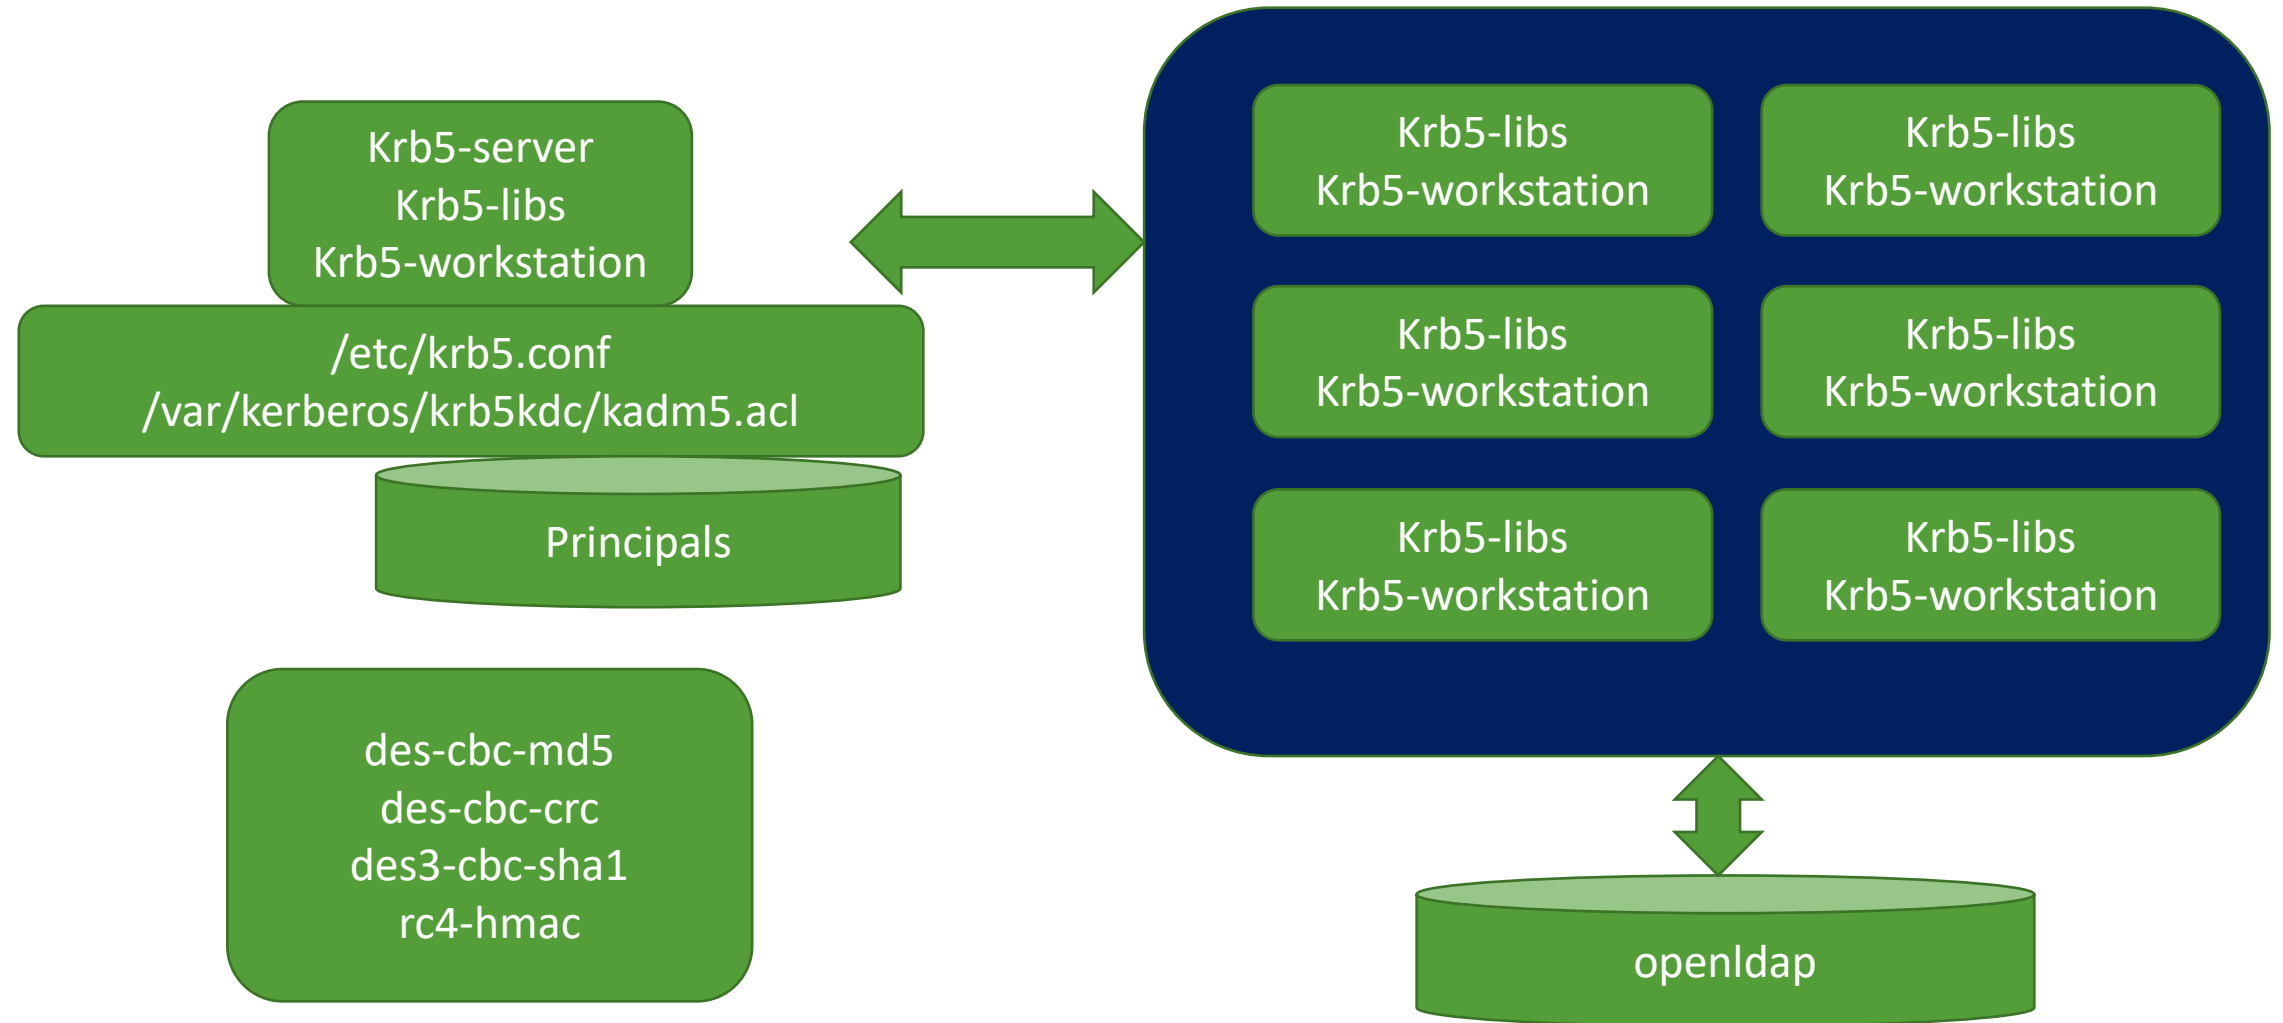
Kerberos Architecture



Sentry Role based Authorization



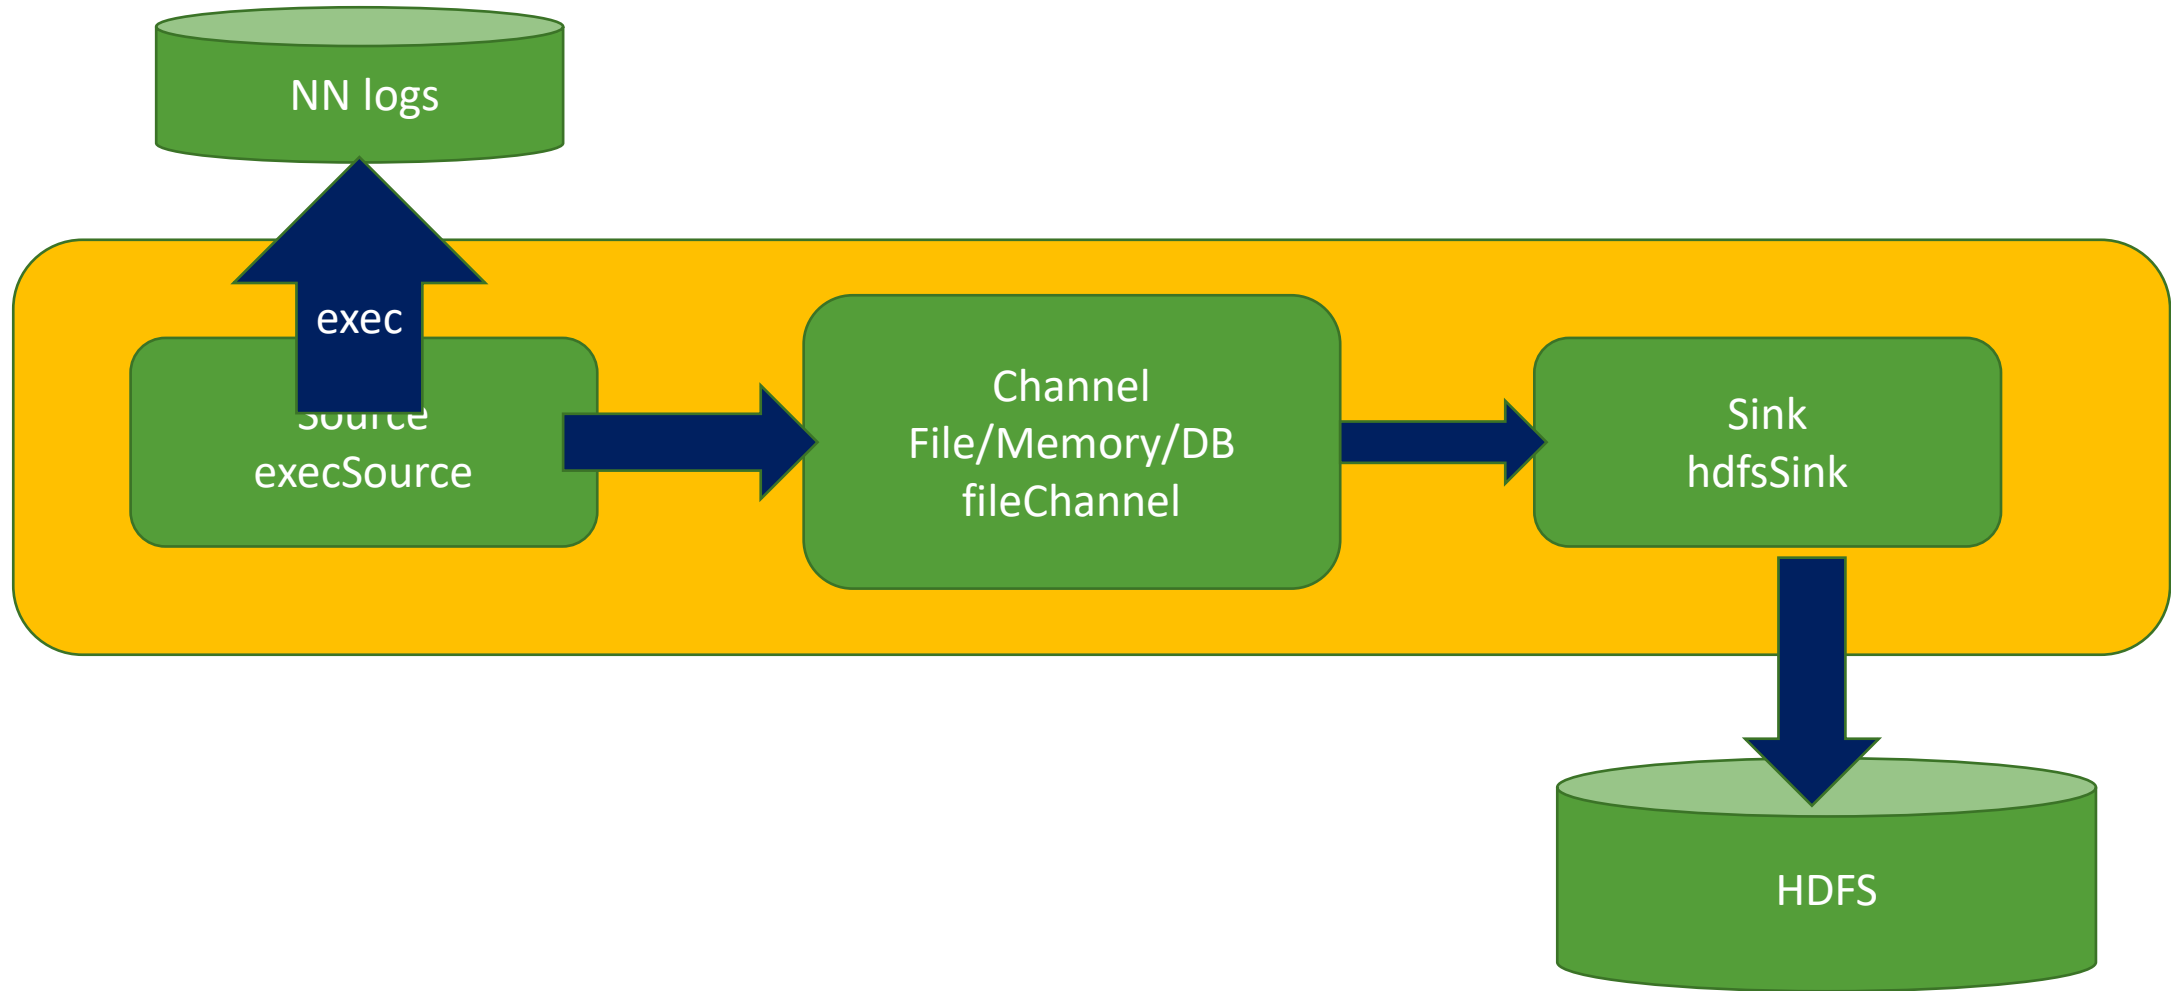
Kerberize Hadoop Cluster



Flume Introduction


- Data ingestion tool for collecting/aggregating/transporting data
- Works with streaming data as logs/Facebook events/ twitter tweets
- Aggregates/sinks data to a single location
- Reliable, Fault tolerant, Scalable, Manageable

Flume




Flume Multi Source



Machine 1 – M1 - 57




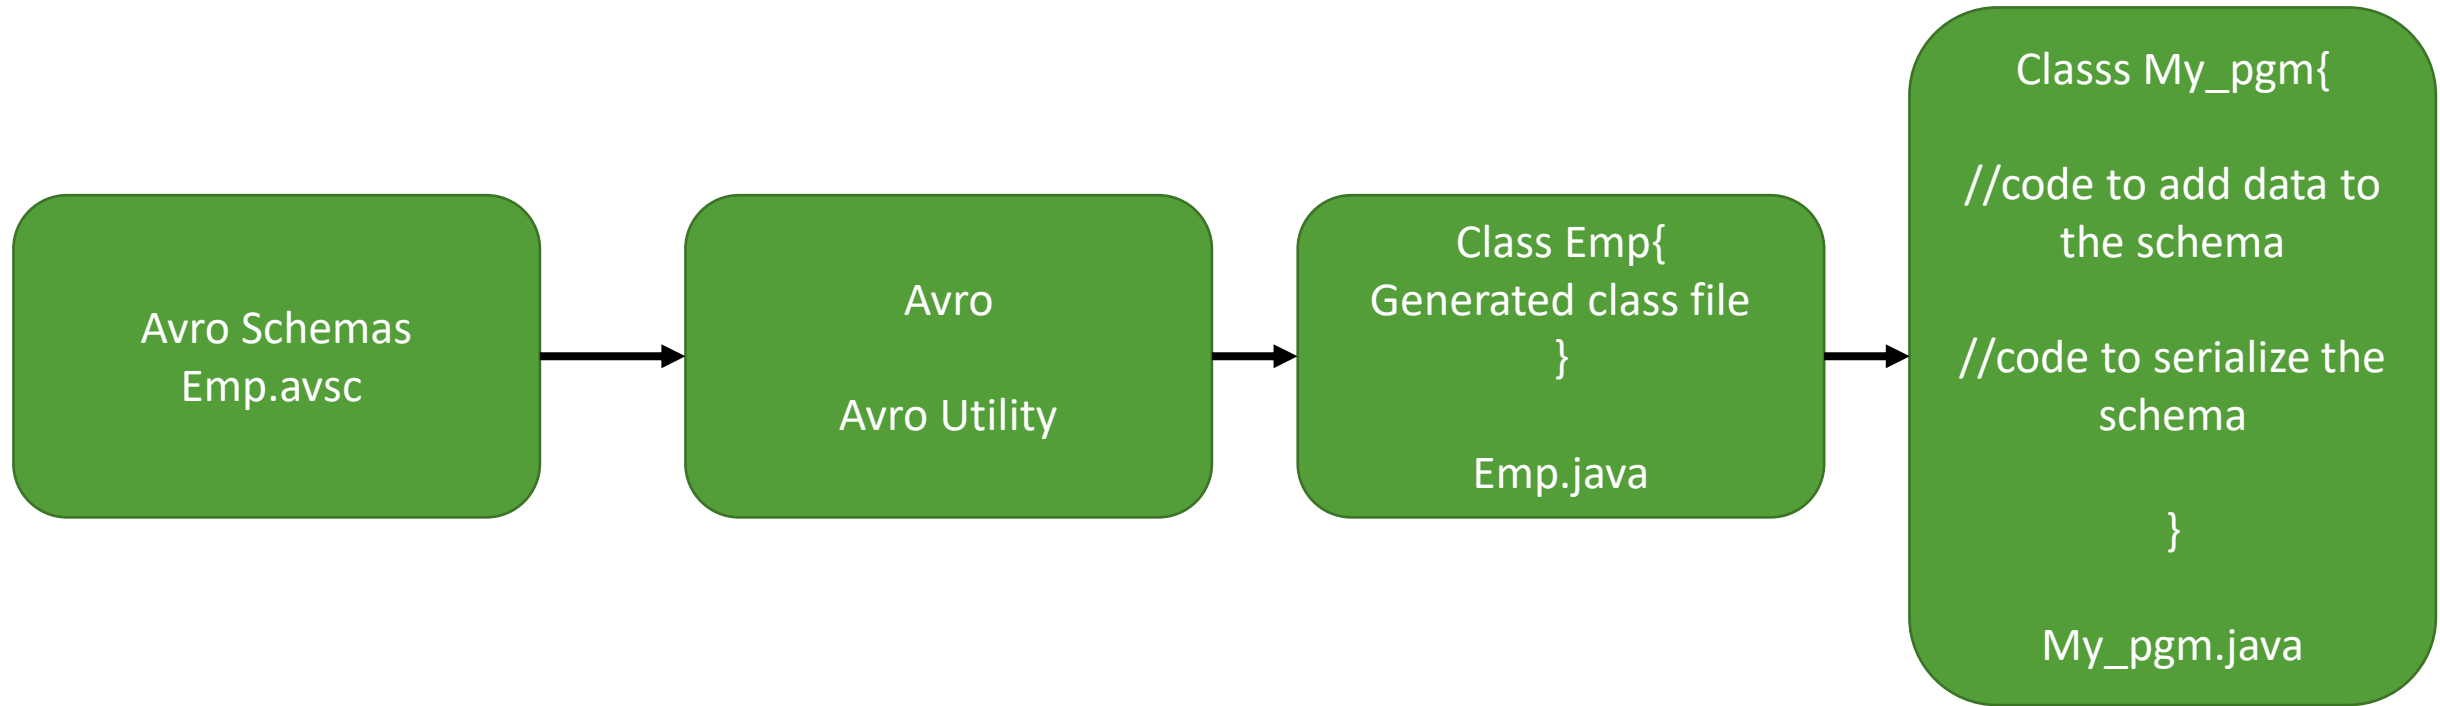
Machine 2 – M2 - 185



Collector – C1 - 96



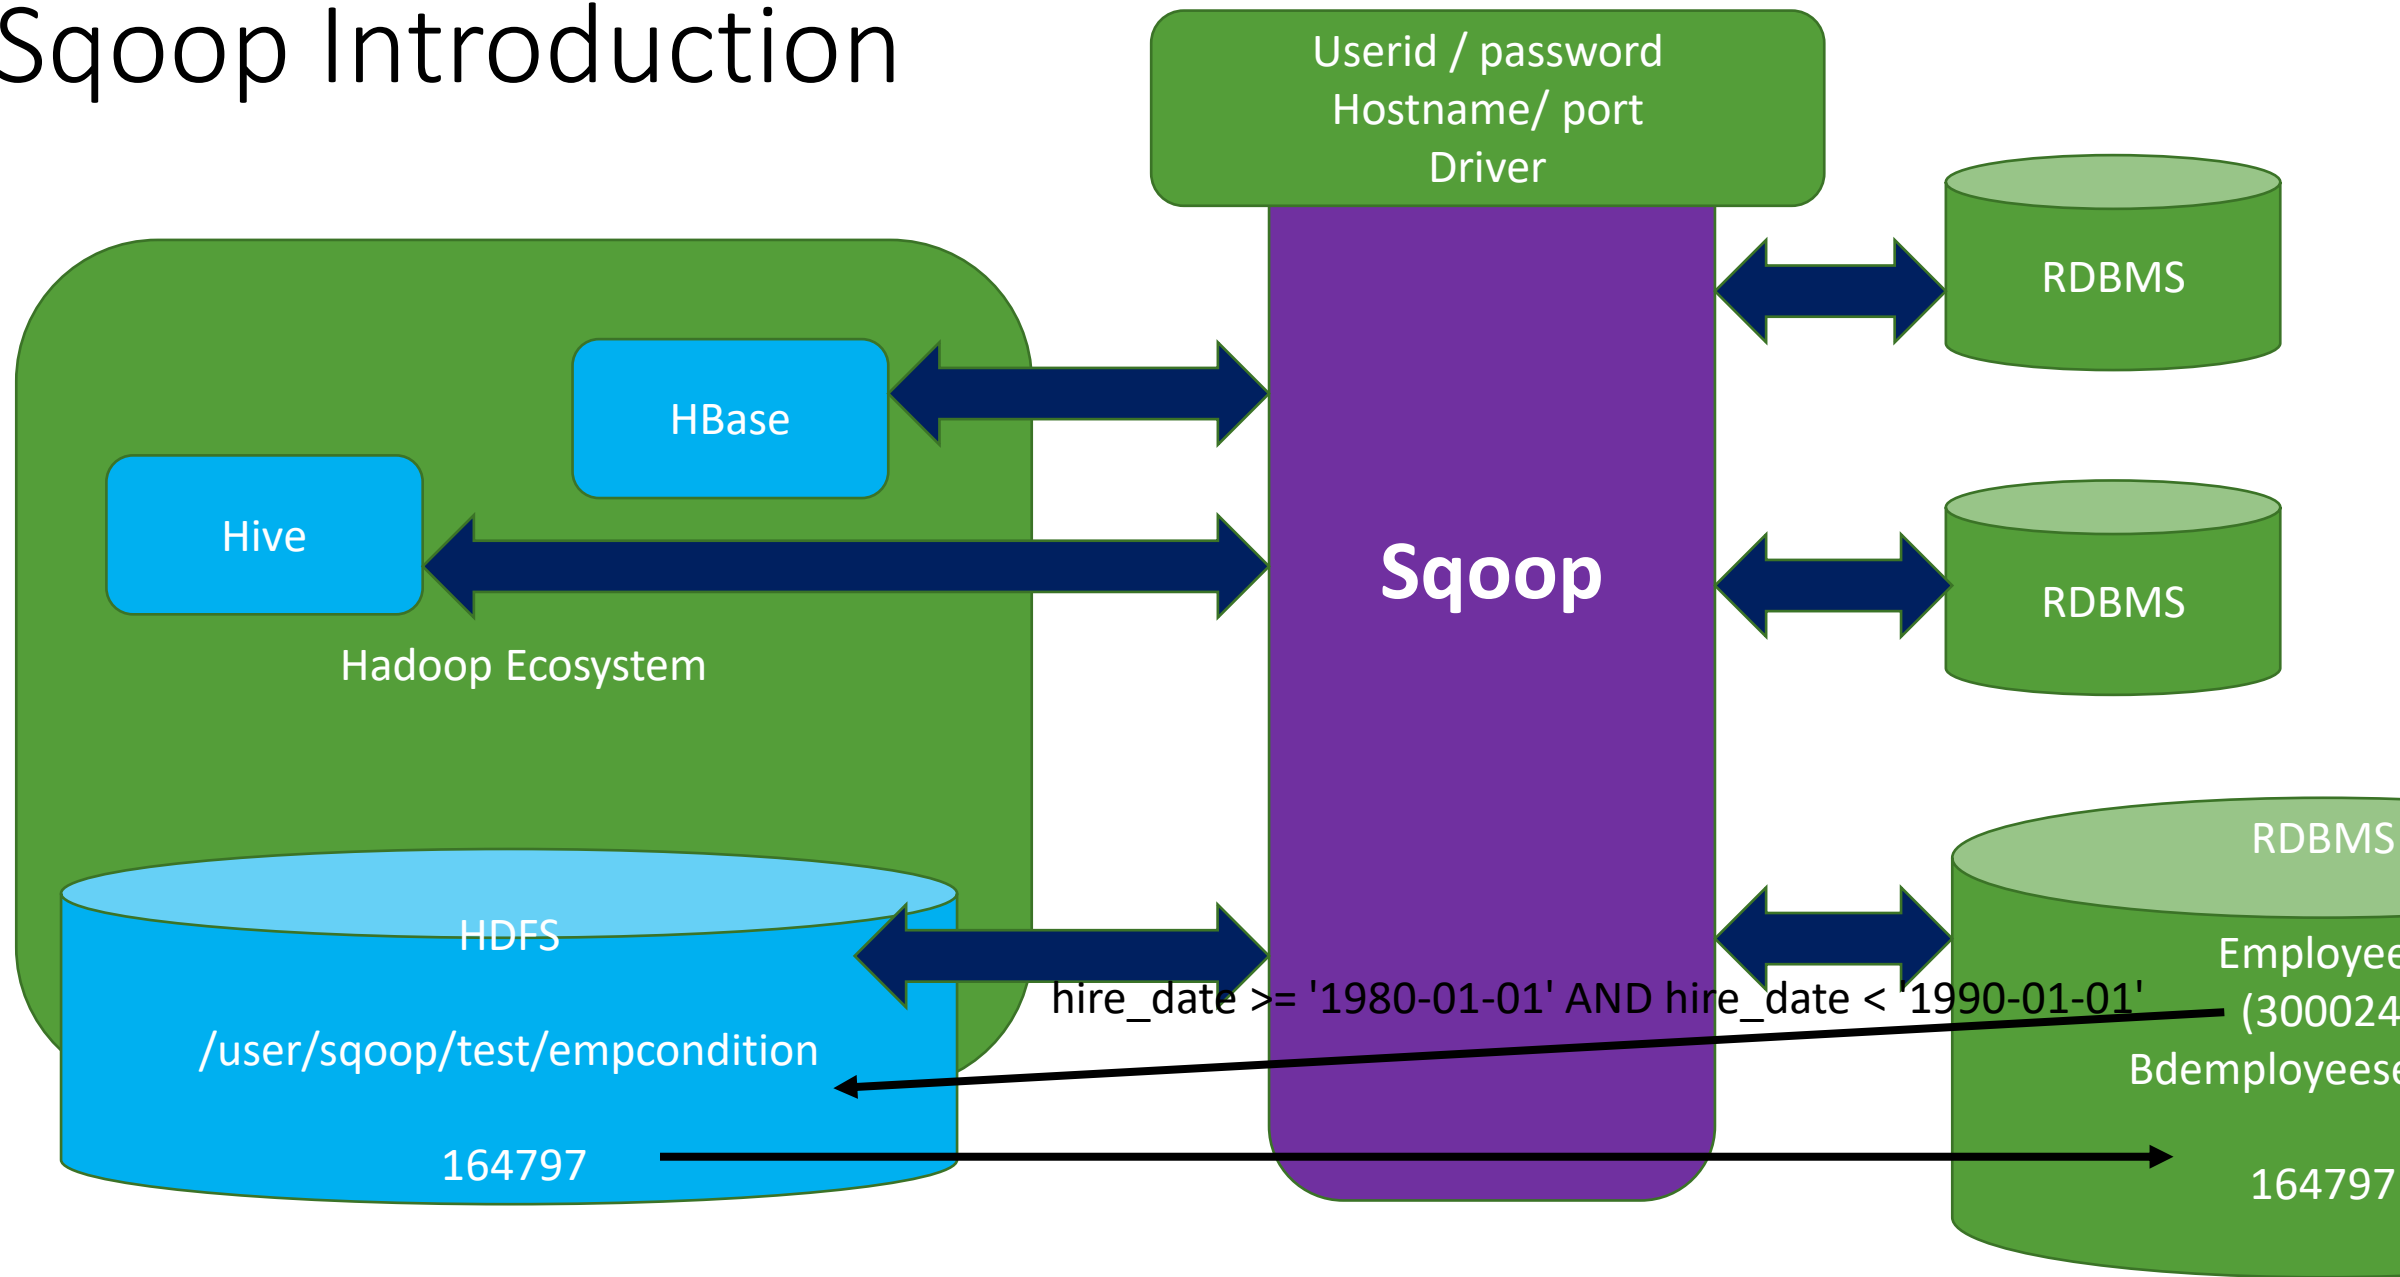




Sqoop2

- Works on client server model
- Connectors for major RDBMS not yet available
- Data transfer from RDBMS to Hive and Hbase not supported
- Data transfer from Hive and Hbase to RDBMS not supported
- Not yet integrated with HUE


Sqoop Introduction



HBase

- Distributed Column oriented NoSQL Database
- Works on top of HDFS
- Similar to Google's big table
- Data exists as key value pair.
- Group of key value pair together exists as column family
- N number of key value pair makes a row.
- Rowid, column family, key makes unique combination of any cell
- Works on master worker architecture

Hbase Architecture



Regional Server

